

System Programming

(ELEC462)

Introduction

Dukyun Nam
HPC Lab @ KNU

Contents

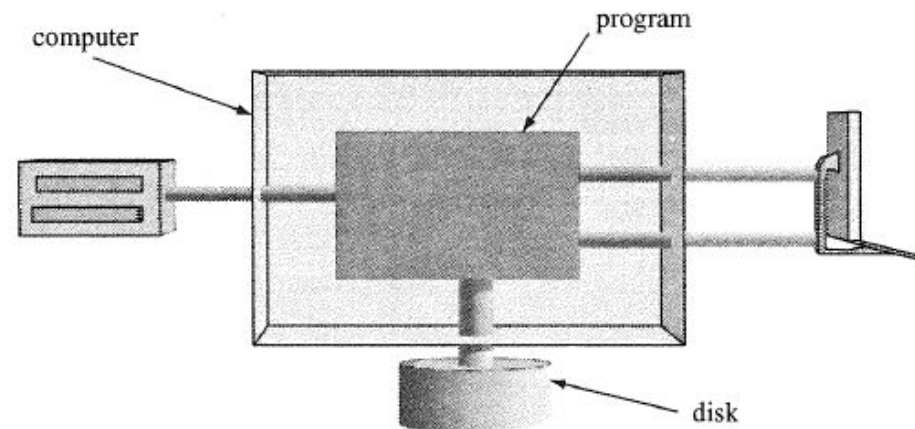
- Introduction
- System Resources
- Unix from the User Perspective
- Unix and Linux
- Summary
- Appendix: `vi` editor

Introduction

- What is ***system programming***, or ***systems programming***?
 - The activity of **programming** computer **system software**
 - Computer *programming* is the process of performing a particular computation
 - *System software* is software designed to provide a platform for other software
- Application programming vs system programming
 - **Application programming** aims to produce software which provides services to the *user* directly
 - e.g., Word processor
 - **System programming** aims to produce software and software platform which provides service to *other softwares*

Introduction (cont.)

- The simple program model
 - Many programs are based on the model as shown below
 - A program: a piece of code that runs in a computer
 - Data go into the program; the program does some tasks with the data, and the output comes out of the program
 - A human may type at a keyboard and monitor; the program may read or write to a disk; the program may send data to a printer; many possibilities!

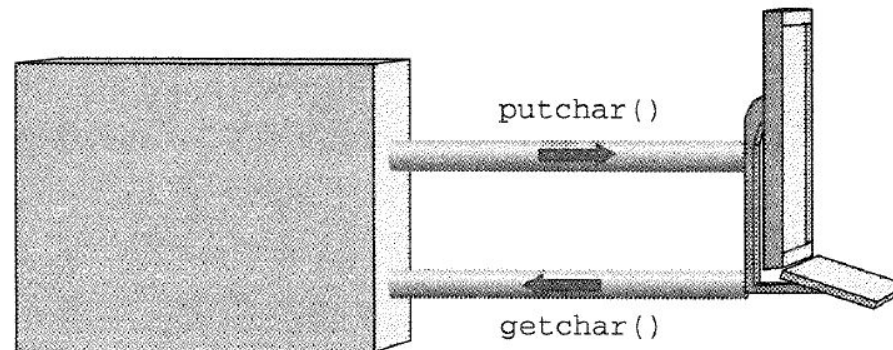


< An application program in a computer >

Introduction (cont.)

- The simple program model
 - The keyboard and screen are both connected to the program
 - Tip) GCC reports “errors” with the source file name and line number

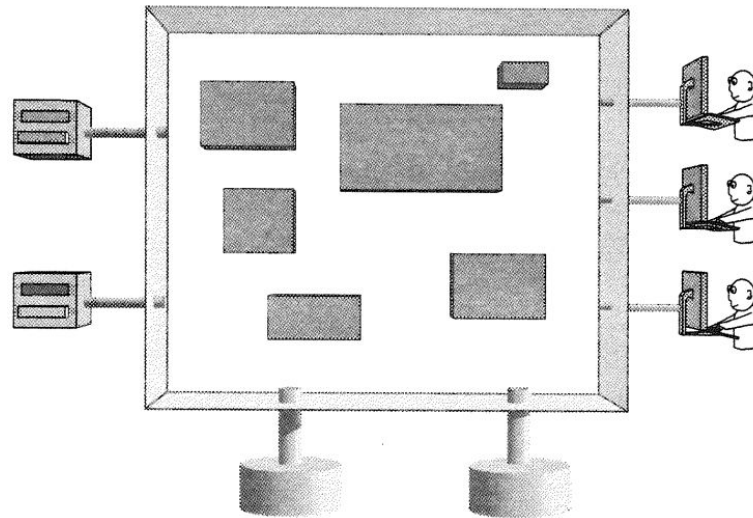
```
/* copy from stdin to stdout */  
void main(){  
    int c;  
    while ( (c = getchar() ) != EOF)  
        putchar(c)  
}
```



< How application programs see user I/O >

Introduction (cont.)

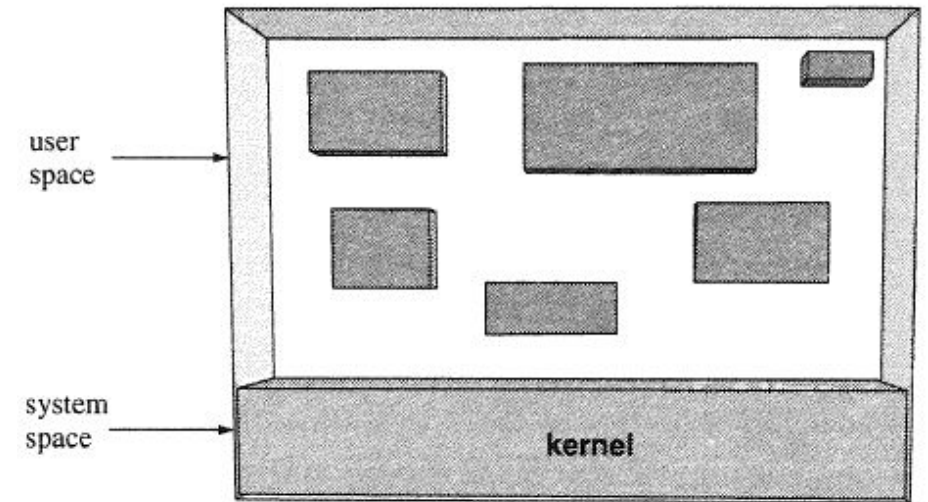
- Face reality
 - There are several keyboards and displays, there may be several disks, one or more printers, and there are certainly several programs running at once
 - Programs that get input from *the keyboard* and send output *the display* or to *the disk* work fine



< Reality: many users, programs, and devices >

Introduction (cont.)

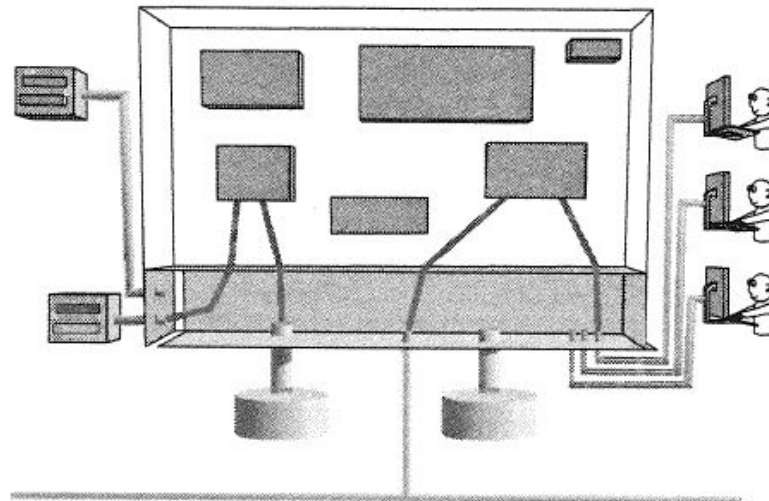
- The role of the operating system
 - To *manage and protect* all the resources and to connect various devices to various programs
 - Operating system: called a **kernel**
 - Memory: space to store programs and data
 - User space
 - Contains users programs
 - System space
 - The operating system is stored



< An operating system is a *program*
>

Introduction (cont.)

- Providing services to programs
 - If a program wants to connect to or control any the devices, it needs to ask the kernel
 - Access to the external objects are *services* the kernel provides to user programs



< The kernel manages all connections >

System Resources

- Processors
 - A program: a set of (machine) instructions compiled from source code
 - A processor (CPU): the hardware executing instructions
 - The kernel assigns processors to programs
 - The kernel starts, pauses, resumes, and ends the execution of a program on a processor
- Input/output
 - All data that flow in and out of a program go through the kernel
 - e.g., Data coming from/going out to users at terminals or disks cannot travel without passing through the kernel

System Resources (cont.)

- Process management
 - Process: a program in execution
 - Consisting of i) memory, ii) open files, and iii) other system resources (e.g., socket, device driver, etc.) needed
 - The kernel creates new processes and schedules them to work cooperatively
- Memory
 - Computer's memory is a resource.
 - The kernel:
 - Keeps track of which processes are using which sections of memory, and
 - Protects the memory of one process from being damaged by another process.

System Resources (cont.)

- Devices
 - One can attach all sorts of devices to a computer
 - e.g., Mice, scanners, printers, keyboards, USB drives, iPhones, ..
 - The kernel provides access to devices and takes care of complexities that may derive from communications with the attached.
- Timers
 - Some programs depend on time.
 - e.g., Periodic tasks (by `cron`), sleep, elapsed time, ...
 - The kernel makes timers available to processes.

System Resources (cont.)

- Interprocess communication (IPC)
 - In a real world, people need to communicate via phone, email, ...
 - In a computer system running simultaneous programs, processes need such communication as well
 - The kernel provides a variety of IPC forms
- Networking
 - A network allows processes on different computers, even on different operating systems for data exchange
 - Network access is a kernel service

Our Goal

- Exploring the following questions:
 - What are the details of each type of service?
 - How does data get from a device to the program and back?
 - ...
- We want to understand:
 - What the kernel is doing
 - How the kernel does it
 - How to write programs that use those services

Our Method

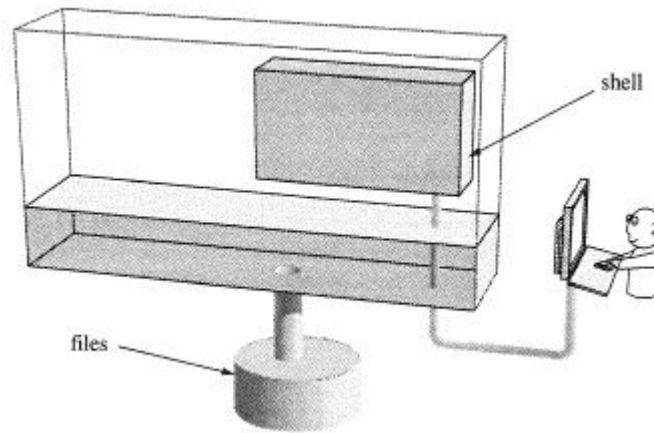
- Three simple steps to learn about system services by
 - 1. Looking at “real” programs
 - We will study standard Unix/Linux programs to see what they do and how they are used in practice
 - We will see what system services these programs use
 - 2. Looking at the system calls
 - We will learn about the system calls we use to invoke these services
 - 3. Writing our own version
 - Once we understand the program, what system services it uses, and how those services are used, we are ready to develop own system programs
 - Our own programs will be extensions of existing programs or ones that use similar principles

Our Method (cont.)

- Three questions that will be asked throughout our course
 - 1. What does that (kernel service) do?
 - 2. How does that work?
 - 3. Can I try to do it?

Unix from the User Perspective

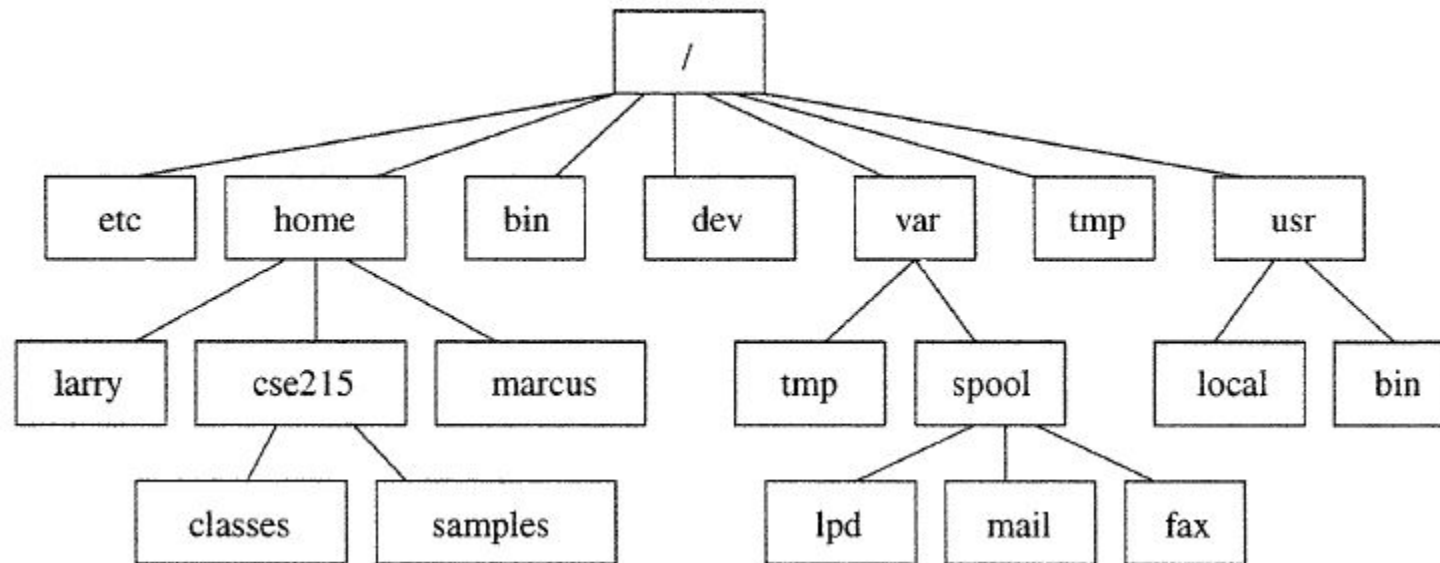
- Log in - Run programs - Log out
 - You login, run some programs, and then log out
 - The computer is the box at the left part of the picture
 - Inside the computer is memory that holds the kernel and user processes
 - The shell prints a prompt to tell the user it is ready to run a program



< The kernel manages all connections >

Unix from the User Perspective (cont.)

- Working with directories
 - A tree of directories



< Part of the directory tree >

Unix and Linux

- Unix is an operating system first developed by AT&T Bell Labs in the late 60s.
 - Primary contributors: Ken Thompson (B lang) and Dennis Ritchie (C Lang)
 - Written in C;
 - Commonly used in internet servers, workstations, etc
 - e.g., OS X, Solaris, BSD (Berkeley Software Distribution)
- Linux:
 - Developed by Linus Torvalds (responsible for kernel), combining GNU (GNU's Not Unix) tools from Richard Stallman
 - A “Unix-like” operating system: user-level interface very similar to Unix
 - Code base is different from original Unix code
 - Several distributions: e.g., Ubuntu, Fedora, SUSE, RedHat, ...

Unix and Linux (cont.)

	Unix	Linux
Cost	Different cost structures according to vendors	Freely distributed, but licensed versions
Development & Distribution	Mostly AT&T	Developed by Open Source Community
Manufacturer	Solaris (Oracle), AIX (IBM), HP-UX (HP), OS X (Apple)	Developed by the community, but Linus Torvalds oversees
Usage	Internet Servers, Workstations, and Finance IT Infra with 24x7 availability	Cell phones, tablet PCs, video game consoles, mainframes/supercomputers.
Security	85 -120 viruses	60-100 viruses
	See how amazing these are compared with that of Windows!	
Inception	Developed by AT&T employees at Bell Labs and Dennis Ritchie in 1969. Written in "C" language. Designed to be a portable, multi-tasking and multi-user system in a time-sharing configuration.	Inspired by MINIX (a Unix-like system), later added many features like GUI, and then Linus Torvalds developed it in 1992.

Summary

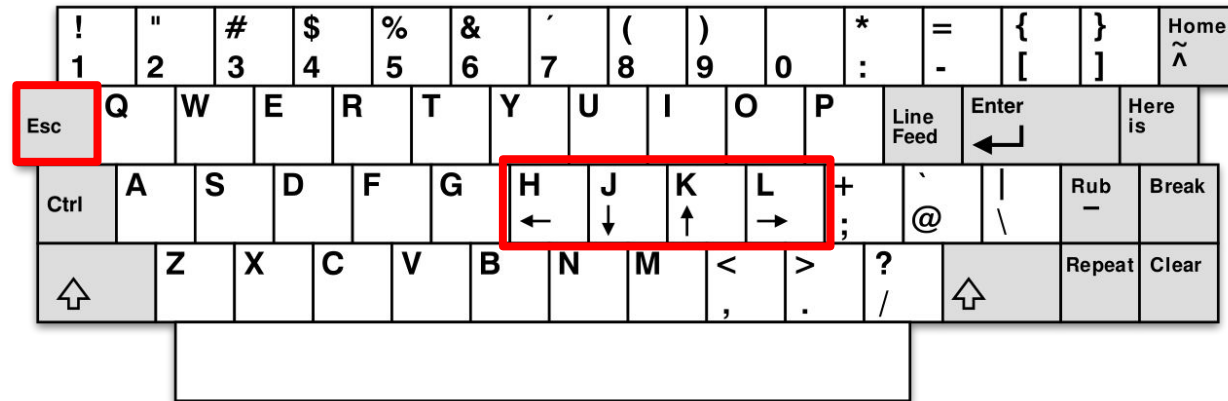
- Chapter 1 introduces the topic and approach of the book
- The purpose of the book is to explain how Unix works and to demonstrate how to write programs that use Unix system services
- The chapter explains that the Unix kernel is a program that manages multiple programs and system resources and connects programs to resources
- The method used to explain these ideas is to study and write common Unix programs

Appendix

`vi` Editor Tutorial

About vi

- vi (pronounced as distinct letters, /ˌviːˈaɪ/)
 - A screen-oriented text editor
 - The standard Unix editor
 - It is short for “visual”
 - There are lots of clones (vim, Elvis, nvi, etc.)



< ADM-3A terminal keyboard layout >



< ADM-3A terminal with keypad >

Entering the `vi` Editor

Command	Description
<code>vi filename</code>	Edits an existing <i>filename</i> or creates a new one if <i>filename</i> does not exist.
<code>vi -r filename</code>	Edits <i>filename</i> using <i>.filename.swp</i> as the input . The file <i>.filename.swp</i> must be deleted, using: <code>rm .filename.swp</code> , after use.
<code>vi -R filename</code>	Edits <i>filename</i> in read-only mode does not allow changes to <i>filename</i> .
<code>vi +/pattern filename</code>	Edits <i>filename</i> at the line containing the <i>pattern</i> .

File Interaction in the `vi` Editor

Command	Description
<code>:w</code>	Write to the file that you are editing without exiting the editor. Good for intermediate saves.
<code>:w def</code>	Write to a new file <i>def</i> , but stay in the existing file.
<code>:w! def</code>	Write to an existing file <i>def</i> , but stay in the existing file.
<code>:r def</code>	add (read) file <i>def</i> into the current file
<code>:f</code>	Provides information on the file currently editing, including the current line number, if it has been modified etc. <code><Ctrl+g></code>

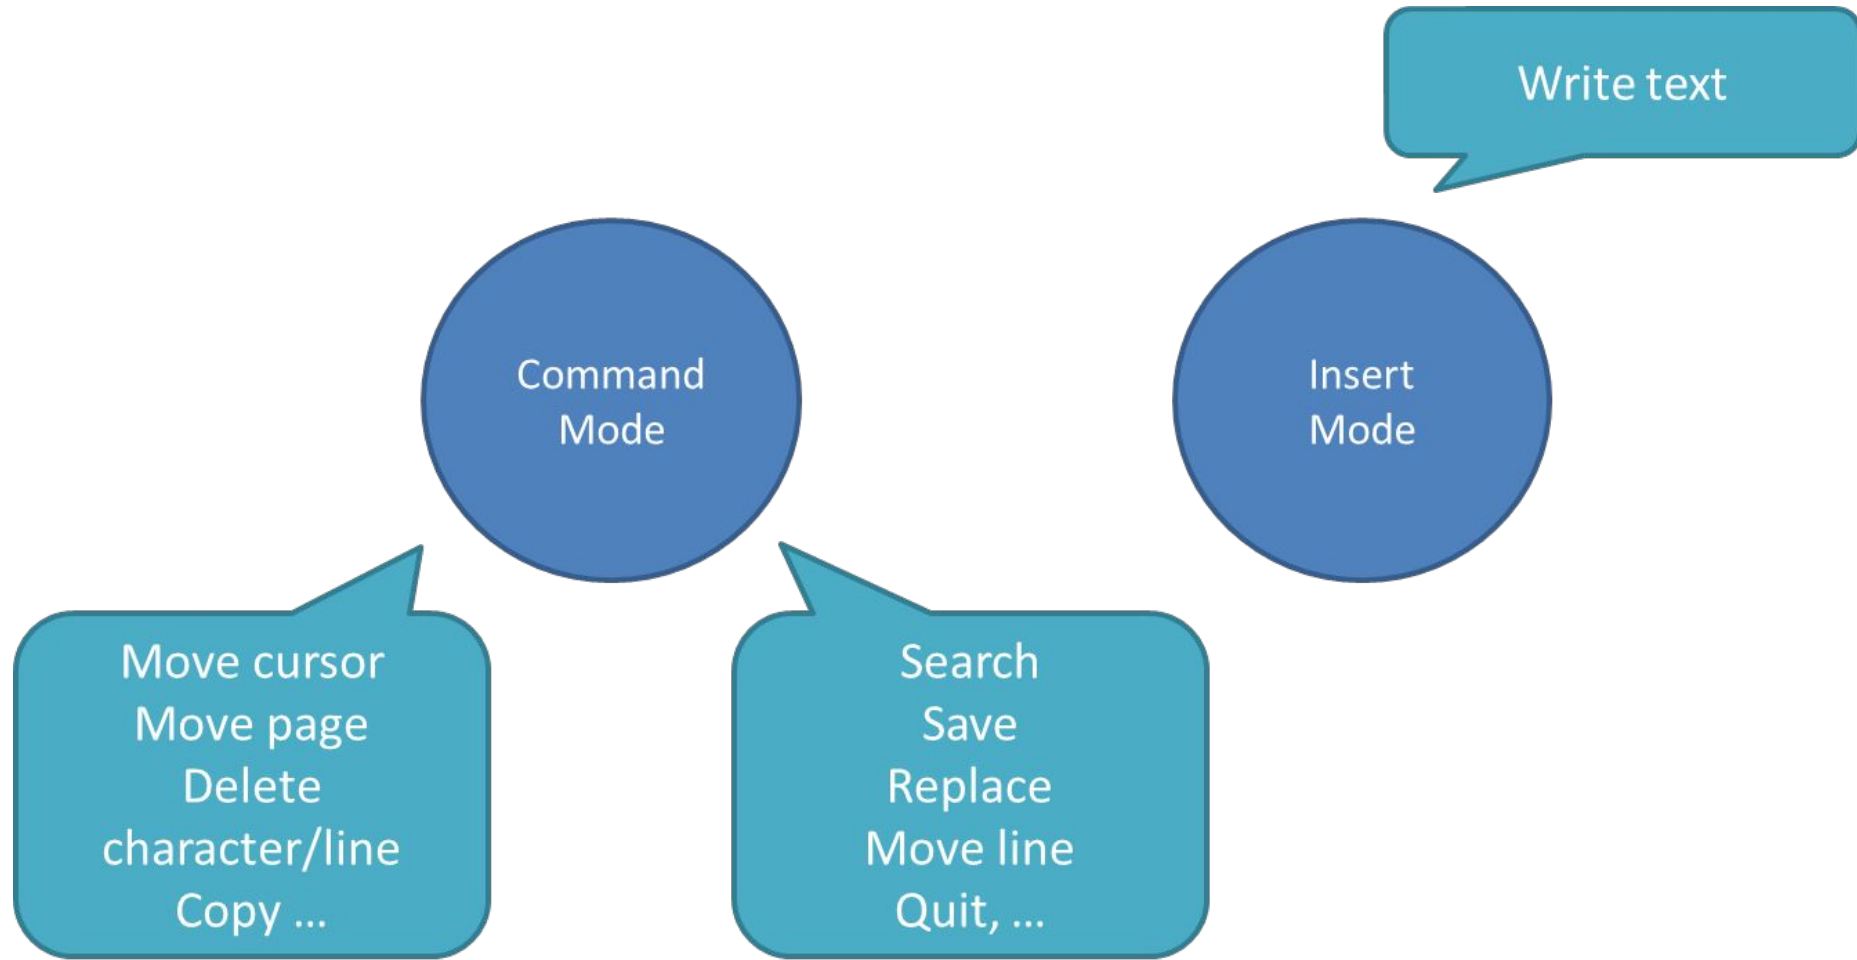
Exiting the `vi` Editor

Command	Description
<code>:q</code>	Exit (quit) without changes. Assumes no changes have been made.
<code>:q!</code>	Exit (quit) without changes. Assumes changes have been made.
<code>ZZ</code> or <code>:x</code>	Save the file only if changes have been made, then exit the editor
<code>:wq</code>	Save the file and exit the editor. This command will save the file even if no changes are made.

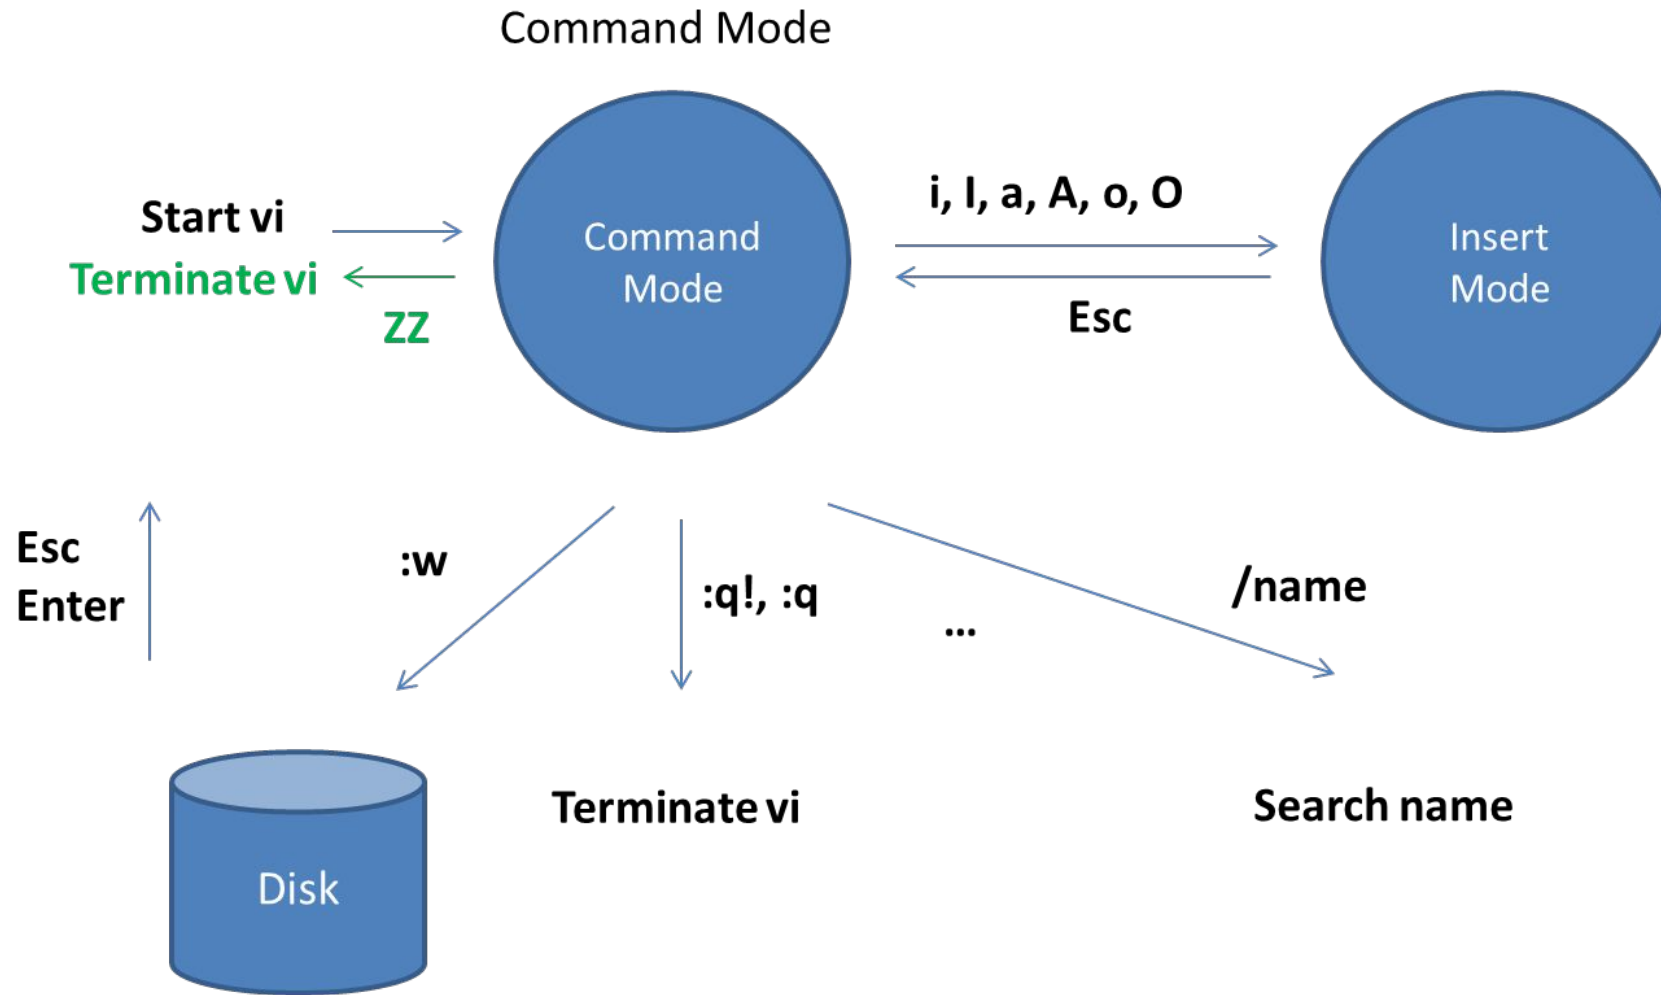
vi Modes

- Command mode:
 - The mode you are in when you start (default mode)
 - The mode in which commands are given to move around in the file, to make changes, and to leave the file
- Insert (or Text) mode
 - The mode in which text is created
 - There is more than one way to get into insert mode but only one way to leave: return to command mode by pressing <ESC>

vi Modes (cont.)



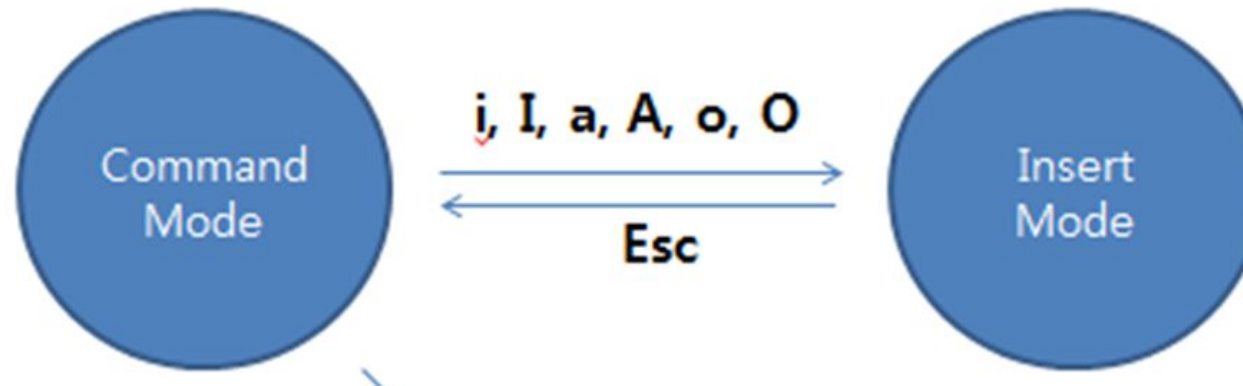
vi Modes (cont.)



Case-sensitive!!

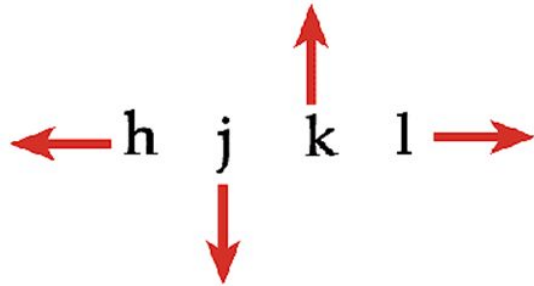
From Command Mode to Insert Mode

Key	Operation
i	Insert text before current character
a	Append text after current character
I	Begin text insertion at the beginning of a line
A	Append text at end of a line
o	Open a new line below current line
O	Open a new line above current line



Command Mode

- Basic cursor movement
 - `h`, `j`, `k`, `l`
 - Arrow key
 - Only support in `vim`



Key	Operation
<code>k</code>	Up one line
<code>j</code>	Down one line
<code>l</code>	Right one character (or use <code><Spacebar></code>)
<code>h</code>	Left one character
<code>w</code>	Right one word
<code>b</code>	Left one word
<code>W</code>	forward one word
<code>B</code>	back one word
<code>gg</code>	go to the first line
<code>G</code>	go to the last line

Command Mode (cont.)

- General command syntax

- The general format for commands is

`ncm`

- Where:

- **`n`** is an optional multiplier value
- **`c`** is the command
- **`m`** is an optional scale modifier

- Examples:

- `3dw` – delete 3 special character delimited words
- `dW` – delete a single, space, delimited word

Command Mode (cont.)

- Delete commands

d	Delete this line
w	Delete special character delimited word
W	Delete whitespace delimited word
}	Delete to next paragraph
^	Delete to beginning of the line

- Delete character commands

n x	Delete current [and n-1] character[s]
n X	Delete previous n character[s]


Command Mode (cont.)

- Splitting and joining lines
 - If you wish to split a line between words, position yourself on a space between the words to be split, then use the replace (r) followed by the return key. This will replace the space with a carriage return forming a new line.
 - If you wish to join or combine two lines go to the upper of the two lines to be joined, then key in the uppercase “J” for join.

- Example:


- We wish to split the line below between the word “should” and “be”.

“This line should be split.”

- 
- 1. Position the cursor here
 - 2. Key in “r” followed by “enter”

- To join the lines:

“This line should
be split.”

- 
- 1. Position the cursor anywhere on this line
 - 2. Key in “J”

Command Mode (cont.)

- Search commands

<i>/text</i>	Search for text forward (wraps to beginning)
<i>?text</i>	Search for text backward
n	Search for same text again
N	Search for text reverse direction of initial search

Command Mode (cont.)

- Copy/Paste commands

[n]Y or	Copy (yank) n lines
[n]ym	Copy (yank) a portion of the file determined by the measurement (e.g. 2yW will yank 2 space separated words.
P,p	Paste yanked or deleted data before (P), after (p) cursor position

- Note: to cut and paste, use the **d**m commands to delete and the p or P to paste

- Visual block command

v	Start Visual mode per character
V	Start Visual mode linewise

Command Mode (cont.)

- Some “nice to know” features
 - To turn line numbers on and off use:
 - `:set nu` and `:set nu!`
 - To execute a single Unix command from inside the editor use:
 - `:!cmd`
 - To go to *linenum*
 - `:linenum`
 - To temporarily return to the shell use:
 - `:sh` (type “exit” to return to editor)
 - To repeat a colon command or to go back to earlier colon commands use:
 - `:↑` or `:↓` (up arrow or down arrow)
 - To show the name of the current file use:
 - `:f`
- Resource: <https://www.fprintf.net/vimCheatSheet.html>