

System Programming (ELEC462)

Directories and File Properties

Dukyun Nam
HPC Lab@KNU

Contents

- Introduction
- Project: `ls` command
- Brief Review of the File System Tree
- Converting File Mode to a String
- The Three Special Bits
- Setting and Modifying the Properties of a File
- Summary
- Appendix: `gdb` Debugger

Introduction

- Recall
 - Last time, we looked at ways to operate on the CONTENTS of a file:


`open, read, write, lseek, close`

- But, there is more to a file than just contents:

- 1) Properties
 - Size, owner, permission, type, etc.
- 2) Location in a directory tree

- To learn about these:

- We shall write a version of `ls`



< Data flow in the `who` command >

What Does `ls` Do?

- Type the command `ls` to see what it does:

```
$ ls
Makefile  docs      ls2.c     s.tar     statdemo.c  tail1.c
chap03    ls1.c     old_src   stat1.c   tail1
$
```

- The default action of '`ls`':
 - To list the names of the files in the “current” directory

What Does `ls` Do? (cont.)

- When given the `'-l'` command-line option
 - `ls` present information about each file and its properties using the *long format*

```
$ ls -l
total 108
-rw-rw-r-- 2 bruce users 345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce users 27521 Aug 1 12:14 chap03
drwxrwxr-x 2 bruce users 1024 Aug 1 12:15 docs
-rw-r--r-- 1 bruce users 723 Feb 9 1998 ls1.c
-rw-r--r-- 1 bruce users 3045 Feb 15 03:51 ls2.c
drwxrwxr-x 2 bruce users 1024 Aug 1 12:14 old_src
-rw-rw-r-- 1 bruce users 30720 Aug 1 12:05 s.tar
-rw-r--r-- 1 bruce support 946 Feb 18 17:15 stat1.c
-rw-r--r-- 1 bruce support 191 Feb 9 1998 statdemo.c
-rwxrwxr-x 1 bruce users 37351 Aug 1 12:13 tail1
-rw-r--r-- 1 bruce users 1416 Aug 1 12:05 tail1.c
$
```

What Does `ls` Do? (cont.)

- Listing other directories

Asking <code>ls</code> about Other Directories and Their Files	
Example	Action
<code>ls /tmp</code>	list names of files in <code>/tmp</code> directory
<code>ls -l docs</code>	show attributes of files in <code>docs</code> directory
<code>ls -l ../Makefile</code>	show attributes of <code>../Makefile</code>
<code>ls *.c</code>	list names of files matching pattern <code>*.c</code>

- If an argument is a “directory,” `ls` lists its contents
- If an argument is a “file,” `ls` lists its name or attributes
 - What command line option is specified affects what `ls` does and what it shows


What Does `ls` Do? (cont.)

- Popular command-line options

Command	Action
<code>ls -a</code>	shows “.”-files
<code>ls -lu</code>	shows last-read time
<code>ls -s</code>	shows size in blocks
<code>ls -t</code>	sorts in time order
<code>ls -F</code>	shows file types

- Use `man` to see more options for `ls`
- Why don't we run `ls` with different options? What can you see?

Brief Review of the File System Tree



- Listing other directories
 - A disk is organized as a tree of directories, each of which contains files or directories
 - In Unix/Linux, every file on the system is located somewhere in a single tree of directories
 - This simplicity simplifies writing `ls`: only thinking about directories and files

How Does `ls` Work?


- What is a directory?

```
open directory
+--> read entry      -end of dir?--
|__ display file info
close directory  <-----+
```

- A directory is *a kind of file* that contains a list of names of files and directories
- Unlike a regular file, a directory is never empty
- Every directory contains two specific items
 - *dot* (.) is the name of the current directory
 - *dotdot* (..) is the name of the directory one level up

How Does `ls` Work? (cont.)

- Reading the contents of a directory



< Reading entries from a directory >

How Does `ls` Work? (cont.)

- The `dirent` structure
 - Typically defined `/usr/include/dirent.h`

```
struct dirent {
    ino_t      d_ino;      /* Inode number */
    off_t      d_off;      /* Not an offset; see below */
    unsigned short d_reclen; /* Length of this record */
    unsigned char d_type;   /* Type of file; not supported
                           by all filesystem types */
    char        d_name[256]; /* Null-terminated filename */
};
```

- An inode (index node) serves as a unique identifier for a specific piece of **metadata** on a given filesystem
- **d_name** contains the null-terminated filename

Can I Write ls? (ls1.c)

- Logic for listing a directory:

```
main()
    opendir
    while ( readdir )
        print d_name
    closedir
```

```
/** ls1.c
**  purpose  list contents of directory or directories
**  action   if no args, use .  else list files in args
**/
#include      <stdio.h>
#include      <sys/types.h>
#include      <dirent.h>

void do_ls(char []);

int main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }

    return 0;
}
```

```
void do_ls( char dirname[] )
/*
 *   list files in directory called dirname
 */
{
    DIR          *dir_ptr;           /* the directory */
    struct dirent *direntp;          /* each entry   */

    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr, "ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            printf("%s\n", direntp->d_name );
        closedir(dir_ptr);
    }
}
```


Writing `ls -l`

- `ls -l` operation
 - Mode: type of file (-, d) + permissions (user, group, everyone (other))
 - Links: reference to a file (often used in Linux)
 - Owner
 - Group
 - Size
 - Last-modified
 - Name

```
$ ls -l
total 108
-rw-rw-r-- 2 bruce users 345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce users 27521 Aug 1 12:14 chap03
drwxrwxr-x 2 bruce users 1024 Aug 1 12:15 docs
-rw-r--r-- 1 bruce users 723 Feb 9 1998 ls1.c
-rw-r--r-- 1 bruce users 3045 Feb 15 03:51 ls2.c
drwxrwxr-x 2 bruce users 1024 Aug 1 12:14 old_src
-rw-rw-r-- 1 bruce users 30720 Aug 1 12:05 s.tar
-rw-r--r-- 1 bruce support 946 Feb 18 17:15 stat1.c
-rw-r--r-- 1 bruce support 191 Feb 9 1998 statdemo.c
```

Writing `ls -l` (cont.)

- Needs a system call, called “`stat`”:



< Reading file properties using `stat` >

- Obtains the information about a given file
- Interface: `stat(name, ptr)`

Writing `ls -l` (cont.)

- System call “`stat`”:

stat		
PUPOSE	Obtain information about a file	
INCLUDE	#include <sys/stat.h>	
USAGE	int result = stat(char *fname, struct stat *bufp)	
AGRS	fname	name of file
	bufp	pointer to buffer
RETURNS	-1	if error
	0	if success

st_mode	type and permissions
st_uid	ID of owner
st_gid	ID of group
st_size	number of bytes in file
st_nlink	number of links to file
st_mtime	last content-modified time
st_atime	last-accessed time
st_ctime	last properties-changed time

Writing ls -l (cont.)

- 'fileinfo.c': to show file properties via stat ()

```
/* statinfo.c - demonstrates using stat() to obtain
 *             file information.
 *             - some members are just numbers...
 */
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>

void show_stat_info(char *, struct stat *);

int main(int ac, char *av[])
{
    struct stat info;          /* buffer for file info */

    if (ac>1)
        if( stat(av[1], &info) != -1 ){
            show_stat_info( av[1], &info );
            return 0;
        }
        else
            perror(av[1]); /* report stat() errors */
    return 0;
}
```

```
void show_stat_info(char *fname, struct stat *buf)
/*
 * displays some info from stat in a name=value format
 */
{
    printf("  mode: %o\n", buf->st_mode);          /* type + mode */
    printf(" links: %ld\n", buf->st_nlink);        /* # links   */
    printf("  user: %d\n", buf->st_uid);          /* user id   */
    printf(" group: %d\n", buf->st_gid);          /* group id  */
    printf("  size: %ld\n", buf->st_size);        /* file size */
    printf("modtime: %ld\n", buf->st_mtime);      /* modified  */
    printf("  name: %s\n", fname );             /* filename  */
}
```



Writing `ls -l` (cont.)

- Results of `fileinfo`

```
dynam@DESKTOP-Q4IJB7:~/lab4$ ./fileinfo fileinfo.c
mode: 100644
links: 1
user: 1000
group: 1000
size: 1155
modtime: 1663595693
name: fileinfo.c
dynam@DESKTOP-Q4IJB7:~/lab4$ ls -l fileinfo.c
-rw-r--r-- 1 dynam dynam 1155 Sep 19 22:54 fileinfo.c
```

```
dynam@DESKTOP-Q4IJB7:~/lab4$ cat /etc/passwd | grep dynam
dynam:x:1000:1000:,,,:/home/dynam:/bin/bash
```

Converting File Mode to a String




- `st_mode` includes **field type** and **permission bits**
- We convert file mode into a 16-bit integer, called *subfield coding*
- To read subfields, we need to use a technique, called **masking**
 - Zeroing out digits (making zeros), only subfield show through
 - Set marks to translate `st_mode` into string displayed by `ls -l`

Converting File Mode to a String (cont.)

- Subfield coding

Examples of Subfield Coding	
617-495-4204	area, exchange, line
027-93-1111	social security number
128.103.33.100	IP numbers

- Masking



- By masking a bit, focus will be reduced; unnecessary information is hidden.

Converting File Mode to a String (cont.)

- Using *masking* to decode file type and permission bits

```
#define S_IFMT      0170000    /* type of file */
#define S_IFREG     0100000    /* regular */
#define S_IFDIR     0040000    /* directory */
#define S_IFBLK     0060000    /* block special */
#define S_IFCHR     0020000    /* character special */
#define S_IFIFO     0010000    /* fifo */
#define S_IFLNK     0120000    /* symbolic link */
#define S_IFSOCK    0140000    /* socket */
```

```
/*
 * File type macros
 */
#define S_ISFIFO(m) (( (m)&(0170000)) == (0010000))
#define S_ISDIR(m) (( (m)&(0170000)) == (0040000))
#define S_ISCHR(m) (( (m)&(0170000)) == (0020000))
#define S_ISBLK(m) (( (m)&(0170000)) == (0060000))
#define S_ISREG(m) (( (m)&(0170000)) == (0100000))
```

```
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );    /* default=no perms */

    if ( S_ISDIR(mode) ) str[0] = 'd';    /* directory? */
    if ( S_ISCHR(mode) ) str[0] = 'c';    /* char devices */
    if ( S_ISBLK(mode) ) str[0] = 'b';    /* block device */

    if ( mode & S_IRUSR ) str[1] = 'r';    /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';    /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';    /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```

Converting User/Group ID to Strings

- Users

- `/etc/passwd` is the list of users
 - But that does not always present a complete list of users
- A library function `getpwuid` provides access to the complete list of users.

- Groups

- `/etc/group` is the list of group
- A user can belong to more than one group
- Function `getgrgid` provides access to list of group

Putting It All Together: ls2.c

```
/* ls2.c
 *      purpose  list contents of directory or directories
 *      action   if no args, use .  else list files in args
 *      note     uses stat and pwd.h and grp.h
 *      BUG: try ls2 /tmp
 */
#include <stdio.h>
#include <sys/types.h>
#include <dirent.h>
#include <sys/stat.h>
#include <string.h>

void do_ls(char[]);
void dostat(char *);
void show_file_info( char *, struct stat *);
void mode_to_letters( int , char [] );
char *uid_to_name( uid_t );
char *gid_to_name( gid_t );

int main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", ++av );
            do_ls( *av );
        }

    return 0;
}
```

```
void do_ls( char dirname[] )
/*
 *      list files in directory called dirname
 */
{
    DIR          *dir_ptr;          /* the directory */
    struct dirent *direntp;         /* each entry   */

    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr, "ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            dostat( direntp->d_name );
        closedir(dir_ptr);
    }
}

void dostat( char *filename )
{
    struct stat info;

    if ( stat(filename, &info) == -1 )          /* cannot stat */
        perror( filename );                    /* say why   */
    else                                         /* else show info */
        show_file_info( filename, &info );
}
```

Putting It All Together: ls2.c (cont.)

```
void show_file_info( char *filename, struct stat *info_p )
/*
 * display the info about 'filename'. The info is stored in struct at *info_p
 */
{
    char    *uid_to_name(), *ctime(), *gid_to_name(), *filemode();
    void    mode_to_letters();
    char    modestr[11];

    mode_to_letters( info_p->st_mode, modestr );

    printf( "%s"      , modestr );
    printf( "%4d "    , (int) info_p->st_nlink);
    printf( "%-8s "   , uid_to_name(info_p->st_uid) );
    printf( "%-8s "   , gid_to_name(info_p->st_gid) );
    printf( "%8ld "   , (long)info_p->st_size);
    printf( "%.12s "  , 4+ctime(&info_p->st_mtime));
    printf( "%s\n"    , filename );
}
```


Putting It All Together: ls2.c (cont.)

```
/*
 * This function takes a mode value and a char array
 * and puts into the char array the file type and the
 * nine letters that correspond to the bits in mode.
 * NOTE: It does not code setuid, setgid, and sticky
 * codes
 */
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );          /* default=no perms */

    if ( S_ISDIR(mode) ) str[0] = 'd';  /* directory?      */
    if ( S_ISCHR(mode) ) str[0] = 'c';  /* char devices    */
    if ( S_ISBLK(mode) ) str[0] = 'b';  /* block device    */

    if ( mode & S_IRUSR ) str[1] = 'r';  /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';  /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';  /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```


Putting It All Together: ls2.c (cont.)

```
#include <pwd.h>

char *uid_to_name( uid_t uid )
/*
 * returns pointer to username associated with uid, uses getpw()
 */
{
    struct passwd *getpwuid(), *pw_ptr;
    static char numstr[10];

    if ( ( pw_ptr = getpwuid( uid ) ) == NULL ){
        sprintf(numstr,"%d", uid);
        return numstr;
    }
    else
        return pw_ptr->pw_name ;
}
```

```
#include <grp.h>

char *gid_to_name( gid_t gid )
/*
 * returns pointer to group number gid. used getgrgid(3)
 */
{
    struct group *getgrgid(), *grp_ptr;
    static char numstr[10];

    if ( ( grp_ptr = getgrgid(gid) ) == NULL ){
        sprintf(numstr,"%d", gid);
        return numstr;
    }
    else
        return grp_ptr->gr_name;
}
```

The Three Special Bits

- The `st_mode` member of the `stat` structure contains 16 bits



- The three ***special*** bits
 - The Set-User-ID bit
 - The Set-Group-ID bit
 - The Sticky bit

1. The Set-User-ID Bit

```
dynam@DESKTOP-Q4IJB7:~/lab4$ ls -asl /usr/bin/passwd  
68 -rwsr-xr-x 1 root root 68208 Mar 14 2022 /usr/bin/passwd
```

```
dynam@DESKTOP-Q4IJB7:~/lab4$ ls -asl /etc/passwd  
4 -rw-r--r-- 1 root root 2939 Sep 13 01:33 /etc/passwd
```

- That SUID bit tells the kernel to run the program as though it were being run by the owner of the program: let's type 'passwd.'
 - Note that this program belongs to "root" as owner and "root" as group.
- The name 'set user ID' comes from the fact that the bit causes the kernel to set the *effective user ID* to the user ID of the owner of the program
 - User `root` owns `/etc/passwd`, so a program running as root can modify the file

2. The Set-Group-ID Bit

```
dynam@DESKTOP-Q4IJB7:~/lab4$ ls -asl /usr/bin/wall  
36 -rwxr-sr-x 1 root tty 35048 Feb 7 2022 /usr/bin/wall
```

- The second special bit sets the *effective group ID* of a program
 - Effective group ID: a number assigned by the kernel to a group
 - If a program belongs to group g and the set group ID bit is set, then the program runs as though it were being run by a member of g
 - This bit grants the program the access rights of members of that group

3. The Sticky Bit

- For files
 - The sticky bit tells the kernel to keep the program on the “swap device” (somewhere on disk) even if nobody was using it right now
 - The program gets never fragmented on the swap device on the disk
 - If the bit is “unset,” then the program might be split into many small sections scattered across the disk
 - So, the loading time of the program may get longer than when the bit is unset
 - Typically, loading a program from the swap device is faster than that from the regular section of the disk
 - By having the bit set to the file associated with the program, we can load the program quicker than otherwise

3. The Sticky Bit (cont.)

```
dynam@DESKTOP-Q4IJB7:~/lab4$ chmod +t /tmp/sticky_bit_test
dynam@DESKTOP-Q4IJB7:~/lab4$ ls -l /tmp
total 32
drwx----- 3 root  root  4096 Sep 20 08:39 snap.lxd
drwxrwxr-t  2 dynam dynam 4096 Sep 22 03:27 sticky_bit_test
```

- For directories
 - Some directories are designed to hold “temporary” files
 - These scratch (temp) directories, notably `/tmp`, are publicly writable, allowing any user to create and delete any files there
 - Setting the sticky bit overrides the publicly writable attribute for a directory
 - To avoid deletion of a folder and its content by other users
 - Files in the directory with the sticky bit set may only be deleted by their owners
 - “No one can delete things in my directory!”

Setting and Modifying the Properties of a File

- 1) Type of a file
 - A file has a type
 - It can be a regular file(-), a directory(d), a device file(b), a socket(s), a symbolic link(l), or a named pipe(p) (connected processes by a pipe)
 - Establishing the “type” of a file
 - The type of the file is determined when the file is created
 - For example, the `creat` system call creates a regular file
 - Different system calls are used to create directories, devices, and the like
 - Changing the type of a file is **impossible** unless the file is recreated

Setting and Modifying the Properties of a File (cont.)

- 2) Permission bits and special bits

- Establishing the mode of a file
- Changing the mode of a file

```
fd = creat( "newfile", 0744 );
```

```
chmod( "/tmp/myfile", 04764 );  
and  
chmod( "/tmp/myfile", S_ISUID | S_IRWXU | S_IRGRP|S_IWGRP | S_IROTH );
```

chmod		
PURPOSE	Change permission and special bits for a file	
INCLUDE	#include <sys/types.h> #include <sys/stat.h>	
USAGE	int result = chmod(char *path, mode_t mode);	
ARGS	path	path to file
	mode	new value for mode
RETURNS	-1	if error
	0	if success

Setting and Modifying the Properties of a File (cont.)

- 3) Number of links to a file
 - Represents the number of times the file is referenced in directories (or somewhere in the filesystem)

Setting and Modifying the Properties of a File (cont.)

- 4) Changing the owner and group of a file
 - A program can modify the owner and group of a file by making the `chown` system call

chown		
PURPOSE	Change owner and or group ID of a file	
INCLUDE	#include <unistd.h>	
USAGE	int chown(char *path, uid_t owner, gid_t group)	
ARGS	path	path to file
	owner	user ID for file
RETURNS	group	group ID for file
	-1	if error
	0	if success

Setting and Modifying the Properties of a File (cont.)

- 5) Size of a file
 - The size of a file, directory, and named pipe represents the number of bytes stored
- 6) Modification and access time

utime		
PURPOSE	Change access and modification time for files	
INCLUDE	#include <sys/time.h> #include <utime.h>	
USAGE	#include <sys/types.h> int utime(char *path, struct utimbuf *newtimes)	
ARGS	path	path to file
	newtimes	pointer to a struct utimbuf see utime.h for details
RETURNS	-1	if error
	0	if success


Setting and Modifying the Properties of a File (cont.)

- 7) Name of a file
 - Changing the name of a file is the 'rename' system call

rename	
PURPOSE	Change name and/or move a file
INCLUDE	#include <stdio.h>
USAGE	int result = rename(char *old, char *new)
ARGS	old old name of file or directory new new pathname for file or directory
RETURNS	-1 if error 0 if success

Summary

- Chapter 3 continues the discussion about files, showing there is more to files than their contents.
- Files are stored in directories, and files have attributes.
- Chapter 3 shows how to read a directory, and explains how to read, modify, and understand file attributes.
- The project for the chapter is to understand and write a version of the `ls` command.



Appendix

`gdb` Debugger

What is gdb?

- “GNU Debugger”
- A debugger for several languages, including C and C++
- It allows you to inspect what the program is doing at a certain point during execution
- Errors like segmentation faults may be easier to find with the help of gdb
- Online manual
 - <https://sourceware.org/gdb/current/onlinedocs/gdb/>

gdb Debugger

- GDB: interactive debugger
 - Allows the user to run a program and interactively examine its execution.
 - Features include:
 - breakpoints (“run until control reaches here, then prompt user”)
 - stack backtrace (chain of calls leading to some point in the code)
 - examination of program variables
- Compile option for debug
 - You have to use “-g” option when you compile a code
 - The (g)cc -g flag tells (g)cc to generate and embed debug info.
- Running gdb
 - `$ gdb program`
 - Caution: argument(program) is a program NOT a source code

gdb Commands

- `run [arglist]`
 - Run a program with a new argument (abbreviated command: r)
- `break { [file:] function | line number }`
 - Stop at a line that a user specifies for break. (abbreviated command: b)
- `continue`
 - Continue to run a program (abbreviated command: c)
- `quit`
 - Quit gdb (abbreviated command: q)

gdb Commands (cont.)

- `step`
 - Execute a line; if the current line has a function call, then run into the function and execute each line of the function at a time (abbreviated command: `s`)
- `next`
 - Continue to the next source line; if that line has a function call, then execute the function without stopping (abbreviated command: `n`)
- `print expr`
 - Perform the '`expr`' expression and then print its value (outcome) (abbreviated: `p`).
- `display`
 - Enable automatic displaying of certain expressions each time GDB stops at a breakpoint or after a step
- `list [first, last]`
 - Print lines from `first` to `last`: 10 lines by default

gdb Hands-on with ls1

- Compile a program with `-g` option

```
$ gcc -g -o ls1 ls1.c
```

- Starting up gdb

```
$ gdb ls1
```

- Running the program

```
(gdb) run
```

- Setting breakpoints

```
(gdb) break ls1.c:37  
(gdb) b do_ls
```