

System Programming (ELEC462)

Connection Control

Dukyun Nam
HPC Lab@KNU

Contents

- Introduction
- Devices Are Just Like Files
- Devices Are Not Like Files
- Attributes of Disk Connections
- Attributes of Terminal Connections
- Programming Other Devices: `ioctl`
- Summary

Introduction

- Objectives
 - Similarities and differences between files and devices
 - How to use those ideas to manage connections to devices.
 - Attributes of connections
 - Race conditions and atomic operations
 - Controlling device drivers
 - Streams
- What to learn and play with
 - System Calls: `fcntl`, `ioctl`
 - Functions: `tcsetattr`, `tcgetattr`
 - Commands: `stty`, `write`

Devices Are Just Like Files

- Files:
 - Contain data
 - Have properties
 - Identified by names in directories
 - We can read and write bytes from and into a file
- Devices
 - Microphone, ear phone, speaker, keyboard, USB, ...
 - Attached to a computer
 - Indeed, *terminal*, with keyboard and display, behaves like a file
 - Keystrokes that we type can be read by a program as input data
 - Characters that a program writes to the terminal are displayed on the screen

Devices Are Just Like Files (cont.)

- Device I/O is like file I/O
 - Devices are just like files
 - You get data from them and send data to them
- In Linux (Unix)
 - Various kinds of devices, such as sound/ graphic cards, terminals, mice, and disks, are treated as the same sort of object
 - Every device is treated as a “file” with:
 - A filename, an inode number, an owner, a set of permission bits, and a last-modified time

Devices Have Filenames

- Every device attached to a Linux machine is represented by a filename
 - You communicate with the device with
`read()`, `write()`, `open()`, `close()`, `lseek()`
- `/dev` directory
 - Files that represent devices are in the `/dev` directory
 - A **device file** is an interface to a *device driver*
 - A **device driver** is a computer program that operates or controls *a particular type of device*

Devices Have Filenames (cont.)

- Types of devices

- Character device

- Device with which the driver communicates by sending and receiving data in unit of character
 - e.g., Terminals, serial ports, parallel ports, sound cards...

- Block device

- Device with which the driver communicates by sending and receiving data in unit of block.
 - e.g., HDD, SSD, ...

```
$ tty
/dev/pts/2
$ cp /etc/passwd /dev/pts/2
$ ls > /dev/pts/2
```

Devices Have Filenames (cont.)

- `tty`

```
dynam@DESKTOP-Q4IJB7:~/lab6$ ls -li /dev/pts/0  
3 crw--w---- 1 dynam tty 136, 0 Oct  4 18:03 /dev/pts/0
```

- Used to print the file name of the terminal connected to standard input
- “136”: an identifier representing the device driver (a program controlling a device, of a terminal)
- “0”: the specific terminal number: type “`ls -li /dev/pts/0`” to see a different minor number
 - inode, mode, link, owner, group, major # (device driver), minor # (device), Last modified, filename
 - `pts`: pseudo-terminal slave (cf. `ptmx`: pseudo-terminal master)
 - Directory(d), Symbolic Link(l), Character Device(c), Block Device(b)
- * `tty` (TeleTYpewriter): Refers to the original printing terminals manufactured by the Teletype Corporation
 - <https://en.wikipedia.org/wiki/Teleprinter>

Terminals Can be Treated Exactly Disk Files

- The `write` Command

- Before instant messaging, chat rooms, or KakaoTalk, Unix users chatted with friends at other terminals with the `write` command.

- Do `man 1 write`:

```
$ man 1 write
WRITE(1)          Linux Programmer's Manual          WRITE(1)

NAME
    write - send a message to another user

SYNOPSIS
    write user [ttyname]

DESCRIPTION
    Write allows you to communicate with other users by copy-
    ing lines from your terminal to theirs.
```

- Sample screen shots

```
dynam@ubuntu:~$ tty
/dev/pts/9
dynam@ubuntu:~$ write dynam /dev/pts/10
Hi
CSE
^Cdynam@ubuntu:~$ |
```

```
dynam@ubuntu:~$ tty
/dev/pts/10
dynam@ubuntu:~$
Message from dynam@ubuntu on pts/9 at 02:08 ...
Hi
CSE
EOF
```

Write the `write` program: `write0.c`

- BUFSIZ: 8192
 - May be subject to change depending on system

```
$ ./write0 /dev/pts/1
```

```
Hi
```

```
KNU
```

```
^C
```

```
/* write0.c
 *
 *      purpose: send messages to another terminal
 *      method: open the other terminal for output then
 *              copy from stdin to that terminal
 *      shows: a terminal is just a file supporting regular i/o
 *      usage: write0 ttyname
 */

#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

int main( int ac, char *av[] )
{
    int    fd;
    char    buf[BUFSIZ];

    /* check args */
    if ( ac != 2 ){
        fprintf(stderr, "usage: write0 ttyname\n");
        exit(1);
    }

    /* open devices */
    fd = open( av[1], O_WRONLY );
    if ( fd == -1 ){
        perror(av[1]); exit(1);
    }

    /* loop until EOF on input */
    while( fgets(buf, BUFSIZ, stdin) != NULL )
        if ( write(fd, buf, strlen(buf)) == -1 )
            break;

    close( fd );


    return 0;
}
```

Device Files and Inodes

- How do these device files work?
How does the Unix/Linux file system of inodes and data blocks support this idea of device files?

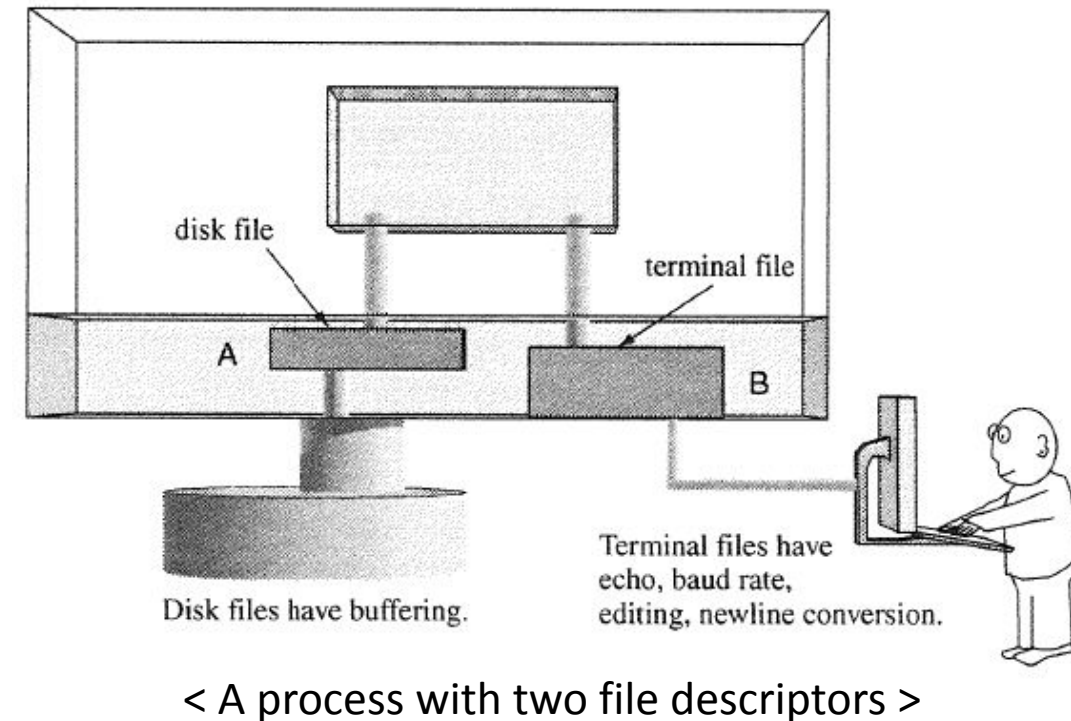
- Consider how `read` works

- The kernel finds the inode for the file descriptor
- The inode tells the kernel the type of the file
- If the file is a disk file
 - the kernel gets data by consulting the block allocation list
- If the file is a device file
 - the kernel reads data by calling the read part of the driver code for that device




Device Files and Inodes (cont.)

- Disk Files vs Device (here, Terminal) Files
 - Devices have their own attributes
 - e.g., Connection attributes available for the Terminal device
 - We can control the disk and connection attributes using data structures and system calls
 - baud rate: the rate at which information is transferred in a communication channel
 - “1024 baud”: “1K bits per sec.”



Attributes of Disk Connections

- The `open` system call
 - Creates a connection between a process and a disk file
 - That connection has several attributes
 - Attribute 1: Buffering
 - Attribute 2: Auto-append mode



< A processing unit in a data stream >


System Call: `fcntl`

- Control a file descriptor by reading and writing that descriptor (in integer)
- Perform operation `cmd` on the open file specified by `fd`

<code>fcntl</code>	
PURPOSE	Control file descriptors
INCLUDE	<code>#include <fcntl.h></code> <code>#include <unistd.h></code> <code>#include <sys/types.h></code>
USAGE	<code>int result = fcntl(int fd, int cmd);</code> <code>int result = fcntl(int fd, int cmd, long arg);</code> <code>int result = fcntl(int fd, int cmd, struct flock *lockp);</code>
ARGS	<code>fd</code> the file descriptor to control <code>cmd</code> the operation to perform <code>arg</code> arguments to the operation <code>lock</code> lock information
RETURNS	<code>-1</code> if error other depends on operation

Changing Settings – Disk File

- Example 1) Buffering can be set ON or OFF



To change driver settings:

1. Get settings,
2. modify them
3. send them back.

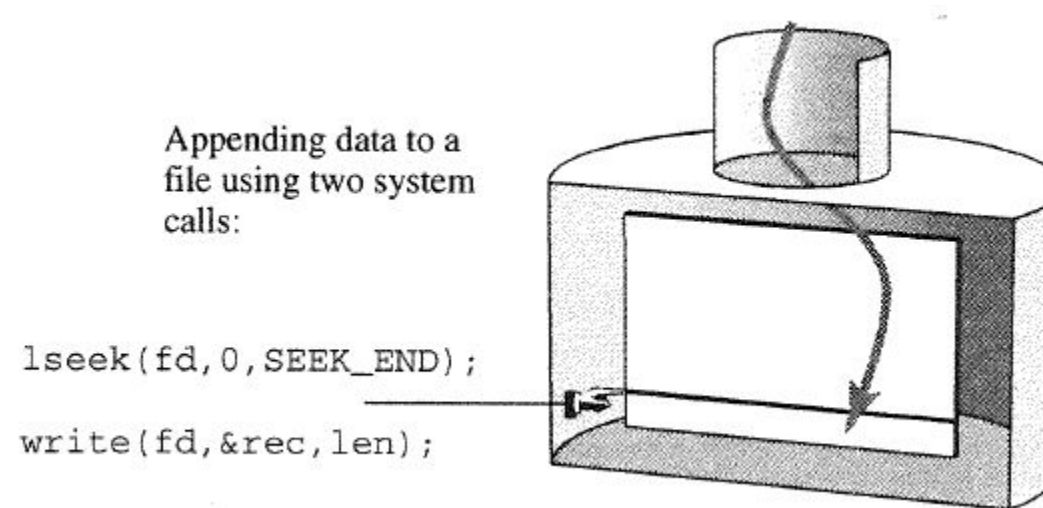
< Modifying the operation of a file descriptor >

- `s`: stores the information of control variables
- `O_SYNC`: set SYNC bit
 - Can be set to turn off kernel buffering
- `write()` s to disk will be synced to disk

```
#include <fcntl.h>
int s;
s = fcntl(fd, F_GETFL); // settings
s |= O_SYNC; // get flags
result = fcntl(fd, F_SETFL, s); // set SYNC bit
if ( result == -1 ) // set flags
    perror("setting SYNC"); // if error
// report
```

Changing Settings – Disk File (cont.)


- Example 2) auto-append mode is another `fd` attribute
 - Consider a logfile, like `wtmp`
 - Each login is recorded by appending a record to the file
 - How does one append to a file?



< Appending with `lseek` and `write` >

Changing Settings – Disk File (cont.)

- What if two people login at the “same” time, how can this *Race Condition* be avoided?



- Time 1: B's login process seeks to end of file
- Time 2: B's time slice up, A's login process seeks to find end of file
- Time 3: A's time slice is up, B's login process writes record
- Time 4: B's time slice is up, A's login process writes record

Changing Settings – Disk File (cont.)

- Solution 1: Set `O_APPEND` attribute in `fd`
 - When the `O_APPEND` bit is set for `fd`, each call to `write` automatically, includes an `lseek` to the end of file
 - The kernel combines `lseek` and `write` into an ***atomic operation***
 - The two are joined into one indivisible unit

```
#include <fcntl.h>

int s;                                // settings
s = fcntl(fd, F_GETFL);               // get flags
s |= O_APPEND;                        // set APPEND bit
result = fcntl(fd, F_SETFL, s);       // set flags
if ( result == -1 )                   // if error
    perror("setting APPEND");         // report
else
    write(fd, &rec, 1);               // write record at end
```

Changing Settings – Disk File (cont.)


- Solution 2: Include flag in `open ()`

```
$ fd = open(WTMP_FILE, O_WRONLY|O_APPEND|O_SYNC);
```

- Opens the `wtmp` file for writing with `O_APPEND` and `O_SYNC` bits on
 - `O_APPEND`: The kernel combines `lseek` and write into an atomic operation, which can NOT be divided
 - `O_SYNC`: The kernel that calls to write should return only when the bytes are written to the actual disk (rather than to the kernel buffer)
- Other flags
 - `O_CREAT`: Create the file if it does not exist.
 - `O_TRUNC`: If the file exists, truncate the file to length zero
 - `O_EXCL`: This flag is intended to prevent two processes from creating the same file. The second call returns -1

Disk Connections

- The kernel transfers data between disks and processes
 - A program can use the `open` and `fcntl` system calls to control the inner workings of these data transfers




< Connections to files have settings >

Attributes of Terminal Connections

- The `open` system call
 - Creates a connection between a process and a terminal
 - All connections to `ttys` have attributes of regular `fd`'s
 - All connections to `ttys` have additional attributes appropriate for human interface

Terminal I/O Not As Simple As It Appears

- A connection between a terminal and a process looks simple...
 - You can use `getchar` and `putchar` to transfer bytes between device and process



< The illusion of a simple, direct connection >

Example) A Simple Experiment

- The process receives no data until user presses Return
- User presses Return (ASCII 13), but processes sees newline (ASCII 10)
- Process sends newline, terminal receives Return-Newline pair

```
/* listchars.c
 *   purpose: list individually all the chars seen on input
 *   output: char and ascii code, one pair per line
 *   input: stdin, until the letter Q
 *   notes: useful to show that buffering/editing exists
 */

#include <stdio.h>

int main()
{
    int c, n = 0;


    while( ( c = getchar() ) != 'Q' )
        printf("char %3d is %c code %d\n", n++, c, c );

    return 0;
}
```

```
dynam@DESKTOP-Q4IJB7:~/lab6$ ./listchars
hello
char  0 is h code 104
char  1 is e code 101
char  2 is l code 108
char  3 is l code 108
char  4 is o code 111
char  5 is
code 10
```

Kernel Processes Terminal Data

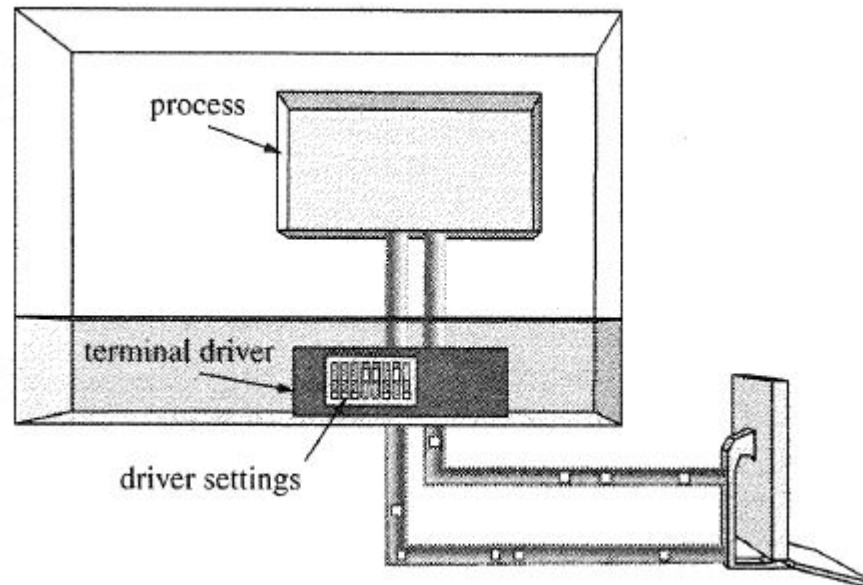
- There MUST be a ***processing layer*** somewhere in the middle of the file descriptor (of the Terminal device)



< Kernel processes terminal data >

The Terminal Driver

- The collection of kernel functions is called the *terminal driver*
 - Process data flowing between a process and the external device
- A connection between a terminal and a process
 - So, we may want to change the driver settings as we wish



< The terminal driver is part of the kernel >

Change Settings – Terminal

- The terminal (`tt_y`) driver is software in the kernel
 - It stands between the external device and the processes
- The driver contains lots of settings that control its operation
- The `stty` command allows you to examine and modify those settings

Change Settings – Terminal (cont.)

- The `stty` command
 - Lets users “read and change” settings in the driver for a terminal
 - `$ stty -a`
 - Some variables with values; e.g., rows, columns, ...
 - Some variables with boolean; `icrnl`, `-olcuc`, etc. with on or off
 - `icrnl` stands for *Input: convert Carriage Return to NewLine*
 - `-olcuc` means to disable the action for *Output: convert LowerCase to UpperCase*
- How to change driver settings? Some examples:
 - `stty erase X` # make ‘X’ the erase key
 - `stty -echo` # type invisibly; turn off keystroke
 - `stty erase @ echo` # change erase char to @ and turn on echo mode

Change Settings – Terminal (cont.)

- The tty driver contains dozens of operations
- These operations are grouped into four categories:
 - 1) Input: what the driver does with chars “coming from” the terminal
 - Includes converting lowercase to uppercase and carriage returns to newlines
 - 2) Output: what the driver does with chars “going to” the terminal
 - Includes replacing tab chars by sequences of spaces, converting newlines to carriage returns, and converting lowercase to uppercase, ...
 - 3) Control: how to represent characters: # of bits, parity, stop bits, etc.
 - Include *even parity*, *odd parity*, etc.
 - 4) Local: what the driver does while characters are inside the driver
 - Includes echoing keystrokes back to the user and buffering input until Return is pressed

Change Settings – Terminal (cont.)

- 3 steps for changing the terminal driver

```
#include <termios.h>
struct termios attribs;           // struct to hold attributes
tcgetattr(fd, &attribs);         // Step 1: Get "attribs" from driver
settings.c_lflag |= O_ECHO;      // Step 2: Modify "attribs"
                                //          Turn on ECHO bit in flagset
tcsetattr(fd, TCSANOW, &attribs); // Step 3: Send "attribs" back to driver
```


```
#include <termios.h>

struct termios settings;

tcgetattr(fd, &settings);

/* test, set, or
   clear bits */

tcsetattr(fd, how, &settings);
```



< Controlling the terminal driver with `tcgetattr` and `tcsetattr`

>

System Call: `tcgetattr()`

- Copies current settings from the terminal driver associated with the open file `fd` into the struct pointed to by `info`

<code>tcgetattr</code>		
PURPOSE	Read attributes from tty driver	
INCLUDE	#include <termios.h> #include <unistd.h>	
USAGE	int result = tcgetattr(int fd, struct termios *info);	
ARGS	fd	file descriptor connected to a terminal
	info	pointer to a struct termios
RETURNS	-1	if error
	0	if success

System Call: `tcsetattr()`

- Copies driver settings from the struct pointed to by `info` to the terminal driver associated to the open file `fd`

<code>tcsetattr</code>		
PURPOSE	Set attributes in tty driver	
INCLUDE	#include <termios.h> #include <unistd.h>	
USAGE	int result = tcsetattr(int fd, int when, struct termios *info);	
ARGS	fd	file descriptor connected to a terminal
	when	when to change the settings
	info	pointer to a struct termios
RETURNS	-1	if error
	0	if success

- TCSANOW: Update driver settings immediately
- TCSADRAIN: Wait until all output already queued in the driver has been transmitted to the terminal. Then update the driver
- TCSAFLUSH: Wait until all output already queued in the driver has been transmitted. Next, discard all queued input data. Then make the changes


Struct termios

- Contains several flagsets and an array of control characters

```
struct termios
{
    tcflag_t c_iflag;          /* input mode flags */
    tcflag_t c_oflag;          /* output mode flags */
    tcflag_t c_cflag;          /* control mode flags */
    tcflag_t c_lflag;          /* local mode flags */
    cc_t      c_cc[NCCS];      /* control characters */
    speed_t   c_ispeed;         /* input speed */
    speed_t   c_ospeed;         /* output speed */
};
```


Bits and chars in termios members

- The first four members are flagsets
 - When the current attributes are read from the driver into a struct termios, all the values in this struct are available
- `c_cc`: an array of control characters
 - Stores “keystrokes” that perform various control functions



Example: echostate.c

```
/* echostate.c
 * reports current state of echo bit in tty driver for fd 0
 * shows how to read attributes from driver and test a bit
 */

#include <stdio.h>
#include <termios.h>
#include <stdlib.h>

int main()
{
    struct termios info;
    int rv;

    rv = tcgetattr( 0, &info );    /* read values from driver */

    if ( rv == -1 ){
        perror( "tcgetattr" );
        exit(1);
    }
    if ( info.c_lflag & ECHO )
        printf(" echo is on , since its bit is 1\n");
    else
        printf(" echo is OFF, since its bit is 0\n");

    return 0;
}
```

```
$ gcc -o echostate echostate.c
$ ./echostate
    echo is on , since its bit is 1
$ stty -echo

$ stty echo
```

Example: setecho.c – change state

```
/* setecho.c
 *  usage:  setecho [y|n]
 *  shows:  how to read, change, reset tty attributes
 */

#include <stdio.h>
#include <termios.h>
#include <stdlib.h>

#define oops(s,x) { perror(s); exit(x); }

int main(int ac, char *av[])
{
    struct termios info;

    if ( ac == 1 )
        exit(0);

    if ( tcgetattr(0,&info) == -1 )           /* get attribs */
        oops("tcgetattr", 1);

    if ( av[1][0] == 'y' )
        info.c_lflag |= ECHO ;               /* turn on bit */
    else
        info.c_lflag &= ~ECHO ;              /* turn off bit */

    if ( tcsetattr(0,TCSANOW,&info) == -1 ) /* set attribs */
        oops("tcsetattr",2);

    return 0;
}
```

```
$ ./echostate; ./setecho n; ./echostate ; stty echo
echo is on , since its bit is 1
echo is OFF, since its bit is 0
$ stty -echo; ./echostate; ./setecho y; ./setecho n
echo is OFF, since its bit is 0
```

- Note that the driver and driver settings are stored in the kernel, NOT in the process.
 - That is, one process can change settings in the driver, and a different process can read or change the settings.

Example: showtty.c

- Display many driver attributes
 - For this example, we write part of the stty command
 - This one reads the settings and creates a report by testing bits and printing informative messages

```
/* showtty.c
 *      displays some current tty settings
 */

#include <stdio.h>
#include <termios.h>
#include <stdlib.h>

struct flaginfo {
    int    fl_value;
    char   *fl_name;
};
```

```
struct flaginfo input_flags[] = {

    IGNBRK ,    "Ignore break condition",
    BRKINT  ,    "Signal interrupt on break",
    IGNPAR  ,    "Ignore chars with parity errors",
    PARMRK  ,    "Mark parity errors",
    INPCK   ,    "Enable input parity check",
    ISTRIP  ,    "Strip character",
    INLCR   ,    "Map NL to CR on input",
    IGNCR   ,    "Ignore CR",
    ICRNL   ,    "Map CR to NL on input",
    IXON    ,    "Enable start/stop output control",
    /* _IXANY ,    "enable any char to restart output", */
    IXOFF   ,    "Enable start/stop input control",
    0       ,    NULL };
```

Example: showtty.c (cont.)

```
struct flaginfo local_flags[] = {
    ISIG      , "Enable signals",
    ICANON    , "Canonical input (erase and kill)",
    /* _XCASE  , "Canonical upper/lower appearance", */
    ECHO      , "Enable echo",
    ECHOE     , "Echo ERASE as BS-SPACE-BS",
    ECHOK     , "Echo KILL by starting new line",
    0         , NULL };
```

```
void show_flagset( int thevalue, struct flaginfo thebitnames[] )
/*
 * check each bit pattern and display descriptive title
 */
{
    int i;

    for ( i=0; thebitnames[i].fl_value ; i++ ) {
        printf( " %s is ", thebitnames[i].fl_name);
        if ( thevalue & thebitnames[i].fl_value )
            printf("ON\n");
        else
            printf("OFF\n");
    }
}
```

```
void show_some_flags( struct termios *ttyp )
/*
 * show the values of two of the flag sets: c_iflag and c_lflag
 * adding c_oflag and c_cflag is pretty routine - just add new
 * tables above and a bit more code below.
 */
{
    show_flagset( ttyp->c_iflag, input_flags );
    show_flagset( ttyp->c_lflag, local_flags );
}
```


Example: showtty.c (cont.)

```
void showbaud( int thespeed )
/*
 *   prints the speed in english
 */
{
    printf("the baud rate is ");
    switch ( thespeed ){
        case B300:    printf("300\n");    break;
        case B600:    printf("600\n");    break;
        case B1200:   printf("1200\n");   break;
        case B1800:   printf("1800\n");   break;
        case B2400:   printf("2400\n");   break;
        case B4800:   printf("4800\n");   break;
        case B9600:   printf("9600\n");   break;
        default:      printf("Fast\n");    break;
    }
}
```

```
int main()
{
    struct termios ttyinfo;    /* this struct holds tty info */

    if ( tcgetattr( 0 , &ttyinfo ) == -1 ){ /* get info */
        perror( "cannot get params about stdin");
        exit(1);
    }

    /* show info */
    showbaud ( cfgetospeed( &ttyinfo ) ); /* get + show baud rate */
    printf("The erase character is ascii %d, Ctrl-%c\n",
           ttyinfo.c_cc[VERASE], ttyinfo.c_cc[VERASE]-1+'A');
    printf("The line kill character is ascii %d, Ctrl-%c\n",
           ttyinfo.c_cc[VKILL], ttyinfo.c_cc[VKILL]-1+'A');

    show_some_flags( &ttyinfo ); /* show misc. flags */

    return 0;
}
```

System call: `ioctl`

- Provides access to the attributes and operations of the device driver connected to `fd`

<code>ioctl</code>		
PURPOSE	Control a Device	
INCLUDE	# include <sys/ioctl.h>	
USAGE	int result = ioctl (int fd, int operation [, arg..])	
ARGS	fd	file descriptor connected to device
	operation	operation to perform
	arg...	any args required for the operation
RETURNS	-1	if error
	other	depends on device

Programming Other Devices: `ioctl`

- Code for showing the settings of a video terminal screen having a size measured in rows and columns or in pixels:

```
#include <stdio.h>
#include <sys/ioctl.h>

void print_screen_dimensions(){
    struct winsize wbuf;


    if(ioctl(0, TIOCGWINSZ, &wbuf) != -1) {
        printf("%d rows x %d cols\n", wbuf.ws_row, wbuf.ws_col);
        printf("%d wide x %d tall\n", wbuf.ws_xpixel, wbuf.ws_ypixel);
    }
}

int main(int argc, char* argv[]){
    print_screen_dimensions();

    return 0;
}
```


Summary

- Connections to disk files is different than connection to device files in terms of data processing and transfer
 - Device driver : kernel code to manage connections to a device
 - `fcntl` and `ioctl` can be used to read and change settings in a device driver
- File descriptors, connections, and drivers
 - Make sure every device attached to a Linux system is controlled by its respective device driver



< File descriptors, connections, and drivers >