

System Programming (ELEC462)

Programming for Humans

Dukyun Nam
HPC Lab@KNU

Contents


- Introduction
- Software Tools vs Device-Specific Programs
- Modes of the Terminal Driver
- Writing a User Program
- Signals
- Prepared for Signals
- Processes are Mortal
- Programming for Devices
- Summary

Introduction

- Objectives
 - Ideas and skills
 - Software tools vs. user programs
 - Reading and changing settings of the terminal driver
 - Modes of the terminal driver
 - Nonblocking input
 - Timeouts on user input
 - Introduction to signals: How Ctrl-C works
 - System Calls
 - `fcntl`, `signal`

Software Tools

- Programs, that see no difference between disk files and devices
 - e.g., `who`, `ls`, `sort`, `uniq`, `grep`, `tr`, `du`, etc.
 - Reads bytes from *standard input*
 - Does some processing, and then
 - Writes a resulting stream of bytes to *standard output*, or sends error message, streams of bytes, to *standard error*



Fact: Most processes automatically have the first three file descriptors open. They do not need to call `open()` to make these connections.

< The three standard file descriptors >

Software Tools (cont.)

- These file descriptors could be connected to files, terminals, mice, printers, and pipes
- Such tools make no assumptions about sources and destinations of data the tools process
- Many of the programs also read from file names on the command line
 - Ex 1) `ls | sort`
 - Ex 2) `ls | uniq`
 - Ex 3) `tr "[:lower:]" "[:upper:]" < file1`

Software Tools (cont.)

- Read stdin or files; write to stdout and stderr
 - `$ sort > file2` # sort input and save to file2
[Ctrl-D] = EOF
 - `$ sort x > /dev/lp` # read x and send to printer
 - `$ ls | tr '[a-z]' '[A-Z]'` # translate `ls`'s output to uppercase

User-Program: A Common Type of Device-Specific Program

- Device-specific programs are written to interact with specific devices
 - e.g., scanner, camera, cd-rom, printer, terminal, ...
- In this chapter, we explore the ideas and techniques of writing a *device-specific programs*, by looking at the most common type of device-specific programs interacting with **terminals** designed to be used by human beings
 - We refer to the terminal-oriented programs as a *user-program*

User-Program (cont.)

- Examples of user programs:
 - `vi`, `emacs`, `more`, `lynx`, `hangman`, ...
- These programs can adjust setting in the terminal driver to control how the **keystrokes** are handled and output is processed
 - The driver has lots of settings
- Among the settings, common concerns of user program:
 - (a) immediate response to keys
 - (b) limited input set
 - (c) timeout on input
 - (d) resistance to Ctrl-C

Modes of the Terminal Driver

- To explore the terminal driver, let's experiment with a toy translation program: `rotate.c`

```
/* rotate.c : map a->b, b->c, .. z->a
 * purpose: useful for showing tty modes
 */

#include <stdio.h>
#include <ctype.h>

int main()
{
    int c;
    while ( ( c=getchar() ) != EOF ){
        if ( c == 'z' )
            c = 'a';
        else if (islower(c))
            c++;
        putchar(c);
    }


    return 0;
}
```

Modes of the Terminal Driver (cont.)

- Canonical (표준) mode: Buffering and Editing

- Run the program using the default settings:

```
$ cc rotate.c -o rotate
$ ./rotate
abx<-cd
bcde
efgCtrl-C
$
```




- Features revealed by this experiment:

< What you type and what the program gets >

- (a) 'x' is never seen by the program; backspace erases it
- (b) Chars appear on the screen as you type them, but
- (c) The program does not receive any input until you press the 'Enter' key
- (d) The Ctrl-C key discards input and stops the program

Modes of the Terminal Driver (cont.)

- Canonical mode: Buffering and Editing (Cont'd)
 - Buffering, echoing, editing, and control key processing are all done by the terminal driver
 - When buffering + editing enabled, the terminal connection is said to be in canonical mode



< Processing layers in the terminal driver >

Modes of the Terminal Driver (cont.)

- Noncanonical processing

Input: abx<-cd

- No buffering

- The command `stty -icanon` turns off *canonical mode* processing in the driver

```
$ stty -icanon ; ./rotate
abbcxy^?cdde
effggh
$ stty icanon
```

- Another experiment:

```
$ stty -icanon -echo ; ./rotate
bcy^?de
fgh
$ stty icanon echo (Note: You won't see this. Why?)
```

- Turn off canonical mode and also turn off echo mode
 - The driver no longer prints back the characters as we type them
 - Output comes only from the program

Summary of Terminal Modes

- Example input

Hello data DEL DEL DEL DEL world

- RAW mode (all processing off) output


Hello data DEL DEL DEL DEL world

- Canonical (cooked) mode output

Hello world

Summary of Terminal Modes (cont.)

- The terminal driver is a complex set of routines in the kernel
 - To understand the practical value of these modes, we develop a user program that uses various driver modes



< Major components of the terminal driver >

Writing a User Program: `play_again.c`

- Many user applications ask users *yes/no* questions
- The following shell script is the main loop for a bank machine:

```
#!/bin/sh
#
# atm.sh - a wrapper for two programs
#
while true
do
    do_a_transaction      # run a program
    if play_again         # run our program
    then
        continue         # if "y" loop back
    fi
    break                 # if "n" break
done
```

Writing a User Program: `play_again.c` (cont.)

- What does `play_again` do?
 - The logic of `play_again.c`:
 - Prompt a user with a question
 - Accept input
 - If “y”, return 0
 - If “n”, return 1

Writing a User Program: `play_again0`

- How: write a function with a loop
- But, a user needs to press RETURN key
 - ATMs don't require that

```
/* play_again0.c
 *   purpose: ask if user wants another transaction
 *   method: ask a question, wait for yes/no answer
 *   returns: 0=>yes, 1=>no
 *   better: eliminate need to press return
 */
#include <stdio.h>
#include <termios.h>

#define QUESTION "Do you want another transaction"

int get_response( char * );

int main()
{
    int response;

    response = get_response(QUESTION);    /* get some answer */
    return response;
}
```

```
int get_response(char *question)
/*
 * purpose: ask a question and wait for a y/n answer
 * method: use getchar and ignore non y/n answers
 * returns: 0=>yes, 1=>no
 */
{
    printf("%s (y/n)?", question);
    while(1){
        switch( getchar() ){
            case 'y':
            case 'Y': return 0;
            case 'n':
            case 'N':
            case EOF: return 1;
        }
    }
}
```

Writing a User Program: `play_again0` (cont.)

- Output

```
dynam@DESKTOP-Q4IJB7:~/lab7$ ./play_again0
Do you want another transaction (y/n)?sure
y
dynam@DESKTOP-Q4IJB7:~/lab7$ ./play_again0
Do you want another transaction (y/n)?no
dynam@DESKTOP-Q4IJB7:~/lab7$
```

- Two problems

- 1) The user has to press 'Enter' before the program can act on input
- 2) When the user presses 'Enter', the program receives and processes an entire line of data

- So, let's turn off canonical input to have the program act based on characters

Writing a User Program: `play_again1`

- Immediate response
 - Idea: process each char as typed
 - How: use `tcsetattr()` to
 - Turn off editing (`&= ~ICANON`)
 - Set input size to 1 char
 - Note: need to reset tty at the end of program
 - Choice: do we set ICANON on or just restore old settings?
 - But, responds to each char, usually with an error

Writing a User Program: play_again1 (cont.)

```
/* play_again1.c
 *   purpose: ask if user wants another transaction
 *   method: set tty into char-by-char mode, read char, return result
 *   returns: 0=>yes, 1=>no
 *   better: do no echo inappropriate input
 */
#include <stdio.h>
#include <termios.h>

#define QUESTION "Do you want another transaction"

int get_response(char * );
void set_crmode();
int tty_mode(int);

int main()
{
    int response;

    tty_mode(0); /* save tty mode */
    set_crmode(); /* set chr-by-chr mode */
    response = get_response(QUESTION); /* get some answer */
    tty_mode(1); /* restore tty mode */
    return response;
}
```

```
int get_response(char *question)
/*
 * purpose: ask a question and wait for a y/n answer
 * method: use getchar and complain about non y/n answers
 * returns: 0=>yes, 1=>no
 */
{
    int input;
    printf("%s (y/n)?", question);
    while(1){
        switch( input = getchar() ){
            case 'y':
            case 'Y': return 0;
            case 'n':
            case 'N':
            case EOF: return 1;
            default:
                printf("\ncannot understand %c, ", input);
                printf("Please type y or no\n");
        }
    }
}
```

Writing a User Program: `play_again1` (cont.)

```
void set_crmode()
/*
 * purpose: put file descriptor 0 (i.e. stdin) into chr-by-chr mode
 * method: use bits in termios
 */
{
    struct termios ttystate;

    tcgetattr( 0, &ttystate);          /* read curr. setting */
    ttystate.c_lflag      &= ~ICANON;  /* no buffering      */
    ttystate.c_cc[VMIN]    = 1;         /* get 1 char at a time */
    tcsetattr( 0, TCSANOW, &ttystate); /* install settings   */
}

/* how == 0 => save current mode,  how == 1 => restore mode */
int tty_mode(int how)
{
    static struct termios original_mode;
    if ( how == 0 )
        tcgetattr(0, &original_mode);
    else
        return tcsetattr(0, TCSANOW, &original_mode);
}
```

Writing a User Program: `play_again1` (cont.)

- Output
 - Type *sure* as a response:

```
$ make play_again1
cc      play_again1.c  -o play_again1
$ ./play_again1
Do you want another transaction (y/n)?s
cannot understand s, Please type y or no
u
cannot understand u, Please type y or no
r
cannot understand r, Please type y or no
e
cannot understand e, Please type y or no
y$
```

Writing a User Program: `play_again2`

- Ignore illegal chars
 - Idea: turn off echo
 - Simply ignore non y/n input
 - How: `&= ~ECHO`
 - No error messages
 - Program echos on legal input
 - Note: need to use `putchar()` to echo
 - But, what if a user wanders away without entering any key?

Writing a User Program: play_again2 (cont.)

```
/* play_again2.c
 *   purpose: ask if user wants another transaction
 *   method: set tty into char-by-char mode and no-echo mode
 *           read char, return result
 *   returns: 0=>yes, 1=>no
 *   better: timeout if user walks away
 */
#include <stdio.h>
#include <termios.h>

#define QUESTION      "Do you want another transaction"

int get_response(char * );
void set_cr_noecho_mode();
int tty_mode(int);

int main()
{
    int    response;

    tty_mode(0);                /* save mode */
    set_cr_noecho_mode();       /* set -icanon, -echo */
    response = get_response(QUESTION); /* get some answer */
    tty_mode(1);                /* restore tty state */
    return response;
}
```

```
int get_response(char *question)
/*
 * purpose: ask a question and wait for a y/n answer
 * method: use getchar and ignore non y/n answers
 * returns: 0=>yes, 1=>no
 */
{
    printf("%s (y/n)?", question);
    while(1){
        switch( getchar() ){
            case 'y':
            case 'Y': return 0;
            case 'n':
            case 'N':
            case EOF: return 1;
        }
    }
}
```


Writing a User Program: `play_again2` (cont.)

```
void set_cr_noecho_mode()
/*
 * purpose: put file descriptor 0 into chr-by-chr mode and noecho mode
 * method: use bits in termios
 */
{
    struct termios ttystate;

    tcgetattr( 0, &ttystate);          /* read curr. setting */
    ttystate.c_lflag      &= ~ICANON;  /* no buffering      */
    ttystate.c_lflag      &= ~ECHO;    /* no echo either    */
    ttystate.c_cc[VMIN]    = 1;        /* get 1 char at a time */
    tcsetattr( 0, TCSANOW, &ttystate); /* install settings   */
}

/* how == 0 => save current mode,  how == 1 => restore mode */
int tty_mode(int how)
{
    static struct termios original_mode;
    if ( how == 0 )
        tcgetattr(0, &original_mode);
    else
        return tcsetattr(0, TCSANOW, &original_mode);
}
```

Writing a User Program: `play_again3`

- Non-blocking mode
 - Blocking mode:
 - e.g., `getchar()` or `read()` → Wait for input
 - How to set non-blocking
 - Use `O_NDELAY` in `open()` or `fcntl()`
 - New `play_again`
 - Timeout feature
 - Telling the terminal driver NOT to wait
 - No input found then sleep for few (2) seconds and look again for input
 - After three attempts, then give up

Writing a User Program: `play_again3` (cont.)

- Idea: if no input
 - wait.. then ask again
 - ...
 - give up
- How: put the fd in 'non-blocking' mode
 - i.e., do not wait (block) for input
 - read return 0 if no chars available
- Note: `read()` usually waits for input
- Fact
 - Non-blocking input is an attribute of a file descriptor and may be set for any open file - disk files AND devices

Writing a User Program: play_again3 (cont.)

```
/* play_again3.c
 *   purpose: ask if user wants another transaction
 *   method: set tty into chr-by-chr, no-echo mode
 *           set tty into no-delay mode
 *           read char, return result
 *   returns: 0=>yes, 1=>no, 2=>timeout
 *   better: reset terminal mode on Interrupt
 */
#include <stdio.h>
#include <termios.h>
#include <fcntl.h>
#include <string.h>
#include <unistd.h>
#include <ctype.h>

#define ASK          "Do you want another transaction"
#define TRIES        3          /* max tries */
#define SLEEPTIME    2          /* time per try */
#define BEEP         putchar('\a') /* alert user */

int get_response( char *, int );
int get_ok_char();
void set_cr_noecho_mode();
void set_nodelay_mode();
void tty_mode(int );

int main()
{
    int    response;

    tty_mode(0);          /* save current mode */
    set_cr_noecho_mode(); /* set -icanon, -echo */
    set_nodelay_mode();   /* noinput => EOF */
    response = get_response(ASK, TRIES); /* get some answer */
    tty_mode(1);          /* restore orig mode */
    return response;
}
```

```
int get_response( char *question , int maxtries)
/*
 * purpose: ask a question and wait for a y/n answer or maxtries
 * method: use getchar and complain about non-y/n input
 * returns: 0=>yes, 1=>no, 2=>timeout
 */
{
    int    input;

    printf("%s (y/n)?", question); /* ask */
    fflush(stdout);                 /* force output */
    while ( 1 ){
        sleep(SLEEPTIME);          /* wait a bit */
        input = tolower(get_ok_char()); /* get next chr */
        if ( input == 'y' )
            return 0;
        if ( input == 'n' )
            return 1;
        if ( maxtries-- == 0 )      /* outatime? */
            return 2;              /* sayso */
        BEEP;
    }
}
```

Writing a User Program: `play_again3` (cont.)

```
void set_cr_noecho_mode()
/*
 * purpose: put file descriptor 0 into chr-by-chr mode and noecho mode
 * method: use bits in termios
 */
{
    struct termios ttystate;

    tcgetattr( 0, &ttystate);          /* read curr. setting */
    ttystate.c_lflag      &= ~ICANON;  /* no buffering      */
    ttystate.c_lflag      &= ~ECHO;    /* no echo either    */
    ttystate.c_cc[VMIN]    = 1;        /* get 1 char at a time */
    tcsetattr( 0, TCSANOW, &ttystate); /* install settings   */
}

void set_nodelay_mode()
/*
 * purpose: put file descriptor 0 into no-delay mode
 * method: use fcntl to set bits
 * notes: tcsetattr() will do something similar, but it is complicated
 */
{
    int termflags;

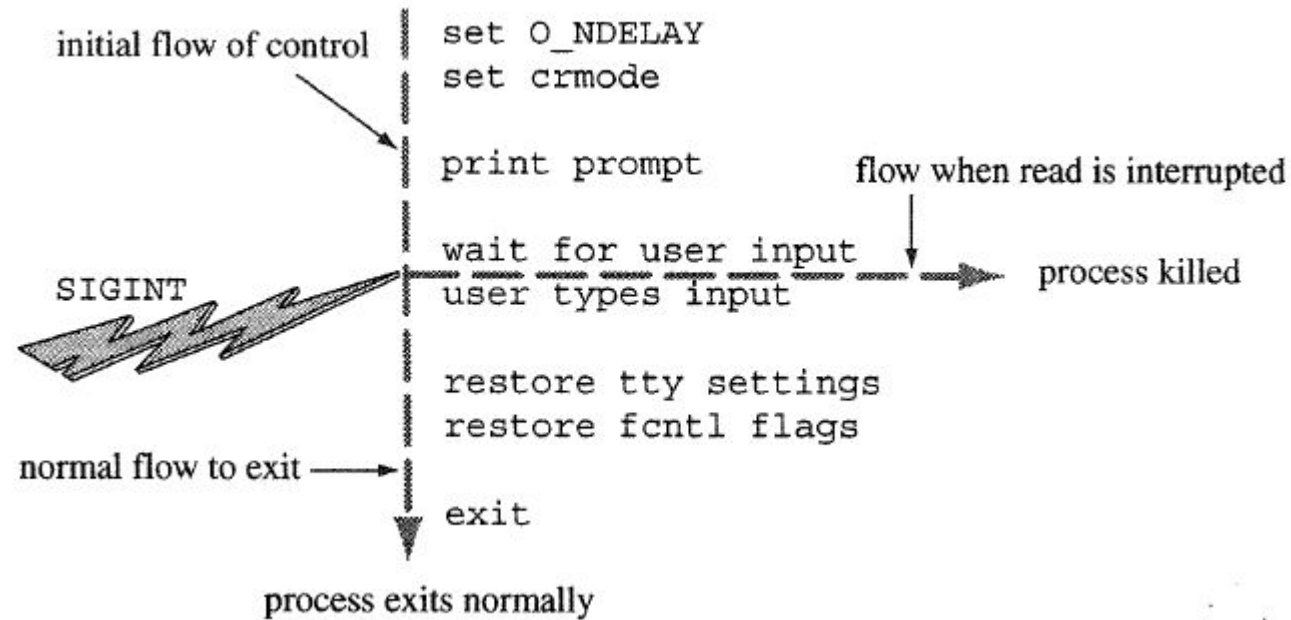
    termflags = fcntl(0, F_GETFL);      /* read curr. settings */
    termflags |= O_NDELAY;              /* flip on nodelay bit */
    fcntl(0, F_SETFL, termflags);      /* and install 'em     */
}
```

```
/* how == 0 => save current mode, how == 1 => restore mode */
/* this version handles termios and fcntl flags */

void tty_mode(int how)
{
    static struct termios original_mode;
    static int original_flags;
    if ( how == 0 ){
        tcgetattr(0, &original_mode);
        original_flags = fcntl(0, F_GETFL);
    }
    else {
        tcsetattr(0, TCSANOW, &original_mode);
        fcntl( 0, F_SETFL, original_flags);
    }
}
```

Signals

- Ctrl-C: Kills a running process



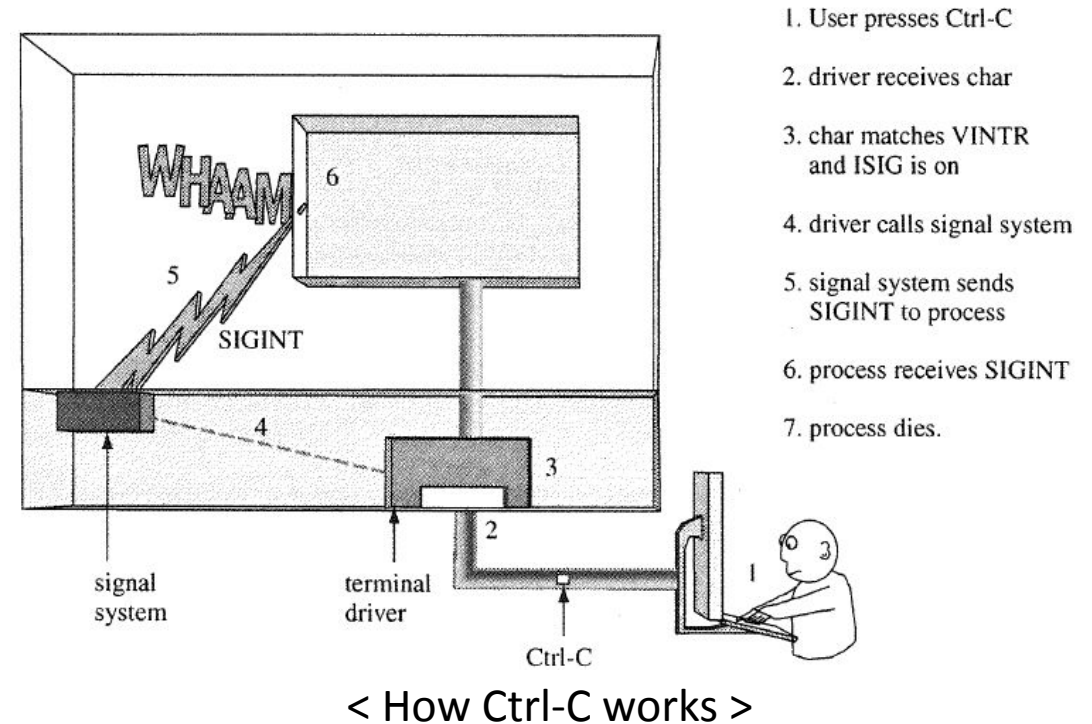
< Ctrl-C kills a program >

Signals (cont.)

- What does **Ctrl-C** do?
 - The Ctrl-C key *interrupts* a running program.
 - Generates a **signal**!
- Signal: defined as a **one-word message**
 - A kernel mechanism of showing how the kernel “controls” processes
 - Simple (by **numbers**) but very powerful and strong
 - e.g., go, stop, out, green light
- Each signal has its own numerical code.
 - Interrupt signal generated by Ctrl-C is No. 2: SIG**INT**

Signals (cont.)

- How does Ctrl-C do?
 - 'VINTR'
 - Interrupt character (INTR)
 - Send a SIGINT signal
 - 'ISIG'
 - When any of the characters INTR, QUIT, SUSP, or DSUSP are received, generate the corresponding signal
 - Refer to `termios`



Signals (cont.)

- Where do signals come from?

- Users


- Press Ctrl-C, Ctrl-\, ...

- Kernel

- A segmentation violation, a floating point exception, or an illegal opcode

- Other processes

- One process can send a signal to another process
 - e.g., `kill` system call



< Three sources of signals >

Signals (cont.)

- List of signals
 - “Typically” defined in `/usr/include/signal.h`

```
#define SIGHUP      1  /* hangup, generated when terminal disconnects */
#define SIGINT      2  /* interrupt, generated from terminal special char */
#define SIGQUIT     3  /* (*) quit, generated from terminal special char */
#define SIGILL      4  /* (*) illegal instruction (not reset when caught) */
#define SIGTRAP     5  /* (*) trace trap (not reset when caught) */
#define SIGABRT     6  /* (*) abort process */
#define SIGEMT      7  /* (*) EMT instruction */
#define SIGFPE      8  /* (*) floating point exception */
#define SIGKILL     9  /* kill (cannot be caught or ignored) */
#define SIGBUS     10  /* (*) bus error (specification exception) */
#define SIGSEGV    11  /* (*) segmentation violation */
#define SIGSYS     12  /* (*) bad argument to system call */
#define SIGPIPE    13  /* write on a pipe with no one to read it */
#define SIGALRM    14  /* alarm clock timeout */
#define SIGTERM    15  /* software termination signal */
```

Signals (cont.)

- When a signal comes
 - A process can tell the kernel, by using the `signal` system call, how it wants to respond to a signal
 - Process have 3 choices:
 - Accept the default action (usually death)
 - `signal(SIGINT, SIG_DFL) // reset signal to its default action`
 - Ignore
 - “Hey kernel, I want to ignore SIGINT”
 - `signal(SIGINT, SIG_IGN) // ignore the signal`
 - Call a function (say, `f`), called signal handler
 - `signal(SIGINT, f)`

Signals (cont.)

- How to call a signal handler
 - `signal(signum, function_name)`

signal	
PURPOSE	Simple signal handling
INCLUDE	<code>#include <signal.h></code>
USAGE	<code>result = signal (int signum, void (*action)(int))</code>
ARGS	signum the signal to respond to action how to respond
RETURNS	-1 if error prevaction if success

Signal Handling: sigdemo1

```
/* sigdemo1.c - shows how a signal handler works.
 *             - run this and press Ctrl-C a few times
 */

#include <stdio.h>
#include <signal.h>
#include <unistd.h>

void f(int);

int main()
{
    void f(int);           /* declare the handler */
    int i;

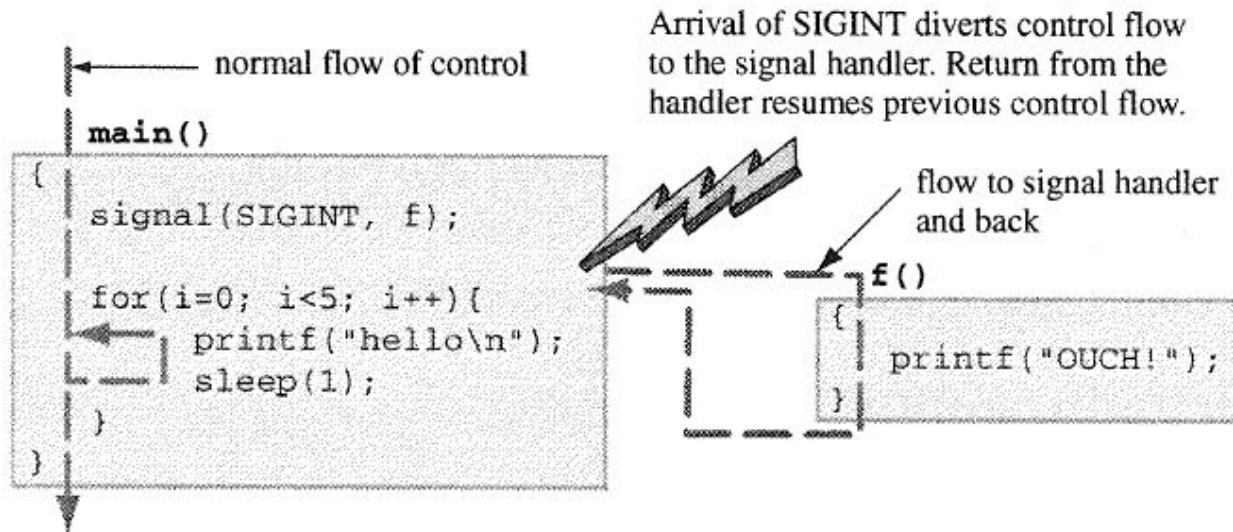
    signal( SIGINT, f );   /* install the handler */
    for(i=0; i<5; i++) {  /* do something else */
        printf("hello\n");
        sleep(1);
    }

    return 0;
}

void f(int signum)        /* this function is called */
{
    printf("OUCH!\n");
}
```

Signal Handling: sigdemo1 (cont.)

- How sigdemo1 works
 - There are two independent flows of control



```
$ ./sigdemo1
hello
hello    press Ctrl-C now
OUCH!
hello    press Ctrl-C now
OUCH!
hello
hello
$
```

< A signal causes a subroutine call >

Signal Handling: sigdemo2

- Ignoring a signal
 - Ctrl-\ : SIGQUIT


```
/* sigdemo2.c - shows how to ignore a signal
 *             - press Ctrl-\ to kill this one
 */

#include <stdio.h>
#include <signal.h>
#include <unistd.h>

int main()
{
    signal( SIGINT, SIG_IGN );

    printf("you can't stop me!\n");
    while( 1 )
    {
        sleep(1);
        printf("haha\n");
    }

    return 0;
}
```



< The effect of `signal (SIGINT, SIG_IGN)` >

```
$ ./sigdemo2
you can't stop me!
haha
haha
haha    press Ctrl-C now
haha    press Ctrl-C nowpress Ctrl-C now
haha
haha
haha    press ^\ now
Quit
$
```

Summary

- Software tools: read stdin or files / write to stdout
 - Software tools view input and output as byte streams
 - Most processes automatically have 3 fd's open
 - 0 (stdin), 1 (stdout), 2 (stderr)
 - The program does not need to call open for these
- User programs
 - Designed to be used by a human at a keyboard and screen
- A signal is a one-word message
 - e.g., green light, stop sign, umpire gesture, etc.

Summary (cont.)

- Device-specific programs have to control the connection to the device
 - Terminal is the most common and popular service
- A terminal driver has many settings
 - A collection of settings: a mode of the terminal driver
- Keys users press fall into the following three categories
 - Regular data: delivered through the driver
 - Editing functions (invoked by the keys)
 - e.g., erase key: removes the previous char from the line buffer and sends the codes to the terminal screen to remove that from the display
 - Process control functions: e.g., Ctrl-C key
- A signal is a short message from the kernel to a process
 - A process tells the kernel how to react upon receipt of a signal

Appendix

Make

* Most slides from CS352 by courtesy of Prof. Saumya Debray at U of Arizona and Prof. Suh at KNU

Structuring Large Applications

- So far, many of our programs have involved a single source file
 - Obviously impractical for large(r) programs
 - Even where practical, may not be good from a design perspective
- If an application is broken up into multiple files, we need to manage the build process:
 - How do we (re)compile the various different files that make up the application(s)?


Structuring Large Applications (cont.)

- When one file is edited, other files may need to be recompiled
 - Changes to typedefs or macros in header files
 - Changes to types of shared variables
- Applications can contain a lot of files
 - e.g., Linux kernel source code: about 5,000 files (totaling 15+M LOC)
- Re-compiling all files whenever any file is changed can be

VERY time-consuming


Structuring Large Applications (cont.)

- Obvious idea: **only** recompile those files that need to be recompiled



Structuring Large Applications (cont.)

- Obvious idea: **only** recompile those files that need to be recompiled




Structuring Large Applications (cont.)

- “Smart recompilation” : issues
 - Need to be able to express, or keep track of dependencies between files
 - “**Dependency**” \approx which files are (might be) affected by a change to a file
 - Need to make sure that all (and only) affected files are recompiled
 - Doing this manually is **tedious** (지루한) and **error-prone** (오류 발생이 쉬운)
 - WANT an **automated** solution!!!
- make
 - a tool to automate recompilation of parts of a project based on a file of dependencies (“*make file*”)

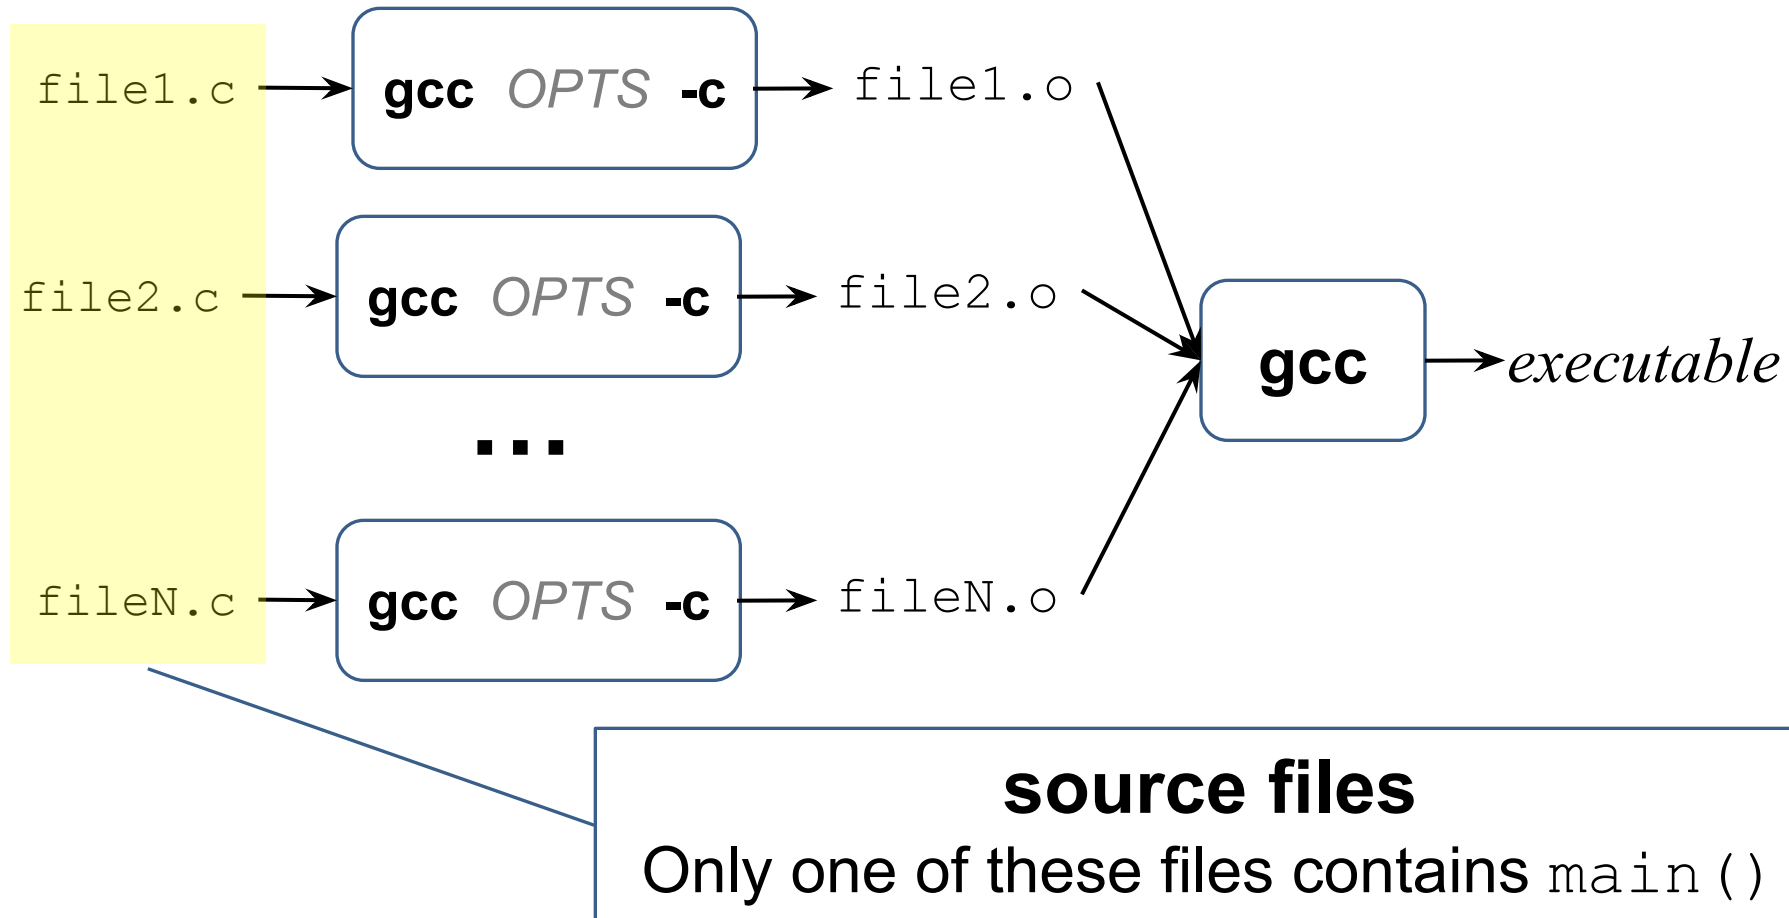
What is make?

- `make(1)` is a command generator and build utility
 - Using a description file (usually `Makefile`) it creates
- A sequence of commands for execution by the shell
 - used to sort out dependency relations among files
 - avoids having to rebuild the entire project after modification of a single source file
 - performs selective rebuilds following a dependency graph
 - allows simplification of rules through use of macros and suffixes, some of which are internally defined
 - different versions of `make(1)` (BSD `make`, GNU `make`, Sys V `make`, ...) may differ in e.g.:
 - variable assignment and expansion/substitution
 - including other files
 - flow control (for-loops, conditionals etc.)

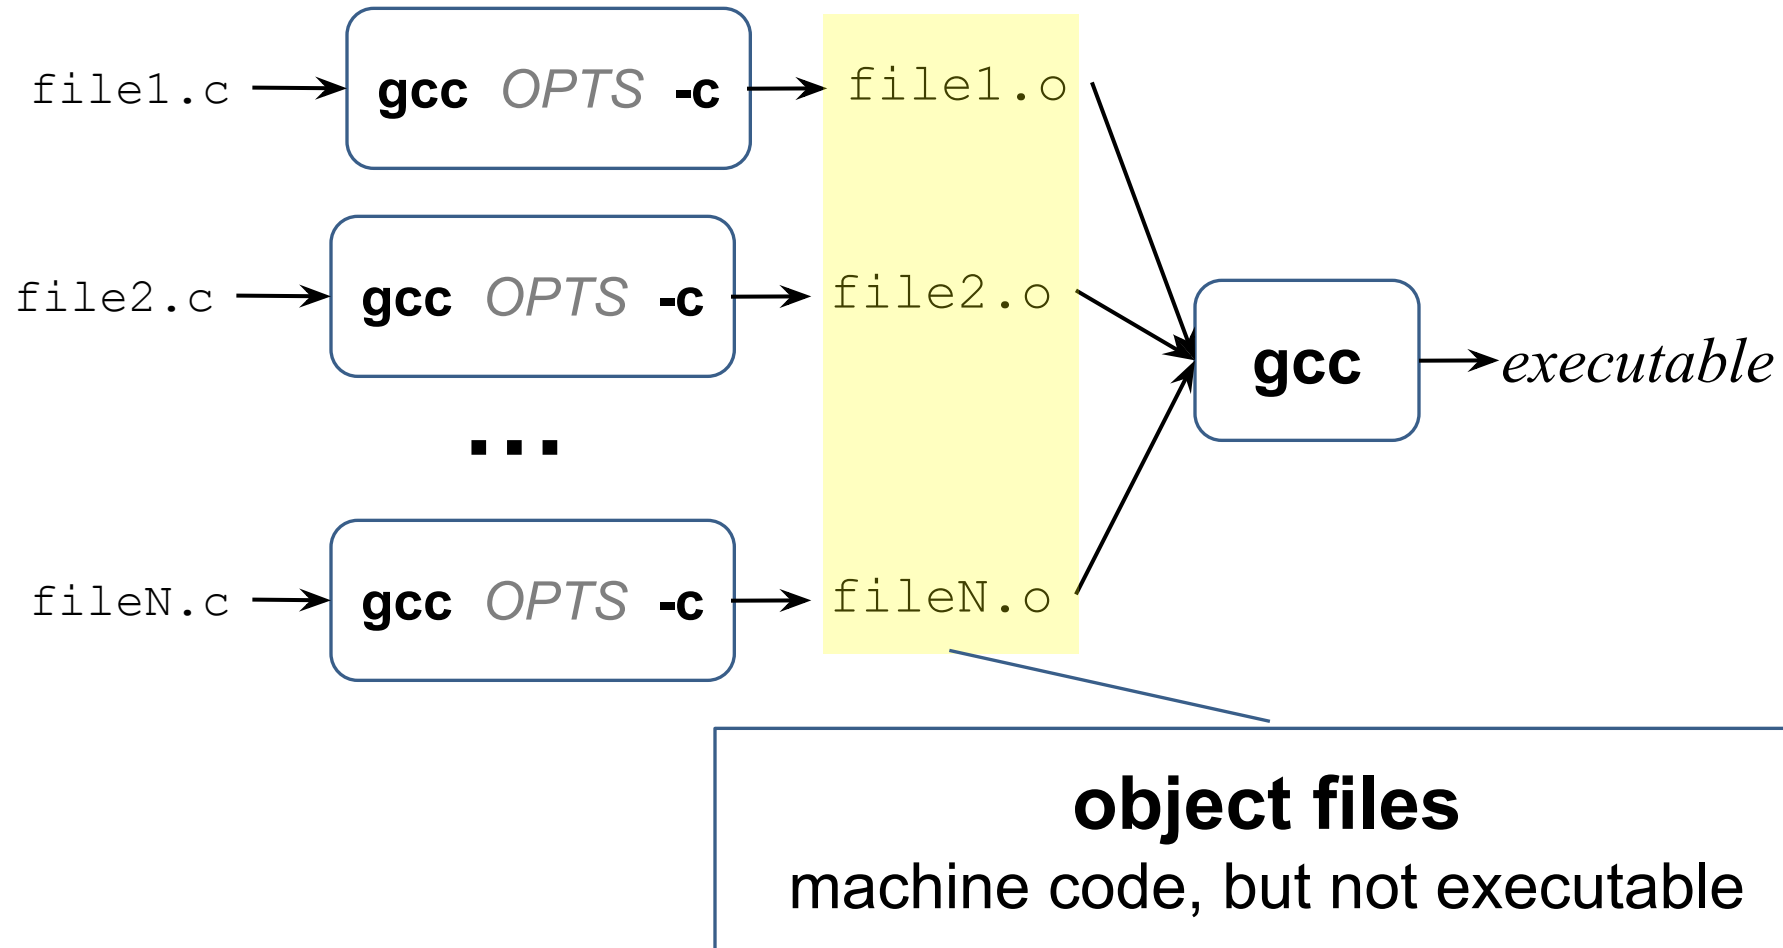
Compiling Multi-File Programs



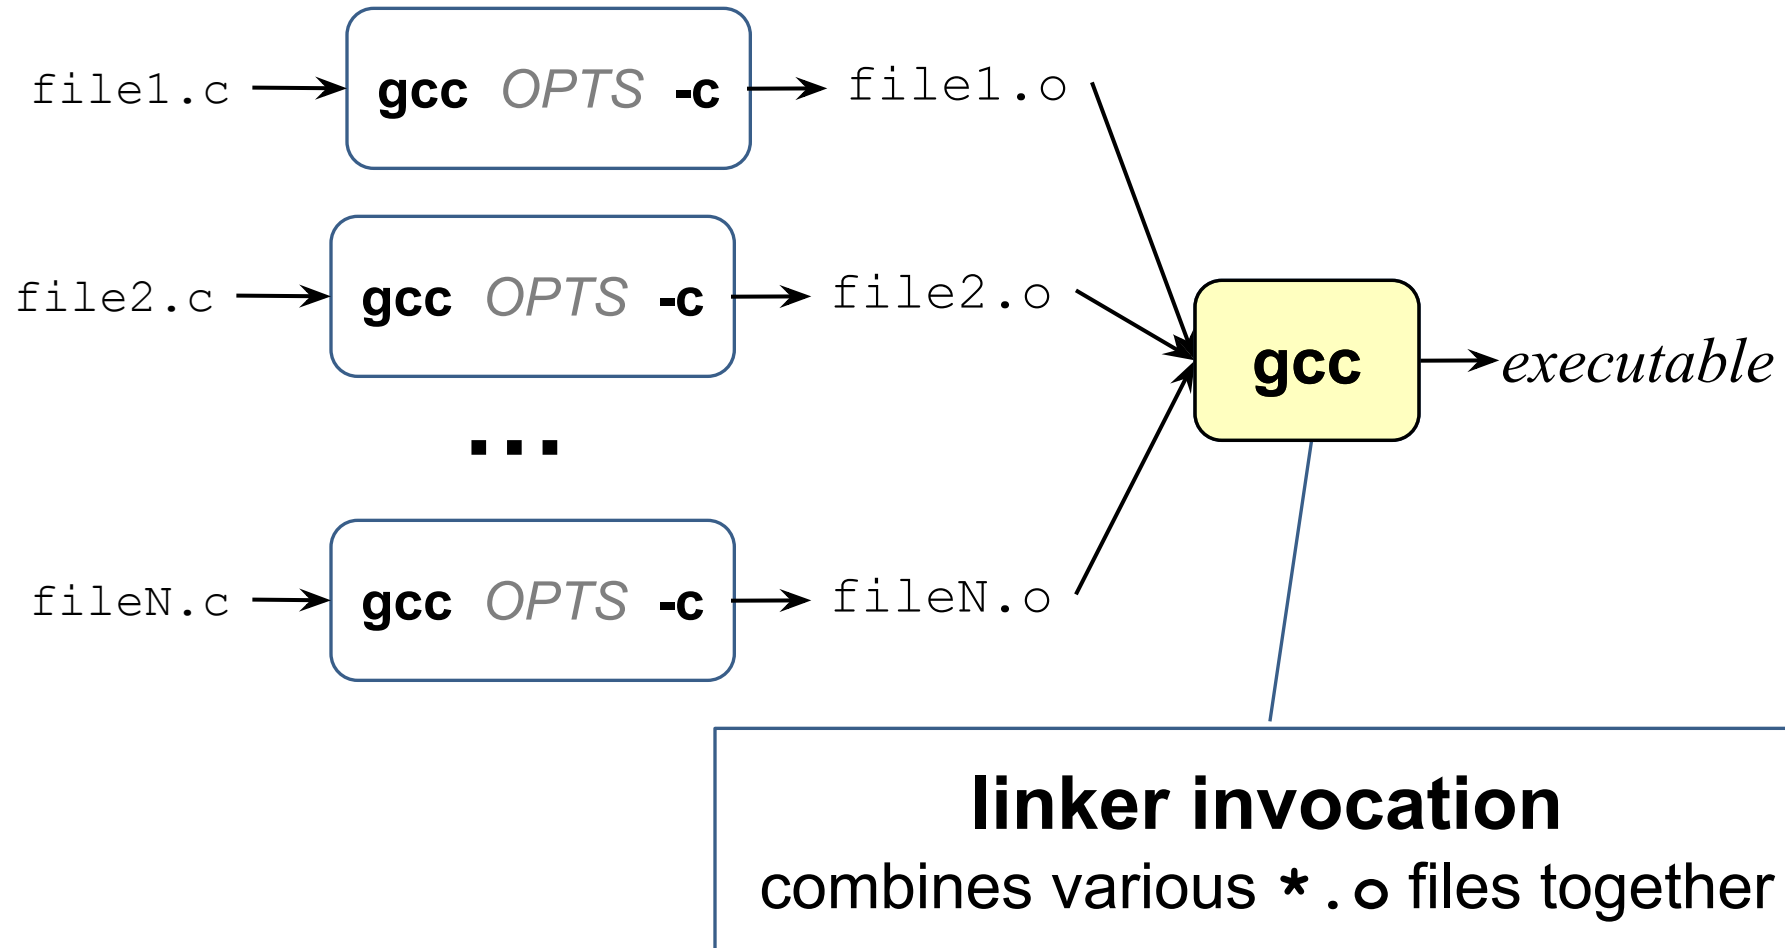
Compiling Multi-File Programs (cont.)



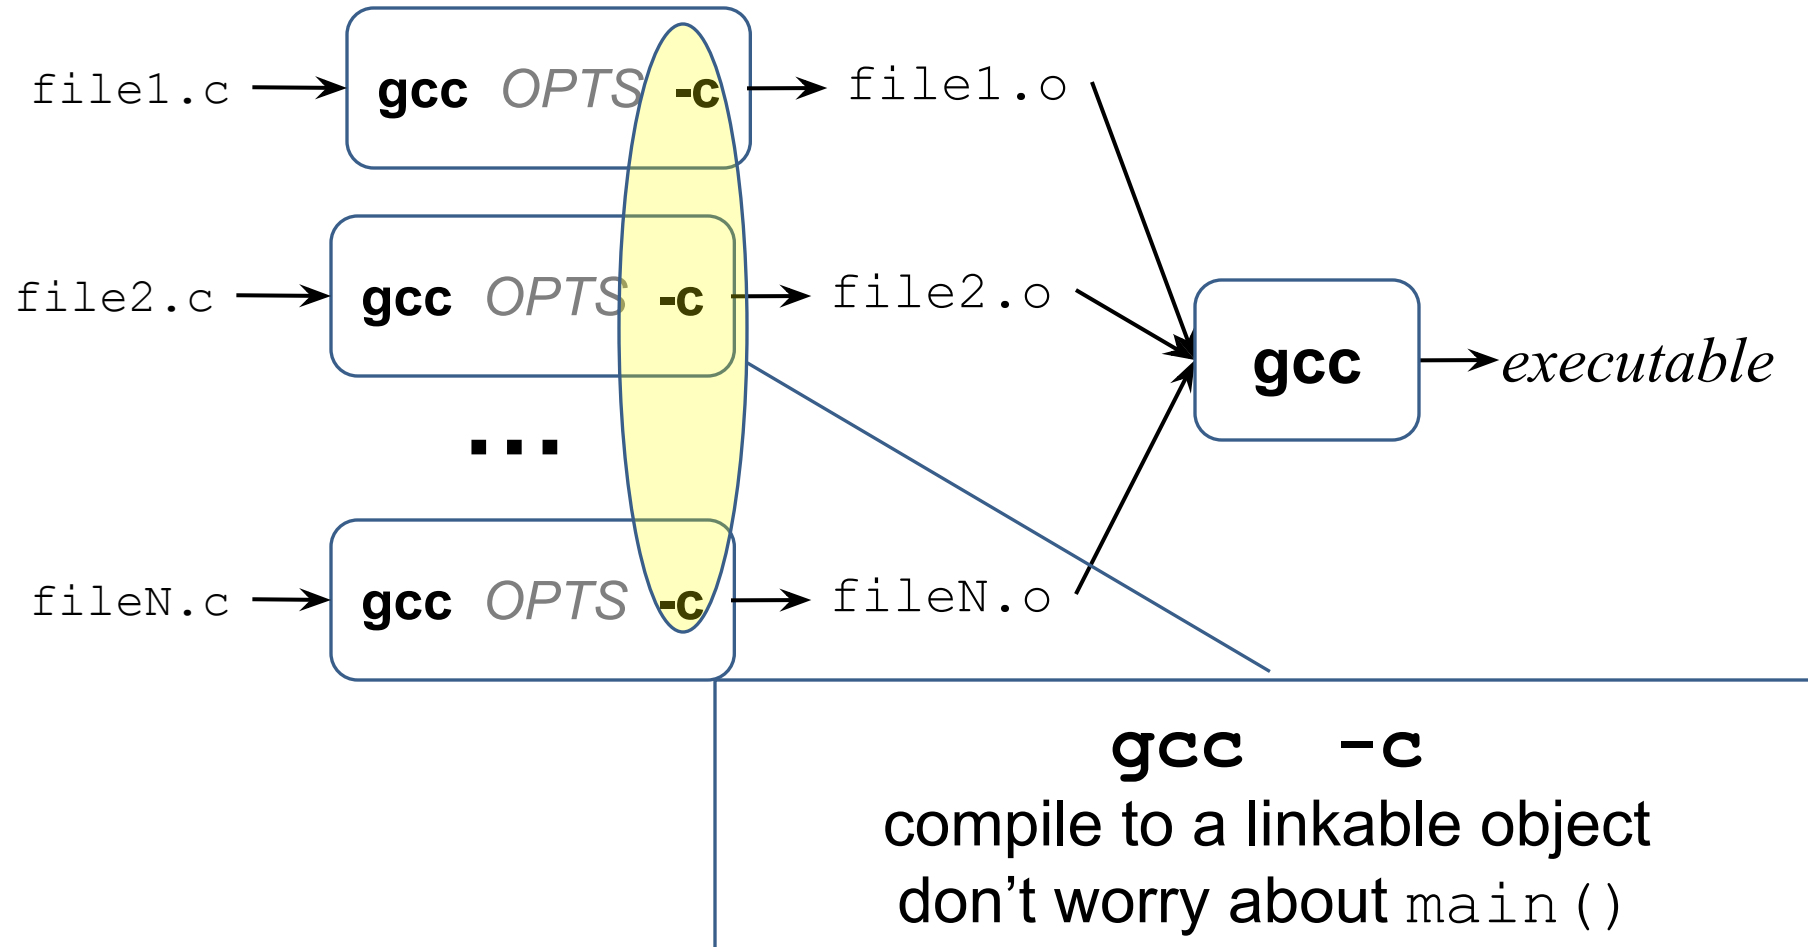
Compiling Multi-File Programs (cont.)



Compiling Multi-File Programs (cont.)

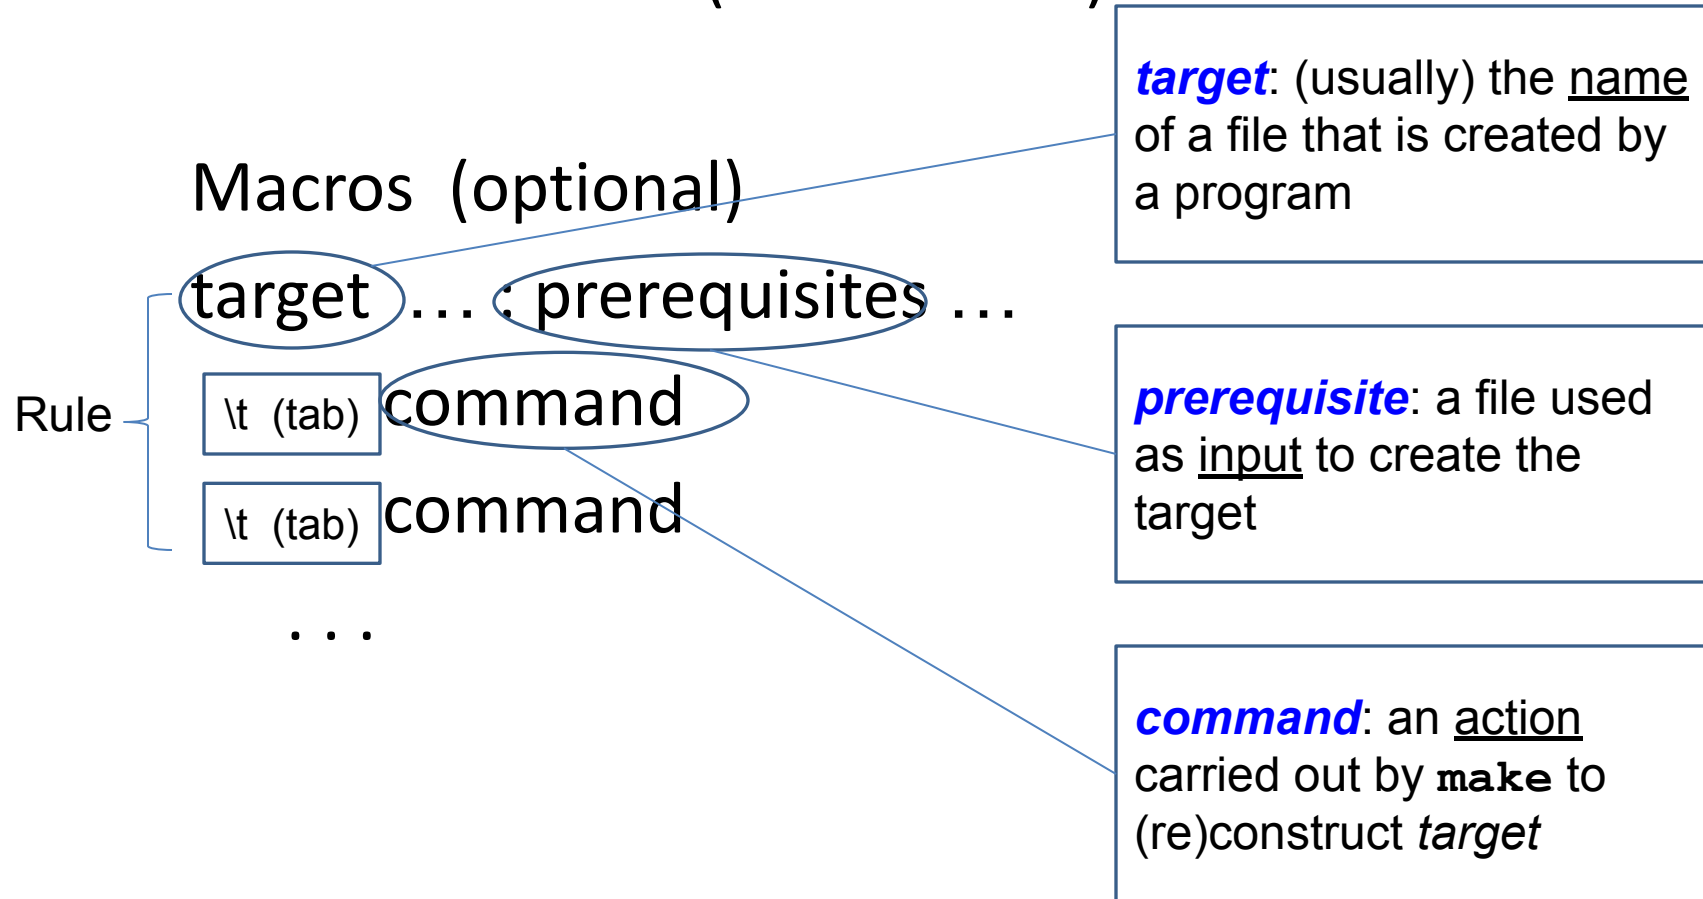


Compiling Multi-File Programs (cont.)




Makefile: Structure

Structure of a makefile (`Makefile`):



Makefile: An Elementary Example

- Dependency structure:



A Makefile file:

```
webserv: webserv.c webserv.h socklib.c socklib.h
gcc -Wall webserv.c socklib.c
```

must be a tab!

Creating a Makefile file

- 1. What are the targets?
 - Figure out which files are *created* from other files and which need to be *re-created* when any of those files change
- 2. For each target, say $f_{○○}$:
 - What are the files which, if changed, would require us to re-create $f_{○○}$?
 - These are the prerequisites for $f_{○○}$ (let's say $bar_1 \dots bar_n$)
- 3. What commands do we use to (re-)create $f_{○○}$?
 - say: $cmd_1 \dots cmd_m$

Creating a Makefile file (cont.)

- The resulting rule for `foo` is:

```
foo: bar1    bar2    ... barn  
  tab cmd1  
  tab cmd2  
    ...  
  tab cmdm
```

or:

```
foo: bar1    bar2    ... barn  
  tab cmd1; cmd2; ... cmdm
```

How to use make?

- Invocation:


`make [-f makeFileName] [target]`

default: make searches (in order) for:
makefile
Makefile

default: builds ***the first target*** in the make file

Makefile: Phony Targets

- A *phony target* is one that is not the name of a file
 - used to run a recipe (i.e., a set of commands) to be executed when an explicit request is made



- “make clean” will remove `a.out` and `*.o` files

Makefile: Phony Targets (cont.)

```
clean:
    rm -f *.o a.out
```

- This won't work if we create a file named "clean" by accident
- Fix:

```
.PHONY: clean
```

```
clean:
    rm *.o a.out
```

cleanup actions will be
executed even if there is a
file named "clean"

Makefile: Macros

- Makes `make` files easier to write, modify
 - Define: `Name = replacement list`
 - Use: `$(Name)`
- Example

```
CC = gcc
OPTLEV = -O2      # optimization level
CFLAGS = -Wall -g -D DEBUG $(OPTLEV) -c
. . .
file1.o : file1.c hdrfile1.h
        $(CC)      $(CFLAGS)  file1.c
```

Makefile: Macros (cont.)

- “gcc -D” defines a macro to be used by the preprocessor

```
// myfile.c
#include <stdio.h>

void main(){
    #ifdef DEBUG
        printf("Debug run\n");
    #else
        printf("Release
run\n");
    #endif
}
```

```
$ gcc -D DEBUG myfile.c -o myfile
$ ./myfile
Debug run
$
```

How make Works


- When invoked, begins processing the appropriate target
- For each **target**, considers the prerequisites it depends on:
 `target : file1 file2 ...`
 - Checks (recursively) whether each of *file_i* (1) exists and (2) is more recent than the files that *file_i* depends on;
 - if not, executes the associated command(s) to update *file_i*
 - Checks whether target exists and is more recent than *file_i*
 - if not, executes the commands associated with *target*
- Commands associated with each rule had better be concatenated by “;”
 - e.g., *cmd₁; cmd₂; cmd₃; ...*
 - If a command returns an error (with a nonzero exit value), `make` abandons that rule; to ignore errors in a command, precede with ‘-’

How make Works (cont.)

Makefile

```
filea: fileb filec  
    cmda  
fileb: filee filed  
    cmdb  
filec: filed filef  
    cmdc  
filed: filef fileg  
    cmdd
```

Dependence structure




How make Works (cont.)

Makefile

▶ `filea: fileb filec`
 `cmda`
`fileb: filee filed`
 `cmdb`
`filec: filed filef`
 `cmdc`
`filed: filef fileg`
 `cmdd`

make execution

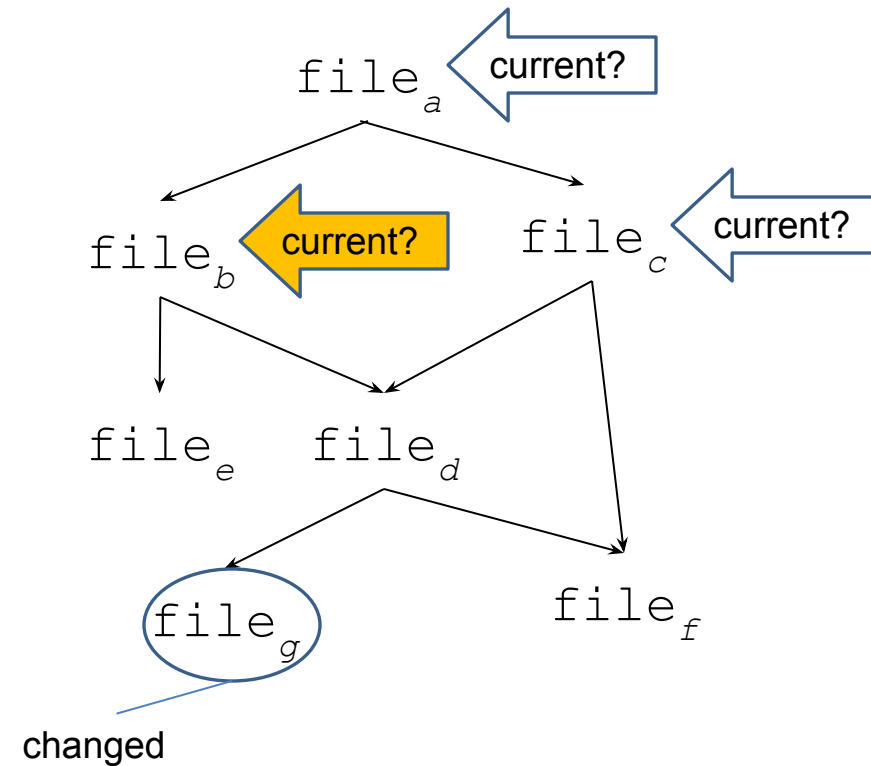


How make Works (cont.)

Makefile

```
filea: fileb filec
cmda
▶ fileb: filee filed
cmdb
filec: filed filef
cmdc
filed: filef fileg
cmdd
```

make **execution**

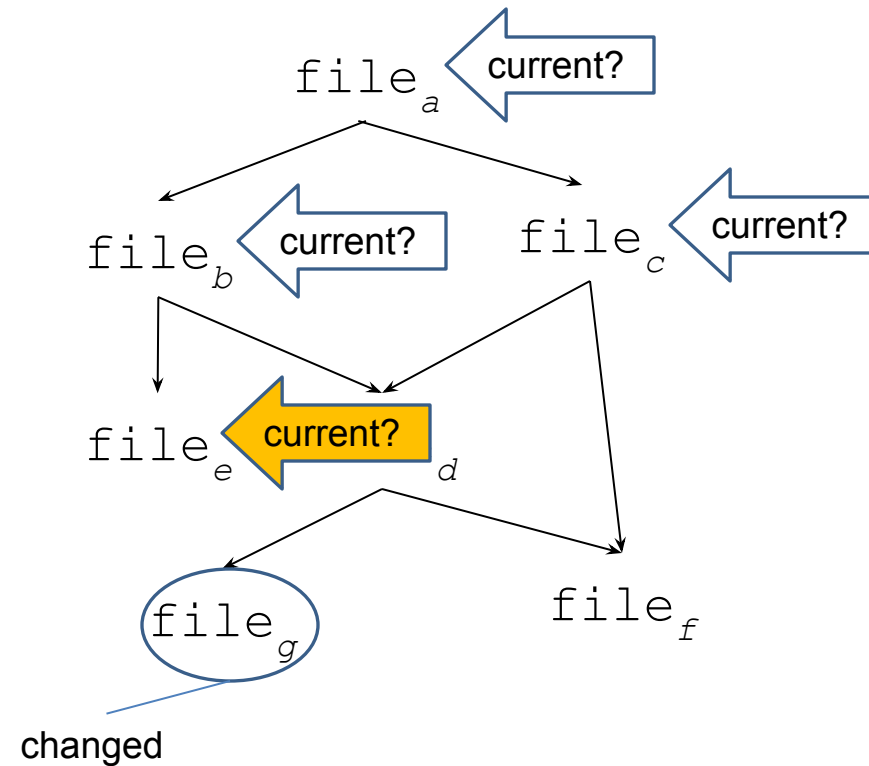


How make Works (cont.)

Makefile

```
filea: fileb filec
cmda
▶ fileb: filee filed
cmdb
filec: filed filef
cmdc
filed: filef fileg
cmdd
```

make **execution**

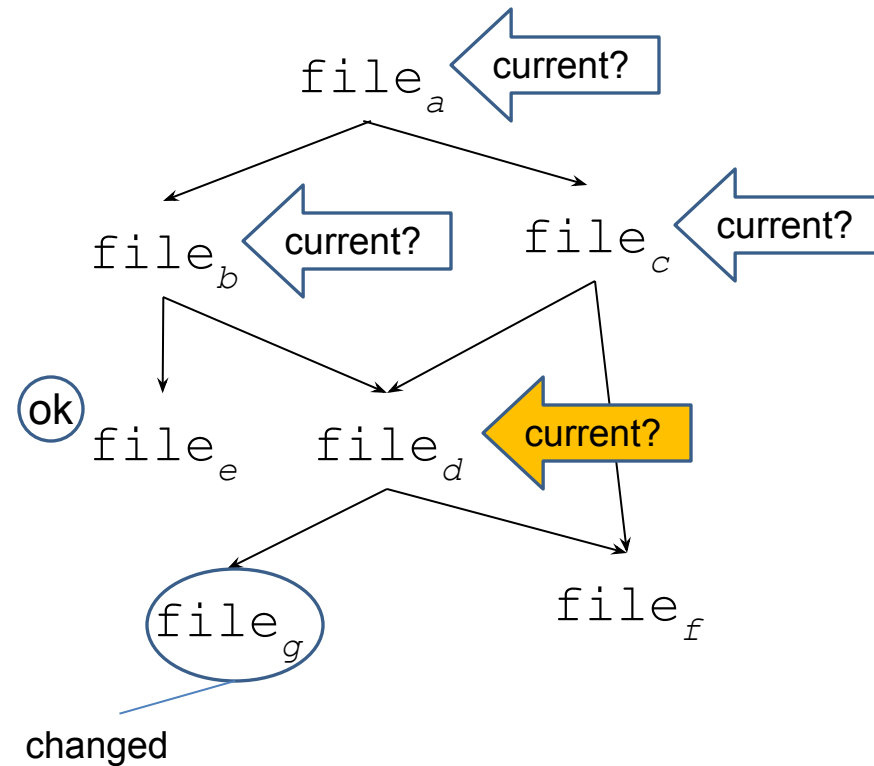


How make Works (cont.)

Makefile

```
filea: fileb filec  
cmda  
fileb: filee filed  
cmdb  
filec: filed filef  
cmdc  
▶ filed: filef fileg  
cmdd
```

make **execution**




How make Works (cont.)

Makefile

```
filea: fileb filec  
  cmda  
fileb: filee filed  
  cmdb  
filec: filed filef  
  cmdc  
filed: filef fileg  
  cmdd
```

make **execution**




How make Works (cont.)

Makefile

```
filea: fileb filec
  cmda
fileb: filee filed
  cmdb
filec: filed filef
  cmdc
filed: filef fileg
  cmdd
```

make **execution**




How make Works (cont.)

Makefile

```
filea: fileb filec  
    cmda  
fileb: filee filed  
    cmdb  
filec: filed filef  
    cmdc  
filed: filef fileg  
    cmdd
```

make **execution**




How make Works (cont.)

Makefile

```
filea: fileb filec  
  cmda  
fileb: filee filed  
  cmdb  
filec: filed filef  
  cmdc  
filed: filef fileg  
  cmdd
```

make **execution**

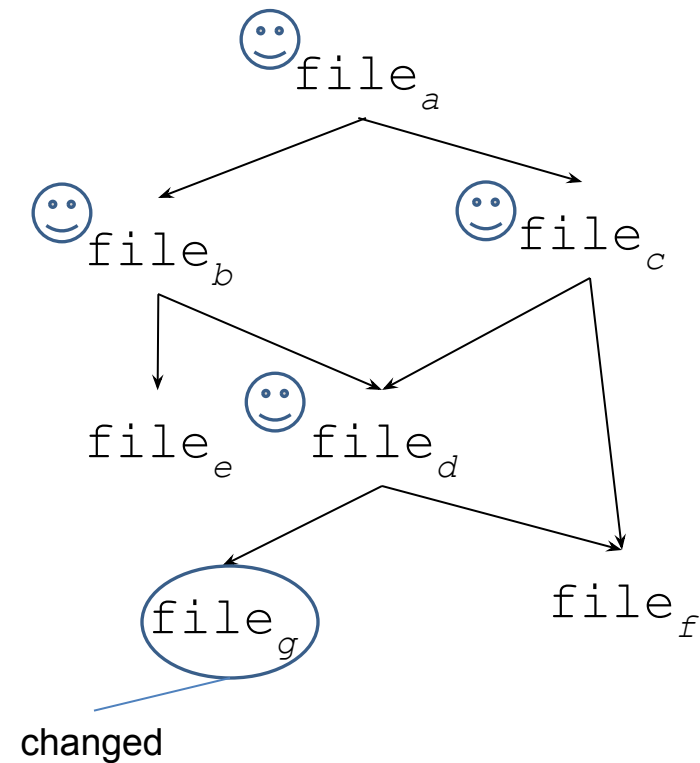


How make Works (cont.)

Makefile

```
filea: fileb filec  
cmda  
fileb: filee filed  
cmdb  
filec: filed filef  
cmdc  
filed: filef fileg  
cmdd
```

make **execution**



Topics Not Covered

- `make` has a lot of functionality we won't get to cover,
 - e.g., implicit rules, implicit variables, conditional parts of make files, recursively running make in subdirectories
- See online `make` tutorials for more information

Summary

- Typically, large applications are written by multiple source files
 - Recompiling all files due to a single update is ridiculous
- `make` is a tool designed to automate a building process for such large programs
 - Recompiles those files that need to be recompiled due to a change(s)
 - Its targets must be determined
 - Phony targets (e.g., `clean`, `install`, etc) may be specified to avoid a conflict and to improve performance
 - Searches by default for `makefile` or `Makefile`
 - Any target that it first meets will be first processed
 - Its macros would help makefiles to be written and modified more easily
 - When invoked, it begins processing the appropriate target