

# System Programming (ELEC462)

*I/O Redirection and Pipes*

Dukyun Nam  
HPC Lab@KNU

# Contents

- Introduction
- A Shell Application: Watch for Users
- Facts About Standard I/O and Redirection
- How to Attach `stdin` to a File
- Redirecting I/O for Another Program: `who > userlist`
- Programming Pipes
- Summary

# Introduction

- Ideas and Skills
  - I/O Redirection: What and Why?
  - Definitions of standard input, output, and error
  - Redirecting standard I/O to files
  - Using fork to redirect I/O for other programs
  - Pipes
  - Using fork with pipes
- System Calls and Functions
  - `dup`, `dup2`
  - `pipe`

# Shell Programming

- How do the following commands work?

```
ls > myfiles
```

```
who | sort > userlist
```

- Questions
  - How does the shell tell a program to send its output to a file instead of the screen?
  - How does the shell connect the output stream of one process to the input stream of another process?
  - What's the real meaning of 'standard input'?
- In this chapter, we'll focus on a particular form of IPC (Interprocess communication)
  - Input/output (I/O) redirection and pipes

# A Shell Application: Watch for Users (watch.sh)

- Consider the following program.
  - You have a list of buddies accessing the same Linux machine
  - A program notifies you when people “log in” or “log out” of the system so you can watch for your peers
- Let’s write a watch shell script!

```
logic
-----
get list of users (call it prev)
while true
    sleep
    get list of users (call it curr)
    compare lists
        in prev, not in curr -> logout

        in curr, not in prev -> login
    make prev = curr
repeat
```

```
#!/bin/sh
#
# watch.sh - a simple version of the watch utility, written in sh
#
who | sort > prev                # get initial user list
while true                       # true is a program: exit(0);
do
    sleep 10                     # wait a while
    who | sort > current          # get current user list
    echo "Logged out:"           # print header
    comm -23 prev current        # and results
    echo "Logged in:"           # header
    comm -13 prev current        # and results
    mv current prev             # make now past
done
```

# A Shell Application: Watch for Users (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab11$ ./watch.sh
Logged out:
Logged in:
Logged out:
dynam pts/1 2022-11-15 14:59
Logged in:
Logged out:
Logged in:
Logged out:
Logged in:
dynam pts/1 2022-11-15 15:00
Logged out:
```

- In WSL, if you have no utmp file, run the following commands

```
$ sudo bash -c "echo '[1] [00049] [~~ ] [runlevel] [~          ] [4.4.0-17115-Micoroso]
[0.0.0.0          ] [Wed Feb 28 13:27:14 2018 STD]' | utmpdump -r > /var/run/utmp"


$ sudo login -f username
```

# A Shell Application: Watch for Users (cont.)

- `who | sort > prev`

- Tells the three things to the shell:

- 1) Run the commands `who` and `sort` at the same time
- 2) Send the output of `who` directly to the input to `sort`



< Connecting output of `who` to input of `sort` >

- Not necessary to finish analyzing the `utmp` file before `sort` begins its task
- The two processes are scheduled to run in small time slices, sharing CPU time with other processes on the system
- 3) Send the output of `sort` into a file, called `prev`: what if it already exists?

# A Shell Application: Watch for Users (cont.)


- **comm**: a command to find lines common to two sorted files
  - Compares two sorted lists and prints out the three columns
  - In the example, it produces exactly the two sets we want
    - Logouts: who did leave? (`comm -23 prev current`)
    - Logins: who are new? (`comm -13 prev current`)

```
COMM(1)                                User Commands                                COMM(1)
NAME
  comm - compare two sorted files line by line
SYNOPSIS
  comm [OPTION]... FILE1 FILE2
DESCRIPTION
  Compare sorted files FILE1 and FILE2 line by line.

  When FILE1 or FILE2 (not both) is -, read standard input.

  With no options, produce three-column output.  Column one contains lines
  unique to FILE1, column two contains lines unique to FILE2, and column three
  contains lines common to both files.

  -1  suppress column 1 (lines unique to FILE1)
  -2  suppress column 2 (lines unique to FILE2)
  -3  suppress column 3 (lines that appear in both files)
```

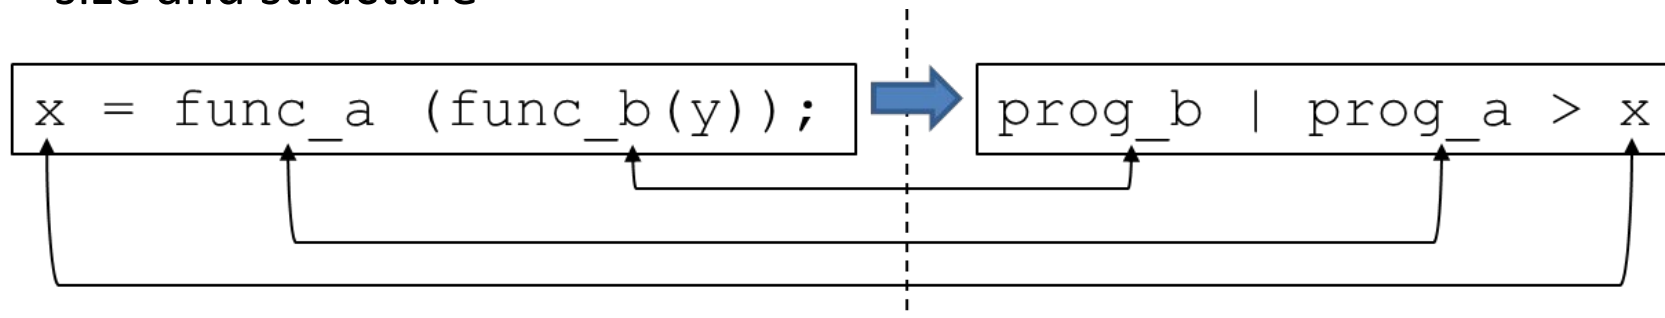


< comm compares two lists and outputs three sets >



# A Shell Application: Watch for Users (cont.)

- Three important ideas behind `watch.sh`
  - (1) Power of shell scripts
    - Easier and quicker than C (or other programming languages requiring compiling)
  - (2) Flexibility of software tools
    - Each tool (or command or program) does one specific, general task
  - (3) Use and value of I/O redirection and pipes
- And one more ...
  - The script shows how to use the '`>`' operator to treat files as variables of arbitrary size and structure




# A Shell Application: Watch for Users (cont.)

- Questions
  - How does all of *these connected programs* jointly work?
  - What role does the *shell* play in connecting processes?
  - What role does the *kernel* play to get the processes to work?
  - What role do the *individual programs* play?

# Facts about Standard I/O & Redirection


- All Linux/Unix I/O redirection based on
  - Principle of **standard streams of data**
  - e.g., The task of `sort`
    - Reads bytes from one stream of data
      - Then it performs the sorting task on the read byte stream
    - Writes the sorted results to another stream
    - Reports any errors to a third stream
  - The three channels for data flow are as follows:
    - **Standard input:** The stream of data to process
    - **Standard output:** The stream of result data
    - **Standard error:** A stream of error messages



< A software tool reads input and writes output and errors >

# Fact 1: 3 Standard File Descriptors

- All Linux/Unix tools make use of the three-stream model
  - Each of the stream is a specific *file descriptor* (fd)
  - Linux/Unix tools find file descriptors 0, 1, and 2 already open for reading, writing, and writing, respectively



# Default Connections: the `tty`

- When you run a Linux tool (`who`, `sort`, `comm`, ...) on your shell
  - The three (`stdin`, `stdout`, and `stderr`) streams are usually “connected” to your terminal
- So the tool `sort`:
  - Reads from the keyboard (`stdin`)
  - Writes output (`stdout`)
  - Writes error messages to the screen (`stderr`)
    - If Ctrl-D is pressed, then `sort` starts to begin the input and writes the result to `stdout`.
- Most tools process data from files or `stdin`
  - Given file names => they read input from those files
  - No files given => they read from standard input



```
dynam@DESKTOP-Q4IJB7:~$ sort
korea
knu
cse
daegu

cse
daegu
knu
korea
```

# The Shell, Not the Program, Redirects I/O

- `cmd > filename`
  - Tells the shell to attach `fd 1` to a file
    - By using the output redirection notation as specified above

```
/* listargs.c
 *
 *      print the number of command line args, list the args,
 *      then print a message to stderr
 */
#include <stdio.h>
#include <unistd.h>

int main( int ac, char *av[] )
{
    int    i;

    printf("Number of args: %d, Args are:\n", ac);
    for(i=0; i<ac; i++)
        printf("args[%d] %s\n", i, av[i]);

    fprintf(stderr, "This message is sent to stderr.\n");

    return 0;
}
```

# The Shell, Not the Program, Redirects I/O (cont.)

- `listargs` prints to standard output the list of command-line arguments
  - Does not print the redirection symbol and filename

```
$ cc listargs.c -o listargs
$ ./listargs testing one two
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
This message is sent to stderr.
$ ./listargs testing one two > xyz
This message is sent to stderr.
$ cat xyz
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
$ ./listargs testing >xyz one two 2> oops
$ cat xyz
args[0] ./listargs
args[1] testing
args[2] one
args[3] two
$ cat oops
This message is sent to stderr.
```

# Some Important Facts ...

- The shell *doesn't* pass the redirection symbol and filename to the command
- The redirection request can appear *anywhere* in the command.
  - Doesn't require spaces around the redirection symbol (>)
    - Even a command like '`> listing ls`' is acceptable
  - Doesn't terminate the command and arguments: just an added request
- Many shells provide notation for redirecting other fds
  - e.g., `2>filename`
    - Redirects fd 2, that is, standard error, to the named file




# Understanding I/O Redirection

- Goal
  - Understand how I/O redirection works
  - AND learn how to write programs that use it
- Method: write programs that do
  - `sort < data`      attach stdin to a file
  - `who > userlist`      attach stdout to a file
  - `who | sort`      attach stdout to stdin

# Fact 2: the “Lowest-Available- $\text{fd}$ ” Principle

- The meaning of a file descriptor? **An array index!**
  - Each process has a collection of files it has open
    - Those “open” files are kept in an array
  - So a file descriptor: simply an index for an item in that array
  - Making a new connection with file descriptors is like receiving a connection on a multiline phone  
the next incoming call => mapped to the lowest available line

Unix always assigns new connections to the lowest available file descriptor.



< The “lowest-available-file-descriptor” rule >

# How to Attach `stdin` to a File

- Standard I/O?
  - Standard input (`stdin`), output (`stdout`), and error (`stderr`) are indicated by file descriptors 0,1 and 2, respectively
  - Three predefined streams:


Integer value	Name	file stream in <code>&lt;stdio.h&gt;</code>
0	standard input	<code>stdin</code>
1	standard output	<code>stdout</code>
2	standard error	<code>stderr</code>

# How to Attach `stdin` to a File (cont.)

- How does a Linux program redirect `stdin` in order for data to come from a file?
  - Linux processes don't read from *files*, but actually from *file descriptors*
  - If `fd 0` is attached to a file, then the attached file becomes the source for standard input
- There are three methods for attaching `stdin` to a file
  - Method 1: *Close-then-open*
  - Method 2: *Open-close-dup-close*
  - Method 3: *Open-dup2-close*

# Method 1: Close-then-Open


- Step 1) *Starting* with the three standard streams connected to the terminal driver
  - File descriptors 0, 1, 2 attached to `/dev/tty`
    - 0 for reading
    - 1 for writing
    - 2 for writing



< Typical starting configuration >


# Method 1: Close-then-Open (cont.)

- Step 2) *Then, close(0)*: hang up the connection to stdin
  - Breaks the connection from standard input to the driver
    - See the “unused” element in the array below
  - If the process closes file descriptor 0,
    - that entry in its array of I/O channels is free



# Method 1: Close-then-Open (cont.)

- Step 3) *Finally*, `open(filename, O_RDONLY)`
  - If the process opens another file,
    - that connection is attached to the FIRST FREE entry in the array of I/O channels



# Method 1: Close-then-Open (cont.)

- stdinredir1.c

```
/* stdinredir1.c
 *   purpose: show how to redirect standard input by replacing file
 *           descriptor 0 with a connection to a file.
 *   action: reads three lines from standard input, then
 *           closes fd 0, opens a disk file, then reads in
 *           three more lines from standard input
 */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>

int main()
{
    int fd ;
    char line[100];

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );

    /* redirect input */

    close(0);
    fd = open("/etc/passwd", O_RDONLY);
    if ( fd != 0 ) {
        fprintf(stderr, "Could not open data as fd 0\n");
        exit(1);
    }

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );

    return 0;
}
```

What's returned?





# Method 1: Close-then-Open (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab11$ ./stdinredir1
line1
line1
testing line2
testing line2
line3
line3
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
```


```
dynam@DESKTOP-Q4IJB7:~/lab11$ head -3 /etc/passwd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
dynam@DESKTOP-Q4IJB7:~/lab11$
```

# Method 2: `Open . . close . . dup . . close`

- The system call `dup` makes a second connection to an existing file descriptor
  - `open (file)` : open a file to which `stdin` should be attached
    - Will return a file descriptor with a non-zero, as 0 is still in use
  - `close (0)` : close fd 0, which becomes now “unused”
  - `dup (fd)` : makes a “duplicate” of `fd`
    - Uses the lowest but not yet used number for fd
    - So what would be the number?
      - The duplicate of the connection to the file: located at spot 0 in the array open files
  - `close (fd)` : invokes `close (fd)` , the original connection to the file
    - Leaving only the connection to file descriptor 0

# Method 2: Open..close..dup..close (cont.)

```
/* redirect input */
fd = open("/etc/passwd", O_RDONLY);    /* open the disk file */
#ifdef CLOSE_DUP
close(0);
newfd = dup(fd);                      /* copy open fd to 0 */
#else
newfd = dup2(fd, 0);                  /* close 0, dup fd to 0 */
#endif
if ( newfd != 0 ){
    fprintf(stderr, "Could not duplicate fd to 0\n");
    exit(1);
}
close(fd);                            /* close original fd */
```



< Using dup to redirect >

# Method 2: Open..close..dup..close (cont.)

- stdinredir2.c

```
/* stdinredir2.c
 *      shows two more methods for redirecting standard input
 *      use #define to set one or the other
 */
#include      <stdio.h>
#include      <stdlib.h>
#include      <unistd.h>
#include      <fcntl.h>

#define CLOSE_DUP          /* open, close, dup, close */
/* #define      USE_DUP2      /* open, dup2, close */

int main()
{
    int      fd ;
    int      newfd;
    char      line[100];

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
```

```
    /* redirect input */
    fd = open("/etc/passwd", O_RDONLY);      /* open the disk file */
#ifdef CLOSE_DUP
    close(0);
    newfd = dup(fd);                        /* copy open fd to 0 */
#else
    newfd = dup2(fd,0);                     /* close 0, dup fd to 0 */
#endif
    if ( newfd != 0 ){
        fprintf(stderr,"Could not duplicate fd to 0\n");
        exit(1);
    }
    close(fd);                             /* close original fd */

    /* read and print three lines */

    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );
    fgets( line, 100, stdin ); printf("%s", line );

    return 0;
}
```

# Method 2: Open..close..dup..close (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab11$ ./stdinredir2  
line1  
line1  
line2  
line2  
line3  
line3  
root:x:0:0:root:/root:/bin/bash  
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin  
bin:x:2:2:bin:/bin:/usr/sbin/nologin
```

# Method 3: Open . . dup2 . . close

- The code for `stdinredir2.c` includes `#ifdef`-ed code
  - to replace the `close(0)` and `dup(fd)` system calls with `dup2 (fd, 0)`

```
/* stdinredir2.c
 *      shows two more methods for redirecting standard input
 *      use #define to set one or the other
 */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>

/* #define CLOSE_DUP          /* open, close, dup, close */
#define USE_DUP2             /* open, dup2, close */
```

# System Call Summary: dup

dup, dup2	
PURPOSE	Copy a file descriptor
INCLUDE	#include <unistd.h>
USAGE	newfd = dup(oldfd); newfd = dup2(oldfd, newfd);
ARGS	oldfd    file descriptor to copy newfd    copy of oldfd
RETURNS	-1        if error newfd    new file descriptor

- The system call `dup` creates a copy of *oldfd* as *newfd*
  - The system call `dup2` gets *newfd* associated with the copy of *oldfd*
- ⇒ The two newfds actually refer to the same open file pointed by oldfd!

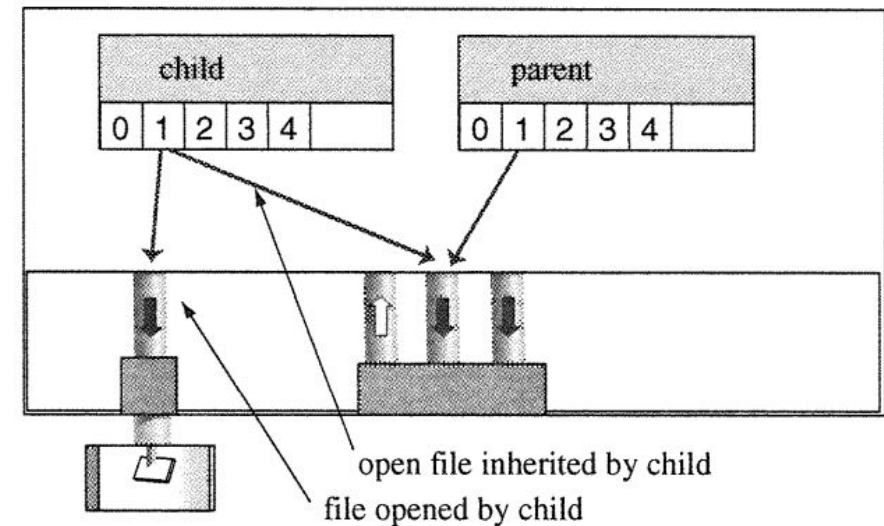


# Redirecting I/O for Another Program

- `who > userlist`
  - The shell runs the command `who`
    - with the standard output of `who` attached to the file called `userlist`
- Key: the *split second* between `fork` and `exec`
  - After `fork`, the child is running the shell code, but is about to execute `exec`
  - `exec` will replace the program running in the process with no change of attributes or connections of the process
  - The process will have the same file descriptors it had before the `exec`

The child inherits from the parent the pointers to open files. The child redirects standard output:

```
close(1);  
creat("f");  
exec();
```




< The shell redirects output for a child ><sup>32</sup>




# Redirecting I/O for Another Program: `who > userlist (cont.)`

- Step 1. Start here: a process that is about to `fork` and its `stdout`
- Step 2. After parent calls `fork()`: `stdout` of a child copied from parent



< A process about to fork and its standard output >

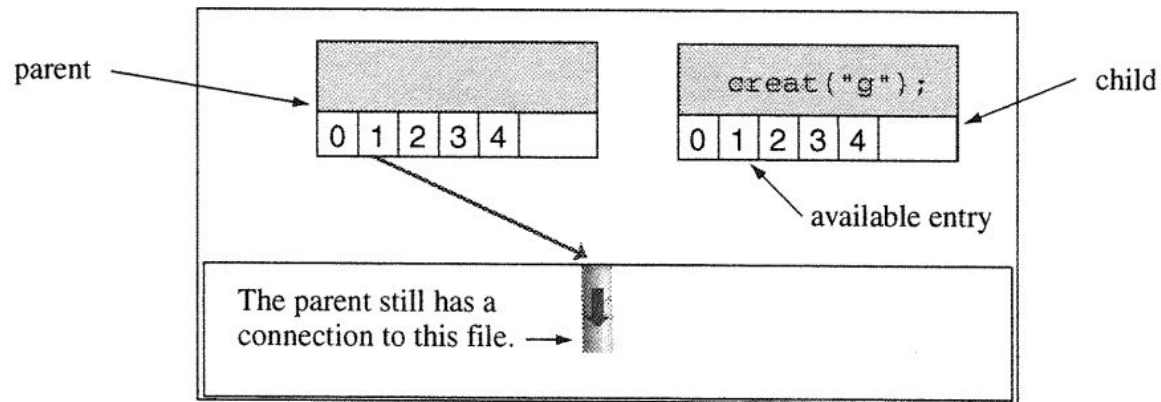


< Standard output of child is copied from parent >

# Redirecting I/O for Another Program: who > userlist (cont.)

- Step 3. After child calls


`close(1)`: child can close its  
`stdout`



< The child can close its standard output >

- Step 4. After child calls

`creat("g", m)`: child opens a  
new file, taking `fd = 1`



< Child opens a new file, getting `fd = 1` >

# Redirecting I/O for Another Program:

## `who > userlist` (cont.)

- Step 5. After child execs a new program (e.g., `who`)

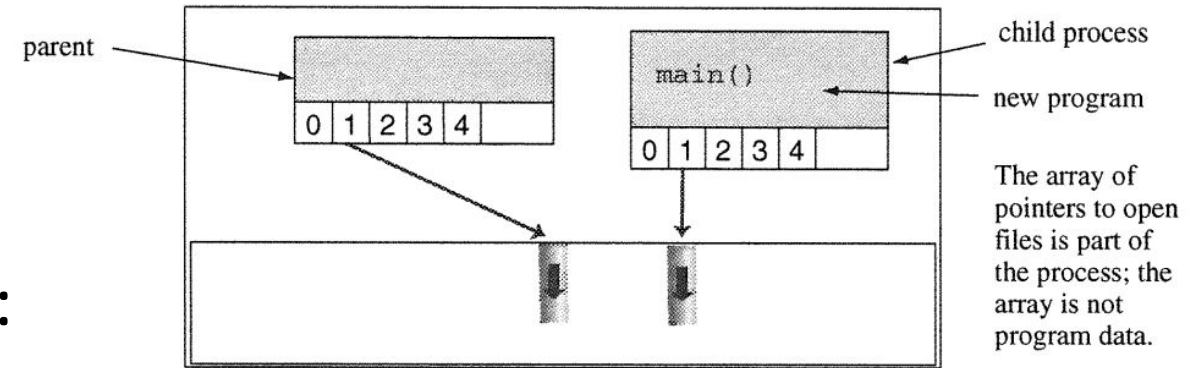
- child runs a program with the new standard output (e.g., `userlist`)

- The code and data for the shell are:

- Removed from the child process
- Replaced by the code and data for `who`

- But the file descriptors are retained across the `exec`

- Note that open files are not part of the code nor data of a program (here, `who`); they are attributes of a process



# Redirecting I/O for Another Program: who > userlist (cont.)

- whotofile.c

```
/* whotofile.c
 *   purpose: show how to redirect output for another program
 *   idea: fork, then in the child, redirect output, then exec
 */
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/wait.h>

int main()
{
    int pid ;
    int fd;

    printf("About to run who into a file\n");

    /* create a new process or quit */
    if( (pid = fork() ) == -1 ){
        perror("fork"); exit(1);
    }
    /* child does the work */
    if ( pid == 0 ){
        close(1);
        fd = creat( "userlist", 0644 );
        execlp( "who", "who", NULL );
        perror("execlp");
        exit(1);
    }
    /* parent waits then reports */
    if ( pid != 0 ){
        wait(NULL);
        printf("Done running who. results in userlist\n");
    }

    return 0;
}
```

# Redirecting I/O for Another Program:

## who > userlist (cont.)

- Execution


```
dynam@DESKTOP-Q4IJB7:~/lab11$ make
cc -o whotofile whotofile.c
dynam@DESKTOP-Q4IJB7:~/lab11$ ./whotofile
About to run who into a file
Done running who.  results in userlist
dynam@DESKTOP-Q4IJB7:~/lab11$ cat userlist
dynam    pts/0          2022-11-15 14:56
dynam    pts/1          2022-11-15 15:11
dynam@DESKTOP-Q4IJB7:~/lab11$ who
dynam    pts/0          2022-11-15 14:56
dynam    pts/1          2022-11-15 15:11
```

# Summary of Redirection to Files

- (1) File descriptors 0, 1, and 2 represent standard input, output, and error, respectively
- (2) The kernel always uses the lowest numbered unused file descriptor
- (3) The set of file descriptors is passed unchanged across `exec` calls
  - To make I/O redirection to another program, the shell uses the interval in the child process between `fork` and `exec`
    - Reason: for the purpose of attaching the standard data streams to (external) files

# What is a Pipe? How It Works?

- Pipe
  - “One-way” data channel in the kernel
  - Has a “writing” end and “reading” end
    - e.g., `who | sort`




< Two processes connected by a pipe >



# How to Create a Pipe? Use the System call, `pipe()`

<b>pipe</b>	
<b>PURPOSE</b>	Create a pipe
<b>INCLUDE</b>	<code>#include &lt;unistd.h&gt;</code>
<b>USAGE</b>	<code>result = pipe(int array[2])</code>
<b>ARGS</b>	array    an array of two ints
<b>RETURNS</b>	-1      if error 0        if success

- `array[1]`: fd of the writing end
- `array[0]`: fd of the reading end






# Creating a Pipe


- Pipe creation by a process
  - Pipe uses the *lowest-numbered* available file descriptors: like `open ( )`

Before pipe



The process has some usual files open.

After pipe



The kernel creates a pipe and sets file descriptors.

< A process creates a pipe >

# Creating a Pipe (cont.)

- pipedemo.c

```
/* pipedemo.c  * Demonstrates: how to create and use a pipe
 *             * Effect: creates a pipe, writes into writing
 *             * end, then runs around and reads from reading
 *             * end.  A little weird, but demonstrates the idea.
 */
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>
#include <string.h>
#include <sys/wait.h>

int main()
{
    int len, i, apipe[2];    /* two file descriptors */
    char buf[BUFSIZ];        /* for reading end */
```

```
/* get a pipe */
if ( pipe ( apipe ) == -1 ){
    perror("could not make pipe");
    exit(1);
}
printf("Got a pipe! It is file descriptors: { %d %d }\n",
      apipe[0], apipe[1]);

/* read from stdin, write into pipe, read from pipe, print */
while ( fgets(buf, BUFSIZ, stdin) ){
    len = strlen( buf );
    if ( write( apipe[1], buf, len) != len ){ /* send */
        perror("writing to pipe");          /* down */
        break;                               /* pipe */
    }
    for ( i = 0 ; i<len ; i++ )              /* wipe */
        buf[i] = 'X' ;
    len = read( apipe[0], buf, BUFSIZ ) ;     /* read */
    if ( len == -1 ){                         /* from */
        perror("reading from pipe");         /* pipe */
        break;
    }
    if ( write( 1 , buf, len ) != len ){      /* send */
        perror("writing to stdout");         /* to */
        break;                               /* stdout */
    }
}

return 0;
}
```


# Creating a Pipe (cont.)

- Execution
  - Creates a pipe and then
  - Uses the pipe to send the data itself

```
dynam@DESKTOP-Q4IJB7:~/lab11$ ./pipedemo
Got a pipe! It is file descriptors: { 3 4 }
hello
hello
^C
```

# Creating a Pipe (cont.)

- Depicts the flow of bytes:
  - From keyboard to process: 1 → 2
  - From process to pipe: 3 → 4
  - From pipe to process, and: 5 → 6
  - From process back to terminal: 6 → 7





< Data flow in `pipdemo.c` >

- Indeed, `pipe` and `fork` can be “combined” to connect two processes
- So the pipe can be *shared* between them

# Using fork to Share a Pipe


- Sharing a pipe
  - A process calls `pipe`
    - The process has connections to both ends of the pipe
  - The kernel creates a pipe and adds to the array of file descriptors pointers to the ends of the pipe
  - The process then calls `fork`
    - The child process also has connections to the pipe
  - The kernel creates a new process, and copies into that process the array of file descriptors from the parent
- Both processes now have access to both ends of one pipe!





# Using fork to Share a Pipe (cont.)

- Sharing a pipe: Interprocess data flow
  - Parent/child can *write* bytes to the *writing end* of the pipe
  - Parent/child can *read* bytes from the *reading end* of the pipe



- A pipe is “most effective” when one process writes data and the other processes reads the data on the same host
  - Of course, processes can read and write together

# Using fork to Share a Pipe (cont.)

- Shows how to combine pipe and fork
  - To create a pair of processes via pipe communication
- `pipdemo2.c`

```
/* pipdemo2.c  * Demonstrates how pipe is duplicated in fork()
 *              * Parent continues to write and read pipe,
 *              * but child also writes to the pipe
 */
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/wait.h>

#define CHILD_MESS    "I want a cookie\n"
#define PAR_MESS     "testing..\n"
#define oops(m,x)    { perror(m); exit(x); }
```

# Using fork to Share a Pipe (cont.)

- pipedemo2.c (cont.)

```
int main()
{
    int    pipefd[2];           /* the pipe    */
    int    len;                 /* for write   */
    char   buf[BUFSIZ];        /* for read    */
    int    read_len;

    if ( pipe( pipefd ) == -1 )
        oops("cannot get a pipe", 1);

    switch( fork() ){
        case -1:
            oops("cannot fork", 2);

            /* child writes to pipe every 5 seconds */
        case 0:
            len = strlen(CHILD_MESS);
            while ( 1 ){
                if (write( pipefd[1], CHILD_MESS, len) != len )
                    oops("write", 3);
                sleep(5);
            }

            /* parent reads from pipe and also writes to pipe */
        default:
            len = strlen( PAR_MESS );
            while ( 1 ){
                if ( write( pipefd[1], PAR_MESS, len)!=len )
                    oops("write", 4);
                sleep(1);
                read_len = read( pipefd[0], buf, BUFSIZ );
                if ( read_len <= 0 )
                    break;
                write( 1 , buf, read_len );
            }
    }

    return 0;
}
```



# Using fork to Share a Pipe (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab11$ ./pipedemo2
testing..
I want a cookie
testing..
testing..
testing..
testing..
I want a cookie
testing..
testing..
testing..
testing..
testing..
I want a cookie
testing..
testing..
testing..
^C
dynam@DESKTOP-Q4IJB7:~/lab11$ |
```

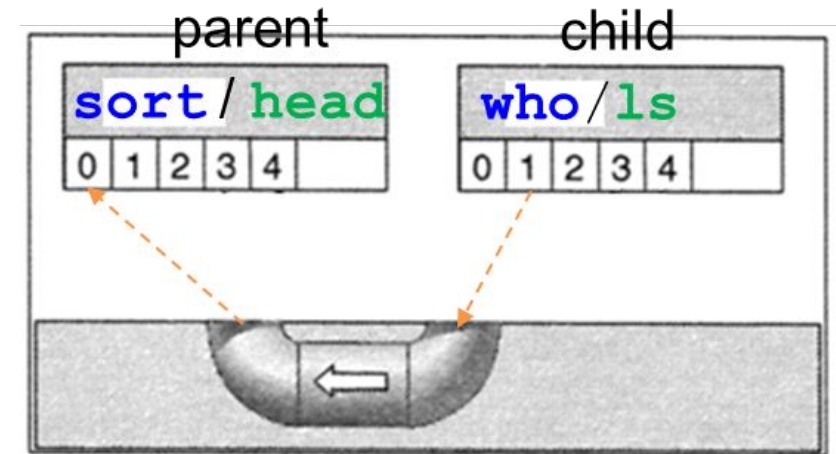
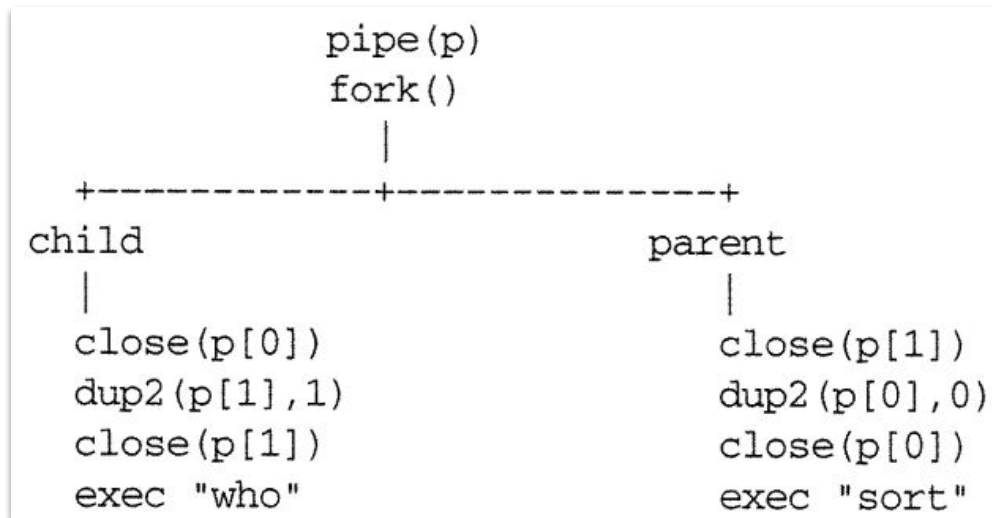
# The Finale: Combining All Skills

- Let's write a general-purpose program, called pipe.
  - It takes the names of two programs as arguments in the following:

```
pipe who sort
```

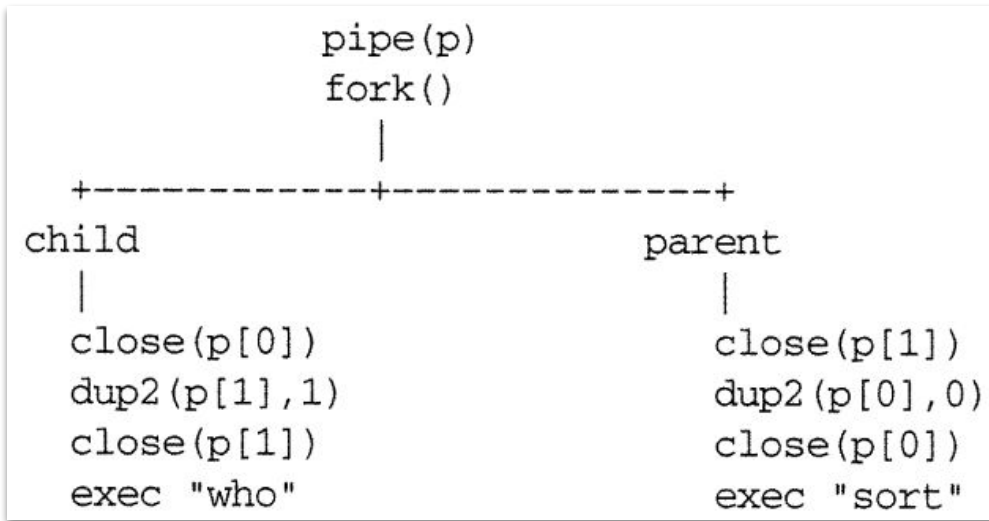
```
pipe ls head
```

- The logic of the program as follows:



# The Finale: Combining All Skills (cont.)

- pipe.c



```
/* pipe.c
 *
 * Demonstrates how to create a pipeline from one process to another
 *
 * Takes two args, each a command, and connects
 *
 * av[1]'s output to input of av[2]
 *
 * usage: pipe command1 command2
 *
 * effect: command1 | command2
 *
 * Limitations: commands do not take arguments
 *
 * uses execlp() since known number of args
 *
 * Note: exchange child and parent and watch fun
 */
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/wait.h>

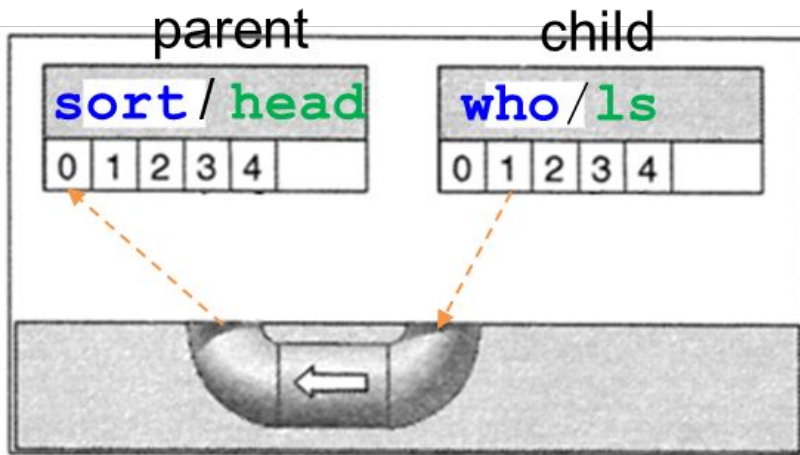
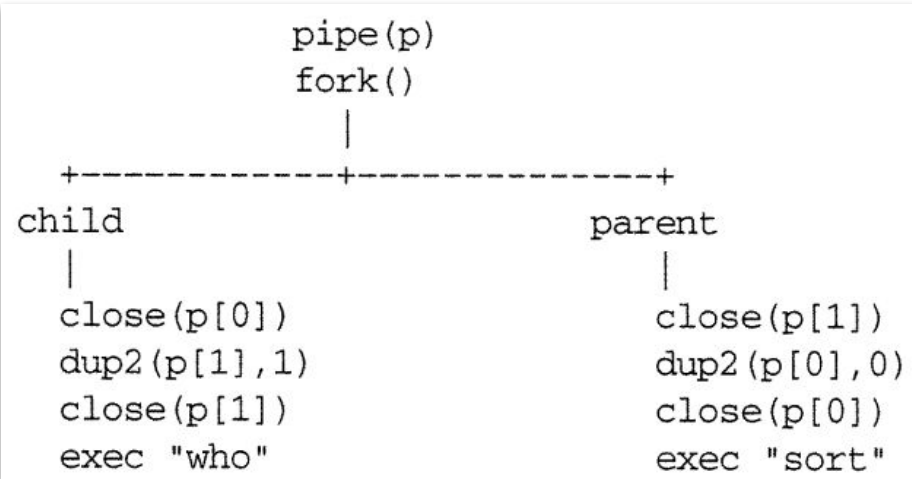
#define oops(m,x) { perror(m); exit(x); }

int main(int ac, char **av)
{
    int thepipe[2], newfd, pid; /* two file descriptors */
                                /* useful for pipes */
                                /* and the pid */

    if ( ac != 3 ){
        fprintf(stderr, "usage: pipe cmd1 cmd2\n");
        exit(1);
    }
    if ( pipe( thepipe ) == -1 ) /* get a pipe */
        oops("Cannot get a pipe", 1);
```

# The Finale: Combining All Skills (cont.)

- pipe.c (cont.)



```

/* ----- */
/*      now we have a pipe, now let's get two processes      */
/* ----- */

if ( (pid = fork()) == -1 )                               /* get a proc */
    oops("Cannot fork", 2);

/* ----- */
/*      Right Here, there are two processes                  */
/*      parent will read from pipe                           */
/* ----- */

if ( pid > 0 ){                                           /* parent will exec av[2] */
    close(thepipe[1]);                                   /* parent doesn't write to pipe */

    if ( dup2(thepipe[0], 0) == -1 )
        oops("could not redirect stdin", 3);

    close(thepipe[0]);                                   /* stdin is duped, close pipe */
    execlp( av[2], av[2], NULL);
    oops(av[2], 4);
}

/*      child execs av[1] and writes into pipe              */

close(thepipe[0]);                                       /* child doesn't read from pipe */

if ( dup2(thepipe[1], 1) == -1 )
    oops("could not redirect stdout", 4);

close(thepipe[1]);                                       /* stdout is duped, close pipe */
execlp( av[1], av[1], NULL);
oops(av[1], 5);

return 0;
}

```

# The Finale: Combining All Skills (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab11$ ./pipe ls sort
Makefile
listargs
listargs.c
pipe
pipe.c
pipedemo
pipedemo.c
pipedemo2
pipedemo2.c
sample.txt
sortfromfile
sortfromfile.c
stdinredir1
stdinredir1.c
stdinredir2
stdinredir2.c
stdinredir3.c
userlist
watch.sh
watch2.sh
whotofile
whotofile.c
whotofile2
whotofile2.c
```

# Similarities between Pipes and Files

- Pipes look like regular files
  - Use `write()` to put data into a pipe
  - Use `read()` to get the data from a pipe
  - Appears as a sequence of bytes without any particular block or record

# Differences between Pipes and Files

- Reading from Pipes

- 1. `read` on a pipe blocks

- When a process tries to `read` from a pipe, the call blocks until some bytes are written into the pipe

- 2. Reading EOF on a pipe

- When all writers close the writing end of the pipe, attempts to `read` from the pipe return 0, which means the end of file

- 3. Multiple readers can cause trouble

- A pipe is queue: first-in-first-out structure
    - When a process reads bytes from a pipe, those bytes (after reading) will be gone in the pipe
    - If two processes try to read from the same pipe, one process will get some of the bytes from the pipe, and the other process get the other bytes

# Differences between Pipes and Files (cont.)

- Writing to Pipes
  - 4. `write` to a pipe blocks until there is space
    - Pipes have a **finite capacity**, far lower than the file-size limit on disk files
    - The `write` call to a pipe will get blocked until enough space is prepared
  - 5. `write` guarantees a minimum chunk size
    - The kernel will not split up chunks of data into blocks ***no smaller than 512 bytes***
    - Linux guarantees an ***unbroken buffer size of 4K bytes*** for pipes
  - 6. `write` fails if no readers
    - If all readers have closed the reading ends of pipe, then an attempt to `write` to the pipe can lead to trouble
    - Kernel's two methods of notifying a process that `write` is no long valid:
      - 1) Sends `SIGPIPE` to that process, which will terminate
      - 2) If the kernel doesn't kill the process, then `write` returns -1 and sets `errno` to `EPIPE`



# Summary

- Input/output redirection allows separate programs to work as a team, each program a specialist
- Linux assumes that programs read input from fd 0 (`stdin`), write results to fd 1 (`stdout`), and report errors fd 2 (`stderr`)
- The log-in procedure sets up fds 0, 1, and 2
  - These connections and all open file descriptors are passed from parent to child and across the `exec` system call

# Summary (cont.)

- System calls creating fds always use the lowest-number free fd
- Redirecting std input/output/error means changing where fds 0, 1, or 2 connect
- Pipe is a data queue in the kernel with each end attached to a fd
  - `pipe` system call can create a pipe
- Both ends of a pipe are copied to a child process when the parent calls `fork`
- Pipes can only connect processes sharing a common parent