

System Programming (ELEC462)

Connecting to Processes Near and Far

Dukyun Nam
HPC Lab@KNU

Contents

- Introduction
- Four Types of Data Sources
- `bc`: a Unix Calculator
- `popen`: Making Processes Look Like Files
- Sockets: Connecting to Remote Processes
- Summary


Introduction

- Ideas and Skills
 - The client/server model
 - Using pipes for two-way communication
 - Coroutines
 - The file/process similarity
 - Sockets: Why, What, How?
 - Network services
 - Using sockets for client/server programs
- System Calls and Functions
 - `fdopen`, `popen`, `socket`
 - `bind`, `listen`, `accept`, `connect`

Four Types of Data Sources

- Unix presents one interface, even though data come from different types of sources

- (1)/(2) Disk/device files
 - Use `open` to connect
 - Use `read` and `write` to transfer data
- (3) Pipes
 - Use `pipe` to create
 - Use `fork` to share
 - Use `read` and `write` to transfer data
- (4) Sockets
 - Use `socket`, `listen`, and `connect` to connect
 - Use `read` and `write` to transfer data
 - Basis on a *client-server model*



< One interface, different sources >

bc: A Unix/Linux Basic Calculator

- bc has variables, loops, and functions
 - Can handle very long numbers
 - The trailing backslashes indicate continuation

```
dynam@DESKTOP-Q4IJB7:~/lab12$ bc
bc 1.07.1
Copyright 1991-1994, 1997, 1998, 2000, 2004, 2006, 2008, 2012-2017 Free Software Foundation, Inc.
This is free software with ABSOLUTELY NO WARRANTY.
For details type 'warranty'.
17^123
22142024630120207359320573764236957523345603216987331732240497016947\
29282299663749675090635587202539117092799463206393818799003722068558\
0536286573569713
2+2
4
2*2+2
6
2+2/2
3
```

bc: A Unix/Linux Basic Calculator (cont.)

- But `bc` is NOT a real calculator but actually runs `dc`

- `dc`: a stack-based (desktop) calculator requiring the user to

(i) enter both values and then

(ii) specify the operation


- e.g., $2 + 2 = 4$

- How `bc` works:

- Reads the expression from `stdin` and parses out the values and the operation
- Sends via a pipe the sequence commands, “2”, “2”, “+”, and “p” to `dc`
- Later reads the result through the pipe it attached to `stdout` of `dc`
- Forwards that message to the user

- How `dc` works:

- Stacks up the received two values, applies the “+” operation, and then prints to `stdout` the value on the top of the stack




< bc and dc as coroutines >

```
dynam@DESKTOP-Q4IJB7:~/lab12$ echo '2 + 2' | bc
4
dynam@DESKTOP-Q4IJB7:~/lab12$ echo '2 2 + p' | dc
4
```


Coding bc: pipe, fork, dup, exec

- Data connections in the kernel from user to bc and bc to dc
- Guidelines
 - Create `todc / fromdc`
 - Create a process *p1* to run `dc` (via `fork`)
 - `bc`: will run in the parent
 - In *p1*, redirect to `stdin` and `stdout` to the pipes, and then `exec dc`
 - In the parent (`bc`),
 - Read and parse user input,
 - Write commands to `dc`
 - Read response from `dc`
 - Send response to user




Tiny bc

- A simple version of `bc`
 - Uses `sscanf` to parse and speaks with `dc` through two pipes
- Program outline
 - a. Get two pipes
 - b. Fork (get another process)
 - c. In the child process to be `dc`, connect `stdin` and `out` to pipes then `execl dc`
 - d. The parent (`tinybc`) receives input and sends it via pipe
 - e. Then close pipe, and `dc` dies



[Remind] Using `fork` to Share a Pipe

- Sharing a pipe: Interprocess data flow
 - Parent/child can *write* bytes to the *writing end* of the pipe
 - Parent/child can *read* bytes from the *reading end* of the pipe



- A pipe is “most effective” when **one process writes** data and **the other processes reads** the data on the same host
 - Of course, processes can read and write together

Writing tiny bc

- tinybc.c

```
/** tinybc.c      * a tiny calculator that uses dc to do its work
**              * demonstrates bidirectional pipes
**              * input looks like number op number which
**              * tinybc converts into number \n number \n op \n p
**              * and passes result back to stdout
**
**
**              +-----+               +-----+
**      stdin  >0               >== pipetodc ==>   |
**              | tinybc      |               | dc - |
**      stdout <1               <== pipefromdc ==<  |
**              +-----+               +-----+
**
**              * program outline
**              a. get two pipes
**              b. fork (get another process)
**              c. in the dc-to-be process,
**                  connect stdin and out to pipes
**                  then execl dc
**              d. in the tinybc-process, no plumbing to do
**                  just talk to human via normal i/o
**                  and send stuff via pipe
**              e. then close pipe and dc dies
**              * note: does not handle multiline answers
**/
```

```
void fatal( char *mess )
{
    fprintf(stderr, "Error: %s\n", mess);
    exit(1);
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

#define oops(m,x)      { perror(m); exit(x); }

void be_dc(int *, int *);
void be_bc(int *, int *);
void fatal( char * );

int main()
{
    int      pid, todc[2], fromdc[2];      /* equipment */

    /* make two pipes */

    if ( pipe(todc) == -1 || pipe(fromdc) == -1 )
        oops("pipe failed", 1);


    /* get a process for user interface */

    if ( (pid = fork()) == -1 )
        oops("cannot fork", 2);
    if ( pid == 0 )                      /* child is dc */
        be_dc(todc, fromdc);
    else {
        be_bc(todc, fromdc);             /* parent is ui */
        wait(NULL);                      /* wait for child */
    }

    return 0;
}
```



Writing tiny bc (cont.)

- `tinybc.c` - `be_bc`




```
int todc[2], fromdc[2];
pipe(todc); pipe(fromdc);
fork()

close(todc[0])
close(fromdc[1])
readline from user
convert
write(todc[1], cmd, len)
read(fromdc[0], reply)
printf to user
```



Writing tiny bc (cont.)

- tinybc.c - be_bc



```
void be_bc(int todc[2], int fromdc[2])
/*
 * read from stdin and convert into to RPN, send down pipe
 * then read from other pipe and print to user
 * Uses fdopen() to convert a file descriptor to a stream
 */
{
    int    num1, num2;
    char    operation[BUFSIZ], message[BUFSIZ], *fgets();
    FILE    *fpout, *fpin, *fdopen();

    /* setup */
    close(todc[0]);                /* won't read from pipe to dc */
    close(fromdc[1]);              /* won't write to pipe from dc */

    fpout = fdopen( todc[1],  "w" ); /* convert file desc- */
    fpin  = fdopen( fromdc[0], "r" ); /* riptors to streams */
    if ( fpout == NULL || fpin == NULL )
        fatal("Error converting pipes to streams");


    /* main loop */
    while ( printf("tinybc: "), fgets(message, BUFSIZ, stdin) != NULL ){

        /* parse input */
        if ( sscanf(message, "%d[%-+*/^]%d", &num1, operation, &num2) != 3 ){
            printf("syntax error\n");
            continue;
        }

        if ( fprintf( fpout, "%d\n%d\n%c\np\n", num1, num2,
                     *operation ) == EOF )
            fatal("Error writing");
        fflush( fpout );
        if ( fgets( message, BUFSIZ, fpin ) == NULL )
            break;
        printf("%d %c %d = %s", num1, *operation, num2, message);
    }
    fclose(fpout);                /* close pipe */
    fclose(fpin);                 /* dc will see EOF */
}
```

Writing tiny bc (cont.)

- tinybc.c - be_dc



```
void be_dc(int in[2], int out[2])
/*
 *   set up stdin and stdout, then execl dc
 */
{
    /* setup stdin from pipein */
    if ( dup2(in[0], 0) == -1 )      /* copy read end to 0 */
        oops("dc: cannot redirect stdin", 3);
    close(in[0]);                    /* moved to fd 0 */
    close(in[1]);                    /* won't write here */

    /* setup stdout to pipeout */
    if ( dup2(out[1], 1) == -1 )     /* dupe write end to 1 */
        oops("dc: cannot redirect stdout", 4);
    close(out[1]);                    /* moved to fd 1 */
    close(out[0]);                    /* won't read from here */

    /* now execl dc with the - option */
    execlp("dc", "dc", "-", NULL );
    oops("Cannot run dc", 5);
}
```

Writing tiny bc (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab12$ ./tinybc
tinybc: 2+2
2 + 2 = 4
tinybc: 2^10
2 ^ 10 = 1024
tinybc: 123*456
123 * 456 = 56088
tinybc: 123/123
123 / 123 = 1
tinybc: dynam@DESKTOP-Q4IJB7:~/lab12$ |
```

fdopen: Making File Descriptors Look Like Files

- fdopen: a library function
 - Works like `fopen`, returning a `FILE*`
 - Takes a *file descriptor* not a filename as argument
 - Used when
 - You have a file descriptor but no filename
 - c.f., `fopen` if you know a filename
 - You want to convert the pipe connection into a `FILE*`
 - So you can use standard, buffered I/O operations
 - Notice how the `tinybc.c` code uses
 - `fprintf` and `fgets` to send data through the pipes to dc
 - Makes a remote access feel even more like a file


Lessons from bc/dc

- Client/Server Model
 - bc/dc: an example of the client/server model of a program design
- Bi-directional Communication
 - Requires one process to communicate with both stdin and stdout of another process
 - Traditionally, pipes can carry data in a unidirectional way
- Persistent service
 - Recall that on a shell each command creates a new process
 - bc keeps a single dc process running
 - bc uses that same instance of dc over and over again:
 - Sends dc commands in response to each line of user input
 - The bc/dc pair is treated as *coroutines*: c.f., subroutines applied to function calls
 - Both “continue to run,” but control passes from one to another
 - e.g., parsing and printing => bc, computing => dc

popen: Making Processes Look Like Files

- fopen: a library function
 - Opens a buffered connection to a *file*


```
FILE *fp;                                /* a pointer to a struct */
fp = fopen( "file1", "r" );              /* args are filename, connection type */
c = getc(fp);                            /* read char by char */
fgets(buf, len, fp);                     /* line by line */
fscanf(fp, "%d%d%s",&x,&y,x);            /* token by token */
fclose(fp);                             /* close when done */
```



fopen("file", "r")

- popen: a library function
 - Opens a buffered connection to a *process*

```
FILE *fp;                                /* same type of struct */
fp = popen("ls", "r");                   /* args are program name, connection type */
fgets(buf, len, fp);                     /* exactly the same functions */
pclose(fp);                             /* close when done */
```



popen("ls", "r")

What popen Does: Use Case

- Using popen to obtain a sorted list of current users
 - By the command of “who | sort”

```
/* popendemo.c
 * demonstrates how to open a program for standard i/o
 * important points:
 *      1. popen() returns a FILE *, just like fopen()
 *      2. the FILE * it returns can be read/written
 *         with all the standard functions
 *      3. you need to use pclose() when done
 */
#include <stdio.h>
#include <stdlib.h>

int main()
{
    FILE *fp;
    char buf[100];
    int i = 0;

    fp = popen( "who|sort", "r" );          /* open the command */

    while ( fgets( buf, 100, fp ) != NULL ) /* read from command */
        printf("%3d %s", i++, buf );        /* print data */

    pclose( fp );                          /* IMPORTANT! */
    return 0;
}
```

What `popen` Does: Use Case (cont.)


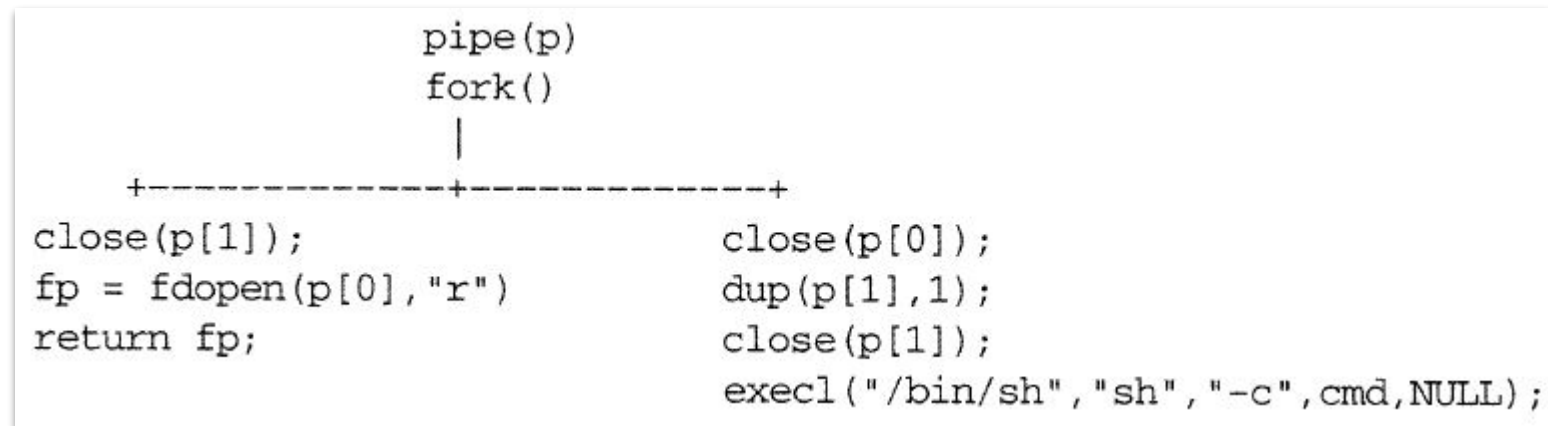
- `pclose` is required when `popen` gets invoked
 - NOT `fclose`
 - A callee process needs to be waited for
 - The callee process becomes a zombie process unless being awaited...
 - So its parent needs to retrieve its exit value
 - `pclose` calls `wait`

How Does `popen` Work? How to Write It?

- `popen`
 - Runs a program in a new process: use `fork`
 - Returns a connection to `stdin` or `stdout` of that program
 - Use `pipe`: for the connection
 - Use `dup2()`: for the redirection
 - Use `fdopen`: to make a file descriptor (fd) into a buffered stream
 - Use `exec`: to run any shell command in that process
 - `/bin/sh`: can execute any program on the shell
 - The `-c` option: tells the shell to run a command and then exit.
 - e.g., `sh -c "who | sort"`: what does this do?

How Does `popen` Work? How to Write It? (cont.)

- `popen` (Cont'd)
 - Flow chart and illustration for writing `popen`



< Reading from a shell command >

How Does `popen` Work? How to Write It? (cont.)

- `popen.c`

```
/* popen.c - a version of the Unix popen() library function
 *      FILE *popen( char *command, char *mode )
 *      command is a regular shell command
 *      mode is "r" or "w"
 *      returns a stream attached to the command, or NULL
 *      execls "sh" "-c" command
 *      todo: what about signal handling for child process?
 */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <signal.h>

#define READ 0
#define WRITE 1

int main()
{
    FILE *fp;
    char buf[BUFSIZ];

    fp = popen("ls", "r");
    while( fgets(buf, BUFSIZ, fp) != NULL)
        fputs(buf, stdout);

    return 0;
}
```

How Does `popen` Work? How to Write It? (cont.)

- `popen.c` (Cont'd)

```
FILE *popen(const char *command, const char *mode)
{
    int    pfp[2], pid;           /* the pipe and the process */
    FILE   *fdopen(), *fp;        /* fdopen makes a fd a stream */
    int     parent_end, child_end; /* of pipe */

    if ( *mode == 'r' ){          /* figure out direction */
        parent_end = READ;
        child_end = WRITE ;
    } else if ( *mode == 'w' ){
        parent_end = WRITE;
        child_end = READ ;
    } else return NULL ;

    if ( pipe(pfp) == -1 )         /* get a pipe */
        return NULL;
    if ( (pid = fork()) == -1 ){    /* and a process */
        close(pfp[0]);             /* or dispose of pipe */
        close(pfp[1]);
        return NULL;
    }
}
```

```
/* ----- parent code here ----- */
/* need to close one end and fdopen other end */

if ( pid > 0 ){
    if (close( pfp[child_end] ) == -1 )
        return NULL;
    return fdopen( pfp[parent_end] , mode); /* same mode */
}

/* ----- child code here ----- */
/* need to redirect stdin or stdout then exec the cmd */

if ( close(pfp[parent_end]) == -1 ) /* close the other end */
    exit(1);                        /* do NOT return */

if ( dup2(pfp[child_end], child_end) == -1 )
    exit(1);

if ( close(pfp[child_end]) == -1 ) /* done with this one */
    exit(1);

/* all set to run cmd */
execl( "/bin/sh", "sh", "-c", command, NULL );
exit(1);
}
```

Access to Data: Files, APIs, and Servers

- Method 1: Getting data (directly) from *files*
 - By reading from a file: who for the utmp file
 - Not a perfect solution, as a client needs to know a file format and specific names in structures

Access to Data: Files, APIs, and Servers (cont.)

- Method 2: Getting data from *functions*
 - A library function can hide all the details behind a standard function interface
 - Application programming interface (API)-based information services are not always a right solution
 - Two methods for using system library functions
 - *Static* library (or static linking) including actual function code
 - but potentially containing a bug or using out-of-date file formats
 - *Dynamic* library (or shared libraries)
 - not always installed on a system or version error

Access to Data: Files, APIs, and Servers (cont.)

- Method 3: Getting data from processes; the bc/dc example
 - May require a network connection
 - Good for a client-server model
 - Server program can be written in any language: C, C++, Java, Perl, Python ...
 - A server can be at a different machine from a client machine
 - How to connect to a process on a different machine?
 - Solution: IP address and port #
- What mechanism to allow us to connect to a process on a different computer?

Recall) What Pipes Can and Cannot Do

- Pros

- Simple, easy, less complicated, no need of network
- Allowing processes to send data to other processes as easily as they send data to files


- Cons

- Created in one process and shared by calling `fork`
 - Can only connect related processes
- Can only connect process on the SAME machine
 - What if you want to send your data to another remote host?

- Linux provides another method of IPC for remote connection: **Sockets!**

Sockets: Connecting to Remote Processes


- Allow processes to create pipeline connections to
 - Not only unrelated processes but also ones on other machines
- We'll study the basic ideas of sockets
- We'll see how to use sockets to connect clients and servers on different machines



< Connecting to a remote process >

An Analogy: “At the Tone, the Time Will Be...”

-



Four Important Concepts

- 1. Client and Server

- Server: a program (rather than a machine) that provides “services”
 - Its process waits for a request, processes that request, and then loops back to take the next request
- Client: a program (rather than a machine) that requests “services”
 - Connects to and exchanges some data with the server, and then continues (its own task) and later terminates
 - Note that it does not loop

- 2. Protocol

- The rules of interaction between the client and the server
- In the time service, the protocol is simple:
 - If the client calls, the server answers, sends the time and then hangs up

Four Important Concepts (cont.)

- 3. Hostname and port
 - Host (identified by Internet Protocol (IP) address)
 - A server on the Internet
 - A running process on a machine, or host
 - Has its assigned name (hostname) and a port number
 - These two determines a server (address), or an end of communication.
 - e.g., `cse.knu.ac.kr`:`cse` as hostname, `80`: port number (hidden)
- 4. Address family
 - A group (or, set) of different addresses for indicating a service
 - Telephone + ext. number (for telephone): maybe denoted as `AF_PHONE`
 - Street address + zip code (postal code) (for mailing): maybe denoted as `AF_MAIL`
 - Longitude + latitude (for GPS): maybe denoted as `AF_GLOBAL`
 - IP address + port number (for network connection): `AF_INET`

Lists of Services: Well-Known Ports

- 119 for emergencies, 112 for spy, 114 for phone, ...
- How can we know what services available on my machine?

```
dynam@DESKTOP-Q4IJB7:~/lab12$ more /etc/services
# Network services, Internet style
#
# Note that it is presently the policy of IANA to assign a single well-known
# port number for both TCP and UDP; hence, officially ports have two entries
# even if the protocol doesn't support UDP operations.
#
# Updated from https://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml .
#
# New ports will be added on request if they have been officially assigned
# by IANA and used in the real-world or are needed by a debian package.
# If you need a huge list of used numbers please install the nmap package.

tcpmux      1/tcp                                # TCP port service multiplexer
echo        7/tcp
echo        7/udp
discard     9/tcp      sink null
discard     9/udp      sink null
sysstat     11/tcp     users
daytime     13/tcp
daytime     13/udp
netstat     15/tcp
qotd        17/tcp      quote
chargen     19/tcp      ttytst source
chargen     19/udp      ttytst source
ftp-data    20/tcp
ftp         21/tcp
fsp         21/udp      fspd
ssh         22/tcp                                # SSH Remote Login Protocol
telnet      23/tcp
smtp        25/tcp      mail
time        37/tcp      timserver
time        37/udp      timserver
whois       43/tcp      nickname
tacacs      49/tcp                                # Login Host Protocol (TACACS)
tacacs      49/udp
domain      53/tcp                                # Domain Name Server
domain      53/udp
bootps      67/udp
bootpc      68/udp
tftp        69/udp
gopher      70/tcp                                # Internet Gopher
finger      79/tcp
http        80/tcp      www                # WorldWideWeb HTTP
kerberos    88/tcp      kerberos5 krb5 kerberos-sec # Kerberos v5
```


How Do We Write Time Server and Time Client?

- Six steps for our telephone-based time server

Action	System call
1. Get a phone line	<code>socket</code>
2. Assign a number	<code>bind</code>
3. Allow incoming calls	<code>listen</code>
4. Wait for a call	<code>accept</code>
5. Transfer data	<code>read/write</code>
6. Hang up	<code>close</code>

- Four steps for our telephone-based time client

Action	System call
1. Get a phone line	<code>socket</code>
2. Call the server	<code>connect</code>
3. Transfer data	<code>read/write</code>
4. Hang up	<code>close</code>

Working Principle of a Time Server

- Step 1: Ask kernel for a socket
 - A **socket**: a place from which calls can be made and a place to which calls can be directed
- Step 2: Bind address to a socket.
 - Address is *hostname* and *port*
- Step 3: Allow incoming calls with queue size=1 on socket
 - A server accepts incoming calls.
- Step 4: Wait for/Accept a Call.
 - Once the socket is created, assigned an address, and set up receive incoming calls, then the program is ready to go!
- Steps 5 and 6: Transfer Data and then Hang Up

Step 1: Ask kernel for a socket

- `socket` creates an endpoint (종단점) for communication and returns an identifier for that socket
 - Various sorts of communication systems, each called *domain* (e.g., Internet)
 - The *type* of a socket specifies the type of data flow
 - `SOCK_STREAM`: a bidirectional type (like TCP)
 - `SOCK_DGRAM`: connectionless (like UDP)
 - *Protocol* used within the network code in the kernel
 - c.f., `/etc/protocols`

socket		
PURPOSE	Create a socket	
INCLUDE	#include <sys/types.h> #include <sys/socket.h>	
USAGE	sockid = socket(int domain, int type, int protocol)	
ARGS	domain	communication domain. PF_INET is for Internet sockets
	type	type of socket. SOCK_STREAM looks like a pipe
	protocol	protocol used within the socket. 0 is default.
RETURNS	-1	if error
	sockid	a socket id if successful

Step 2: Bind address to a socket

- `bind` assigns a network address to a socket
 - The Internet address family (`AF_INET`) uses host and port
 - Port 13000 will be used; port 13 reserved for the *real* time server

bind		
PURPOSE	Bind an address to a socket	
INCLUDE	#include <sys/types.h> #include <sys/socket.h>	
USAGE	result = bind(int sockid, struct sockaddr *addrp, socklen_t addrlen)	
ARGS	sockid	the id of the socket
	addrp	a pointer to a struct containing the address
	addrlen	the length of the struct
RETURNS	-1	if error
	0	if success

Step 3: Allow Incoming calls with Queue size=1 on Socket

- `listen` asks the kernel to allow the specified socket to receive incoming calls.
 - Applied to `SOCK_STREAM` (not to `SOCK_DGRAM`)
 - Queue for incoming calls
 - Queue size=1 means a queue of one call
 - Maximum queue size depends on the socket implementation

listen		
PURPOSE	Listen for connections on a socket	
INCLUDE	#include <sys/socket.h>	
USAGE	result = listen(int sockid, int qsize)	
ARGS	sockid	socket that will accept calls
	qsize	backlog of incoming connections
RETURNS	-1	if error
	0	if success

Step 4: Wait for / Accept a Call

- `accept` suspends the current process until an incoming connection on the specified socket is established
 - The socket has an address, consisting of a hostname and port number
 - Returns a file descriptor (*fd*) opened for reading and writing
 - *fd*: a connection to a file descriptor in the calling process

accept		
PURPOSE	Accept a connection on a socket	
INCLUDE	#include <sys/types.h> #include <sys/socket.h>	
USAGE	fd = accept(int sockid, struct sockaddr *callerid, socklen_t *addrlenp)	
ARGS	sockid	accept a call on this socket
	callerid	pointer to struct for address of caller
	addrlenp	pointer to storage for length of address of caller
RETURNS	-1	if error
	fd	a file descriptor open for reading and writing

The Remaining Steps

- Step 5: Transfer Data
 - The `fd` returned by `accept` is a regular file descriptor.
 - Use `fdopen` to make the *fd* into a buffered data stream for `fprintf`
- Step 6: Close Connection
 - The *fd* returned by `accept` should be closed with `close`
 - When one process closes one end,
 - The other end will see EOF for a data read (as seen in pipes)

A Time Server: timeserv.c

-

Action	System call
1. Get a phone line	socket
2. Assign a number	bind
3. Allow incoming calls	listen
4. Wait for a call	accept
5. Transfer data	read/write
6. Hang up	close

```
/* timeserv.c - a socket-based time of day server
*/

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <time.h>
#include <strings.h>

#define PORTNUM 13000 /* our time service phone number */
#define HOSTLEN 256
#define oops(msg) { perror(msg) ; exit(1) ; }

int main(int ac, char *av[])
{
    struct sockaddr_in saddr; /* build our address here */
    struct hostent *hp; /* this is part of our */
    char hostname[HOSTLEN]; /* address */
    int sock_id, sock_fd; /* line id, file desc */
    FILE *sock_fp; /* use socket as stream */
    char *ctime(); /* convert secs to string */
    time_t thetime; /* the time we report */
}
```


A Time Server: `timeserv.c` (cont.)

-

Action	System call
1. Get a phone line	<code>socket</code>
2. Assign a number	<code>bind</code>
3. Allow incoming calls	<code>listen</code>
4. Wait for a call	<code>accept</code>
5. Transfer data	<code>read/write</code>
6. Hang up	<code>close</code>

```
/*
 * Step 1: ask kernel for a socket
 */

sock_id = socket( PF_INET, SOCK_STREAM, 0 );    /* get a socket */
if ( sock_id == -1 )
    oops( "socket" );

/*
 * Step 2: bind address to socket.  Address is host,port
 */

bzero( (void *)&saddr, sizeof(saddr) ); /* clear out struct */

gethostname( hostname, HOSTLEN );          /* where am I ? */
hp = gethostbyname( hostname );             /* get info about host */
                                           /* fill in host part */
bcopy( (void *)hp->h_addr, (void *)&saddr.sin_addr, hp->h_length);
saddr.sin_port = htons(PORTNUM);           /* fill in socket port */
saddr.sin_family = AF_INET ;               /* fill in addr family */

if ( bind(sock_id, (struct sockaddr *)&saddr, sizeof(saddr)) != 0 )
    oops( "bind" );

/*
 * Step 3: allow incoming calls with Qsize=1 on socket
 */

if ( listen(sock_id, 1) != 0 )
    oops( "listen" );
```

A Time Server: `timeserv.c` (cont.)

Action	System call
1. Get a phone line	<code>socket</code>
2. Assign a number	<code>bind</code>
3. Allow incoming calls	<code>listen</code>
4. Wait for a call	<code>accept</code>
5. Transfer data	<code>read/write</code>
6. Hang up	<code>close</code>

```
/*
 * main loop: accept(), write(), close()
 */

while ( 1 ){
    sock_fd = accept(sock_id, NULL, NULL); /* wait for call */
    printf("Wow! got a call!\n");
    if ( sock_fd == -1 )
        oops( "accept" );                /* error getting calls */

    sock_fp = fdopen(sock_fd, "w"); /* we'll write to the */
    if ( sock_fp == NULL )          /* socket as a stream */
        oops( "fdopen" );          /* unless we can't */

    thetime = time(NULL);           /* get time */
                                    /* and convert to string */
    fprintf( sock_fp, "The time here is .." );
    fprintf( sock_fp, "%s", ctime(&thetime) );
    fclose( sock_fp );              /* release connection */
}
}
```

Working Principle of a Time Client

- Step 1: Ask Kernel for a Socket
 - Needs a socket to connect to the network
 - Like a client needing a phone line in the phone network
- Step 2: Connect to Server
 - Uses the `connect` system call
- Step 3: Transfer Data
 - Reads one line from the server through the connected socket
- Step 4: Hang Up
 - `closes` the file descriptor for the connected socket and exits

More Details about Step 2

- connect attempts to connect the socket specified by `sockid` to the socket address pointed by `serv_addrp`
 - If the attempt succeeds, `result` will get 0.
 - `sockid` now then becomes a fd open for reading and writing
 - Data written into are sent to and read from the socket's ends

connect		
PURPOSE	Connect to a socket	
INCLUDE	#include <sys/types.h> #include <sys/socket.h>	
USAGE	result = connect(int sockid, struct sockaddr *serv_addrp, socklen_t addrlen);	
ARGS	sockid	socket to use for connection
	serv_addrp	pointer to struct containing server address
	addrlen	length of that struct
RETURNS	-1	if error
	0	if success

A Time Client: `timeclnt.c`

Action	System call
1. Get a phone line	<code>socket</code>
2. Call the server	<code>connect</code>
3. Transfer data	<code>read/write</code>
4. Hang up	<code>close</code>

```
/* timeclnt.c - a client for timeserv.c
 *          usage: timeclnt hostname portnumber
 */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <strings.h>

#define oops(msg) { perror(msg); exit(1); }

int main(int ac, char *av[])
{
    struct sockaddr_in servadd; /* the number to call */
    struct hostent *hp;        /* used to get number */
    int sock_id, sock_fd;      /* the socket and fd */
    char message[BUFSIZ];      /* to receive message */
    int messlen;               /* for message length */
```

```
/*
 * Step 1: Get a socket
 */

    sock_id = socket( AF_INET, SOCK_STREAM, 0 ); /* get a line */
    if ( sock_id == -1 )
        oops( "socket" ); /* or fail */

/*
 * Step 2: connect to server
 *          need to build address (host,port) of server first
 */

    bzero( &servadd, sizeof( servadd ) ); /* zero the address */

    hp = gethostbyname( av[1] ); /* lookup host's ip # */
    if ( hp == NULL )
        oops( av[1] ); /* or die */
    bcopy( hp->h_addr, (struct sockaddr *)&servadd.sin_addr, hp->h_length);

    servadd.sin_port = htons(atoi(av[2])); /* fill in port number */
    servadd.sin_family = AF_INET; /* fill in socket type */

    /* now dial */
    if ( connect( sock_id, (struct sockaddr *)&servadd, sizeof( servadd ) ) != 0 )
        oops( "connect" );


/*
 * Step 3: transfer data from server, then hangup
 */

    messlen = read( sock_id, message, BUFSIZ ); /* read stuff */
    if ( messlen == -1 )
        oops( "read" );
    if ( write( 1, message, messlen ) != messlen ) /* and write to */
        oops( "write" ); /* stdout */
    close( sock_id );

    return 0;
}
```


Testing timeserv and timeclnt

- The server process runs on one machine
- A client process on another machine connects to the server over the network
- Then the server sends data to client by write
- The client receives that message by read



< Processes on different computers >

Testing timeserv and timeclnt (cont.)

- Execution

- Server

- \$ `./timeserv &`

```
dynam@DESKTOP-Q4IJB7:~/lab12$ ./timeserv &  
[1] 53  
dynam@DESKTOP-Q4IJB7:~/lab12$ Wow! got a call!
```

- Client

- \$ `./timeclnt localhost 13000`

```
dynam@DESKTOP-Q4IJB7:~/lab12$ ./timeclnt DESKTOP-Q4IJB7.localdomain 13000  
The time here is ..Thu Nov 24 02:32:57 2022  
dynam@DESKTOP-Q4IJB7:~/lab12$ |
```


Example of Another Server and Client: Remote `ls`

- Listing files on a remote computer
 - Could log in to that machine and run `ls`
 - e.g., `$./r1s 155.230.100.100 /home/username`
 - `r1s` needs a server process running on the other machine
 - To receive the request, do the work, and return the result
 - The server runs on one computer
 - A client on another computer connects to the server
 - Sends the name of a requesting directory: e.g., `'/home/username'`
 - The server sends back to the client a list of the files in that requested directory
 - The client displays the list by writing to `stdout`
- This two-process system really provides access to directories on a different machine!

Remote ls

- Three things to implement the remote ls system

- 1) A protocol


- Consists a request and a reply between a client and a server program

- 2) A client program

- Sends a single-line containing the name of a requested directory
- Reads the list of files line by line until the server closes the connection

- 3) A server program

- The server opens and reads that directory and sends back to the client the list of files
- When closing the connection, it generates an EOF condition



< A remote ls system >

Remote `ls` (cont.)

- The client: `rls`
 - Differences of this client from the time client
 - 1. Writing the directory name into the socket
 - 2. Entering a loop, copying data from the socket to `stdout` until EOF
 - The loop uses a standard buffer size for efficiency
 - 3. Using `write` and `read` for data transfer (with the server)


Remote ls (cont.)

- The server: `r1sd`
 - Has to get a socket
 - Bind and listen, and then accept a call.
 - Read then the name of a requested directory from the socket
 - Lists the contents of that directory
 - Use `popen` to read the output from the regular version of `ls`
- Notes
 - The server uses standard buffered streams for reading and writing
 - Use `fgets` to read the directory name from the client
 - After `popen`, it transfers data using `getc/putc` for a file copy
 - Actually copying data from one process to a process on another machine

Remote ls (cont.)

- Additional notes

- The string the program receives
 - Does not overflow the input buffer
 - Does not overflow the buffer for the command
 - Doesn't allow special characters in the directory name
- `popen`: indeed too risky for network services
 - Because it passes a string to a shell
 - It's a *poor* idea to write any server-passing strings to a shell!
 - Two reasons of why to use this example
 - For showing another use of `popen`
 - For alerting you guys to this danger



< Using `popen("ls")` to list remote directories >

Writing remote ls: rlsd.c

```
/* rlsd.c - a remote ls server
 */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <time.h>
#include <ctype.h>
#include <strings.h>

#define PORTNUM 15000 /* our remote ls server port */
#define HOSTLEN 256
#define oops(msg) { perror(msg) ; exit(1) ; }

void sanitize(char * );

int main(int ac, char *av[])
{
    struct sockaddr_in saddr; /* build our address here */
    struct hostent *hp; /* this is part of our */
    char hostname[HOSTLEN]; /* address */
    int sock_id, sock_fd; /* line id, file desc */
    FILE *sock_fpi, *sock_fpo; /* streams for in and out */
    FILE *pipe_fp; /* use popen to run ls */
    char dirname[BUFSIZ]; /* from client */
    char command[BUFSIZ]; /* for popen() */
    int dirlen, c;
```

```
/** Step 1: ask kernel for a socket **/
```

```
    sock_id = socket( PF_INET, SOCK_STREAM, 0 ); /* get a socket */
    if ( sock_id == -1 )
        oops( "socket" );
```

```
/** Step 2: bind address to socket. Address is host,port **/
```

```
    bzero( (void *)&saddr, sizeof(saddr) ); /* clear out struct */
    gethostname( hostname, HOSTLEN ); /* where am I ? */
    hp = gethostbyname( hostname ); /* get info about host */
    bcopy( (void *)hp->h_addr, (void *)&saddr.sin_addr, hp->h_length);
    saddr.sin_port = htons(PORTNUM); /* fill in socket port */
    saddr.sin_family = AF_INET; /* fill in addr family */
    if ( bind(sock_id, (struct sockaddr *)&saddr, sizeof(saddr)) != 0 )
        oops( "bind" );
```

```
/** Step 3: allow incoming calls with Qsize=1 on socket **/
```

```
    if ( listen(sock_id, 1) != 0 )
        oops( "listen" );
```


Writing remote ls: rlsd.c (cont.)

```
/*
 * main loop: accept(), write(), close()
 */

while ( 1 ){
    sock_fd = accept(sock_id, NULL, NULL); /* wait for call */
    if ( sock_fd == -1 )
        oops("accept");

    /* open reading direction as buffered stream */
    if( (sock_fpi = fdopen(sock_fd, "r")) == NULL )
        oops("fdopen reading");

    if ( fgets(dirname, BUFSIZ-5, sock_fpi) == NULL )
        oops("reading dirname");
    sanitize(dirname);

    /* open writing direction as buffered stream */
    if ( (sock_fpo = fdopen(sock_fd, "w")) == NULL )
        oops("fdopen writing");

    sprintf(command, "ls %s", dirname);
    if ( (pipe_fp = popen(command, "r")) == NULL )
        oops("popen");

    /* transfer data from ls to socket */
    while( (c = getc(pipe_fp)) != EOF )
        putc(c, sock_fpo);
    pclose(pipe_fp);
    fclose(sock_fpo);
    fclose(sock_fpi);
}

return 0;
}
```

```
void sanitize(char *str)
/*
 * it would be very bad if someone passed us a dirname like
 * "; rm *" and we naively created a command "ls ; rm *"
 *
 * so..we remove everything but slashes and alphanumerics
 * There are nicer solutions, see exercises
 */
{
    char *src, *dest;

    for ( src = dest = str ; *src ; src++ )
        if ( *src == '/' || isalnum(*src) )
            *dest++ = *src;

    *dest = '\0';
}
```


Writing remote ls: rls.c

```
/* rls.c - a client for a remote directory listing service
 *      usage: rls hostname directory
 */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <string.h>

#define oops(msg) { perror(msg); exit(1); }
#define PORTNUM 15000

int main(int ac, char *av[])
{
    struct sockaddr_in servadd; /* the number to call */
    struct hostent *hp; /* used to get number */
    int sock_id, sock_fd; /* the socket and fd */
    char buffer[BUFSIZ]; /* to receive message */
    int n_read; /* for message length */

    if ( ac != 3 ) exit(1);

    /** Step 1: Get a socket **/

    sock_id = socket( AF_INET, SOCK_STREAM, 0 ); /* get a line */
    if ( sock_id == -1 )
        oops( "socket" ); /* or fail */
```

```
/** Step 2: connect to server **/

    bzero( &servadd, sizeof(servadd) ); /* zero the address */
    hp = gethostbyname( av[1] ); /* lookup host's ip # */
    if (hp == NULL)
        oops(av[1]); /* or die */
    bcopy(hp->h_addr, (struct sockaddr *)&servadd.sin_addr, hp->h_length);
    servadd.sin_port = htons(PORTNUM); /* fill in port number */
    servadd.sin_family = AF_INET; /* fill in socket type */

    if ( connect(sock_id, (struct sockaddr *)&servadd, sizeof(servadd)) != 0 )
        oops( "connect" );

    /** Step 3: send directory name, then read back results **/

    if ( write(sock_id, av[2], strlen(av[2])) == -1 )
        oops("write");
    if ( write(sock_id, "\n", 1) == -1 )
        oops("write");

    while( (n_read = read(sock_id, buffer, BUFSIZ)) > 0 )
        if ( write(1, buffer, n_read) == -1 )
            oops("write");
    close( sock_id );

    return 0;
}
```

Execution of Remote ls

```
dynam@DESKTOP-Q4IJB7:~/lab12$ ./rlsd &  
[1] 193  
dynam@DESKTOP-Q4IJB7:~/lab12$ ./rls DESKTOP-Q4IJB7.localdomain /home/dynam  
GoogleDrive  
a.out  
demodir  
elec462
```

Software Daemons

- Many Linux server programs ending in the letter 'd'
 - e.g., httpd, inetd, syslogd, atd, ntpd, sshd, ...
 - 'd' stands for daemon
 - Daemon: a supernatural helper floating around waiting you to help out
 - Provides a variety of services and performs system maintenance
 - Alerting, flushing, logging, printing, accepting network connections, ...
 - Most daemon processes get started at the boot-up time
 - /etc/rcX.d, where X depends on system
 - Starts these servers in the background for providing services, with being detached from any terminals

Summary

- Some programs are written as separate processes for data transfer
 - A server process responsible for processing or data delivery in the CS model
- A CS system consists of a communication system and a protocol
 - Protocol: a set of rules for the structure of a conversation
 - Clients/servers can communicate through pipes or sockets
- `popen` can make any shell command into a server program.
 - Makes access to the server look like having access to buffered files

Summary (cont.)

- A pipe: a connected pair of fds
 - Socket: an unconnected communication endpoint (potential fd)
 - A client creates a comm. link by connecting its socket to a server socket
- Connections between sockets from one machine to another.
 - Each socket is identified by an *IP address* and a *port number*
- Connections to pipes and sockets use fds
 - Fds provide programs with a single interface for communication with different objects: files, devices, and other processes