

System Programming

(ELEC462)

Threads

Dukyun Nam
HPC Lab@KNU

Contents

- Introduction
- Threads of Execution
- Interthread Cooperation
- Comparing Threads with Processes
- Interthread Notification
- Summary

Introduction

- Ideas and skills
 - Threads of execution
 - Multithreaded programs
 - Creating and destroying threads
 - Passing multiple arguments to a thread
 - Sharing data between threads safely using “mutex locks”
 - Synchronizing data transfer using condition variables among threads
- System calls and functions
 - `pthread_create`, `pthread_join`
 - `pthread_mutex_lock`, `pthread_mutex_unlock`
 - `pthread_cond_wait`, `pthread_cond_signal`

Threads of Execution

- A toy C program: `hello_single.c`
 - Write `print_msg()` to print a given string (e.g., `m`) and then sleep for 1 sec as many times as specified by `NUM`
 - This program prints “Hello” and then “world\n” and sleeps for 1 second, which repeats five times

```
dynam@DESKTOP-Q4IJB7:~/lab13$ ./hello_single
HelloHelloHelloHelloHelloworld
world
world
world
world
```

```
#include <stdio.h>
#include <unistd.h>
#define NUM 5

int main()
{
    void print_msg(char *);


    print_msg("Hello");
    print_msg("world\n");

    return 0;
}

void print_msg(char *m)
{
    int i;
    for(i=0 ; i<NUM ; i++){
        printf("%s", m);
        fflush(stdout);
        sleep(1);
    }
}
```

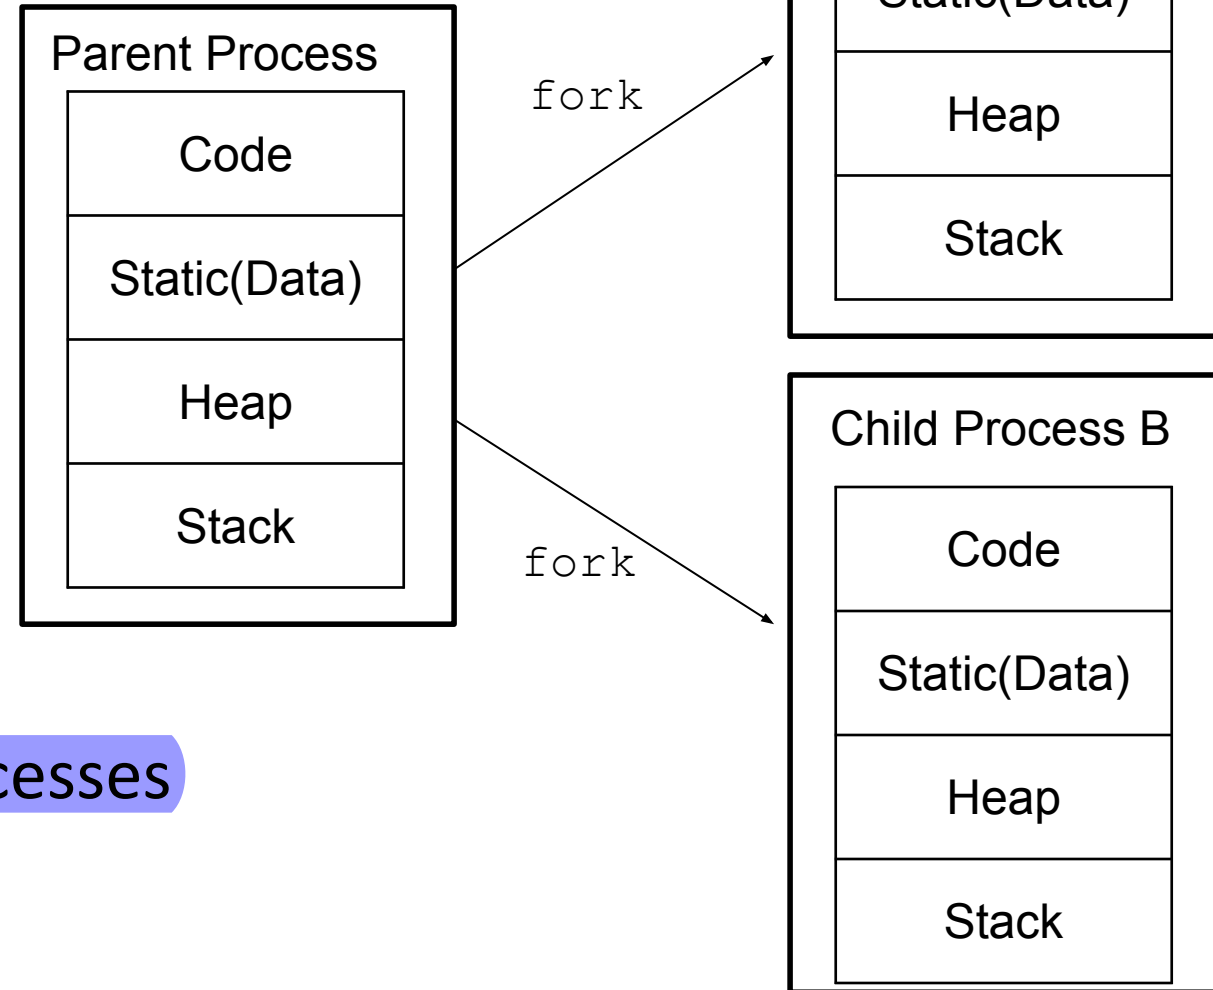
Threads of Execution (cont.)

- An Execution of the Program: a Process
 - Thread of execution
 - “an unbroken path” tracing the order of execution of instructions
 - Control starts in main, then flows into repeat two times, then exits from the end of main



Process (a single-thread program)

- Gets a certain memory address space
 - Each (same) running process has its own memory space
- The parent and the two children have their own memory realms
- Environment variables can be inherited from the parent
- “ps” will tell you those three processes running independently

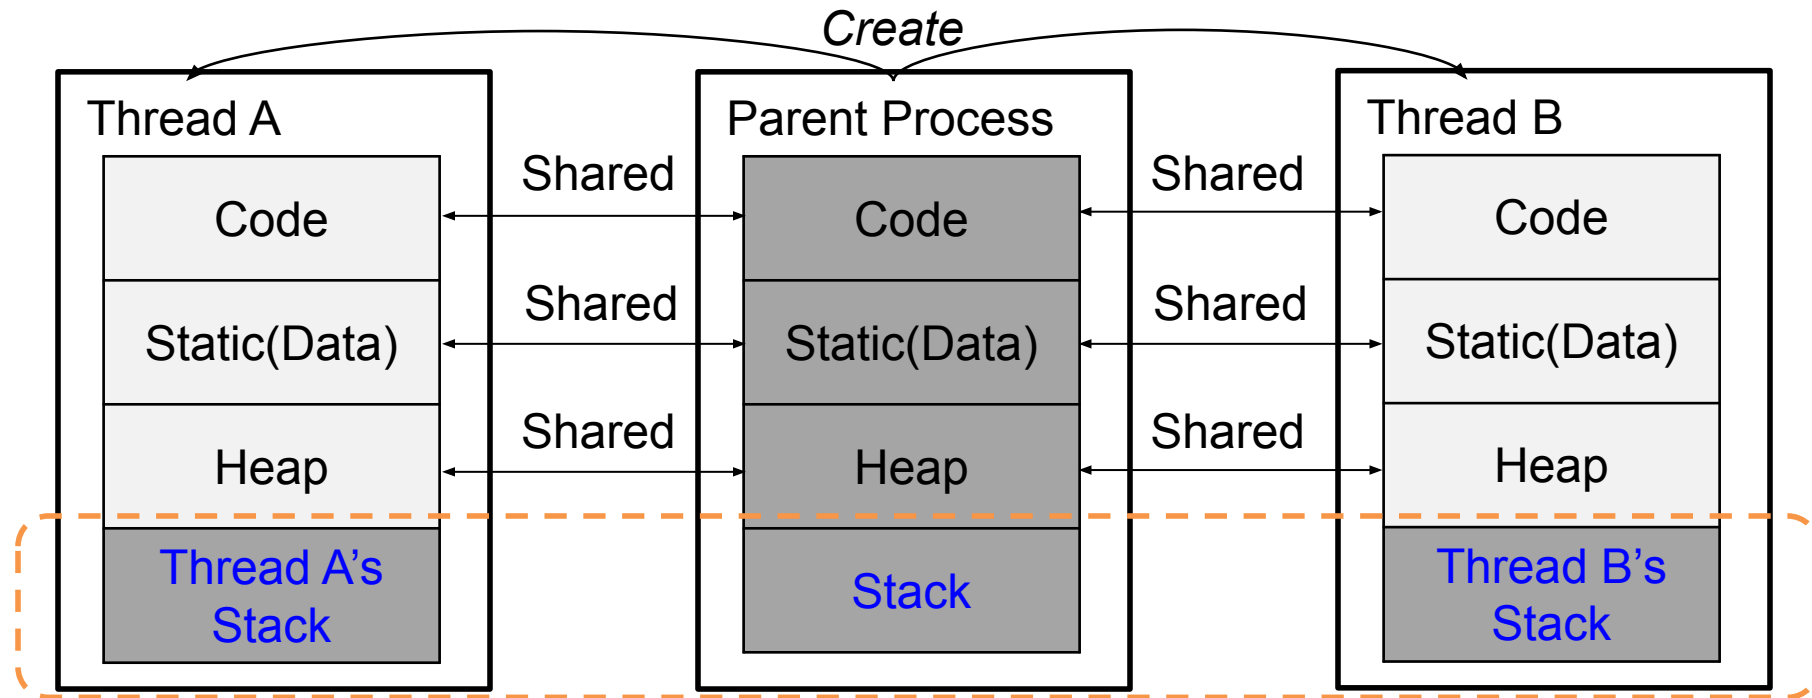


Thread(s)

- Definition: an (multiple) execution(s) within a *single* process
 - Functions what a process is to programs, an environment in which to run.
 - Sometimes referred as a *lightweight* process (LWP)
 - The “smallest unit” of processing that a kernel works on
- Shares the same memory address space as its parent (process)
 - Shares all global variables and file descriptors of the parent process
- Allows the programmer to separate multiple tasks easily in a single process

Thread(s) (cont.)

- Parent and two children




Why Thread(s)?

- Requires “fewer system resources” compared to forking
 - No extra memory space but stack
- Consumes less process time to create and to terminate than a process (compared to `fork`)
- Enables “resource” sharing
 - Some memory realms shared: code, data, heap, static
- Less communication overhead
 - Neither pipe nor socket needed!
- There are also some shortcomings ...
 - We’ll get to discuss similarities and differences between processes and threads in later slides

A Multithreaded Program

- If we want to run `print_msg` concurrently, how?
 - A process may create a new thread by specifying a function to run and an argument



< Multiple threads of execution >

```
int main()
{
    pthread_t t1, t2;           /* two threads */

    void *print_msg(void *);

    pthread_create(&t1, NULL, print_msg, (void *)"hello");
    pthread_create(&t2, NULL, print_msg, (void *)"world\n");
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);

    return 0;
}
```

Relevant Function 1: pthread_create

pthread_create

PURPOSE Create a new thread

INCLUDE #include <pthread.h>

USAGE int pthread_create(pthread_t *thread,
 pthread_attr_t *attr,
 void *(*func)(void *),
 void *arg);

ARGS thread a pointer to a variable of type pthread_t
 attr a pointer to a variable of type pthread_attr_t
 or NULL.
 func the function this new thread will run
 arg the argument to be passed to func

RETURNS 0 if successful
 errcode if not successful

```
int main()
{
    pthread_t t1, t2;           /* two threads */

    void *print_msg(void *);

    pthread_create(&t1, NULL, print_msg, (void *)"hello");
    pthread_create(&t2, NULL, print_msg, (void *)"world\n");
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);

    return 0;
}
```

Relevant Function 2: pthread_join

| pthread_join | | |
|--------------|---|---|
| PURPOSE | Wait for termination of a thread | |
| INCLUDE | #include <pthread.h> | |
| USAGE | int pthread_join(pthread_t thread, void **retval) | |
| ARGS | thread retval | the thread to wait for points to a variable to receive the return value from the thread |
| RETURNS | 0 errcode | if thread terminates if an error |

```
int main()
{
    pthread_t t1, t2;           /* two threads */

    void *print_msg(void *);

    pthread_create(&t1, NULL, print_msg, (void *)"hello");
    pthread_create(&t2, NULL, print_msg, (void *)"world\n");
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);

    return 0;
}
```

A Multithreaded Program: hello_multi.c

```
/* hello_multi.c - a multi-threaded hello world program */

#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

#define NUM      5

int main()
{
    pthread_t t1, t2;                /* two threads */

    void *print_msg(void *);

    pthread_create(&t1, NULL, print_msg, (void *)"hello");
    pthread_create(&t2, NULL, print_msg, (void *)"world\n");
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);

    return 0;
}
```

```
void *print_msg(void *m)
{
    char *cp = (char *) m;
    int i;
    for(i=0 ; i<NUM ; i++){
        printf("%s", cp);
        fflush(stdout);
        sleep(1);
    }
    return NULL;
}
```

A Multithreaded Program: `hello_multi.c` (cont.)

- Execution
 - `-lpthread`: tells `gcc` to (dynamically) link the pthread library to its source code (`-llibname`: link a library named *libname*)


```
dynam@DESKTOP-Q4IJB7:~/lab13$ make
cc -o hello_multi hello_multi.c -lpthread
dynam@DESKTOP-Q4IJB7:~/lab13$ ./hello_multi
helloworld
helloworld
world
helloworld
helloworld
helloworld
dynam@DESKTOP-Q4IJB7:~/lab13$ |
```

Interthread Communication

- Threads execute *functions* in a single process
- They share “**global variables,**” or their communication channel
 - c.f., Processes communicate with each other using pipes, sockets, signals, exit/wait, and the environment
- Simultaneous access to memory: a powerful but risky feature
- Two big questions for threads and data
 - 1. How can threads share variables safely?
 - 2. How can threads transfer data effectively?

Example 1: `incrprint.c`

- All threads in a process share the same set of functions and global variables
 - Two “executions” (main process and one thread) share a global variable



< Two threads share a global variable >

Example 1: `incrprint.c` (cont.)

- In this program, one thread increments a variable, and the other thread prints the value

```
void *print_count(void *m)
{
    int i;
    for(i=0 ; i<NUM ; i++){
        printf("count = %d\n", counter);
        sleep(1);
    }
    return NULL;
}
```

```
/* incrprint.c - one thread increments, the other prints */

#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

#define NUM 5

int counter = 0;

int main()
{
    pthread_t t1; /* one thread */
    void *print_count(void *); /* its function */
    int i;

    pthread_create(&t1, NULL, print_count, NULL);
    for( i = 0 ; i<NUM ; i++ ){
        counter++;
        sleep(1);
    }

    pthread_join(t1, NULL);

    return 0;
}
```

Example 1: `incrprint.c` (cont.)

- When `main` changes the value of `counter`, `print_counter` sees that new value at once
 - There is no need to send the new value through a pipe or socket

```
dynam@DESKTOP-Q4IJB7:~/lab13$ make
cc -o incrprint incrprint.c -lpthread
dynam@DESKTOP-Q4IJB7:~/lab13$ ./incrprint
count = 1
count = 2
count = 3
count = 4
count = 5
```

- A *single-threaded* program works! – Not practical though...
 - Many concurrent programs runs multiple threads in parallel!

A Target Program for Multi-Threading: `wc`


- Counts lines, words, and chars in one or more files in a *single-threaded*

```
dynam@DESKTOP-Q4IJB7:~/lab13$ wc incprint.c
 36  82 545 incprint.c
dynam@DESKTOP-Q4IJB7:~/lab13$ wc -l incprint.c
36 incprint.c
dynam@DESKTOP-Q4IJB7:~/lab13$ wc -w incprint.c
82 incprint.c
dynam@DESKTOP-Q4IJB7:~/lab13$ wc -c incprint.c
545 incprint.c
```

- Now let's make the `wc` program be *threaded*

Example 2-1: twordcount1.c

- Version 1: “Two Threads with **One Counter**”
 - This real “**multi-threaded**” program counts the number of words in two files and prints the number of words
 - We create a separate thread to count each file
 - Both threads increment the counter while detecting words



< A common counter for two threads >

Example 2-1: twordcount1.c (cont.)

```
/* twordcount1.c - threaded word counter for two files. Version 1 */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <ctype.h>

int      total_words ;

int main(int ac, char *av[])
{
    pthread_t t1, t2;          /* two threads */
    void      *count_words(void *);

    if ( ac != 3 ){
        printf("usage: %s file1 file2\n", av[0]);
        exit(1);
    }
    total_words = 0;
    pthread_create(&t1, NULL, count_words, (void *) av[1]);
    pthread_create(&t2, NULL, count_words, (void *) av[2]);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("%5d: total words\n", total_words);

    return 0;
}
```

```
void *count_words(void *f)
{
    char *filename = (char *) f;
    FILE *fp;
    int  c, prevc = '\0';

    if ( (fp = fopen(filename, "r")) != NULL ){
        while( ( c = getc(fp)) != EOF ){
            if ( !isalnum(c) && isalnum(prevc) )
                total_words++;
            prevc = c;
        }
        fclose(fp);
    } else
        perror(filename);
    return NULL;
}
```

Example 2-1: twordcount1.c (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
132005: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /dev/null
249: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /usr/share/dict/words /dev/null
131756: total words
```


- What happens if both threads read then update a variable at the same time?

- Answer) Unpredictable confusion

```
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
131977: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
131969: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
131955: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
131966: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
132005: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
131966: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc1 /etc/group /usr/share/dict/words
131950: total words
```


Example 2-1: twordcount1.c (cont.)


- **Drawback** of this program is ...
 - Two increments take place although we have only “one counter” ...
 - What happens if both threads increment the counter using **fetch-add-store** sequence at the **same** time?
 - In this case, both threads were not well synchronized on the same resource



< Two threads increment the same counter >

Example 2-2: twordcount2.c

- Version 2: “Two Threads, One Counter, **One Mutex**”
 - A mutex is a special kind of object that can be ‘locked’ and ‘unlocked’
 - Only one thread may lock a mutex at a time
 - Attempts to lock a mutex block until the mutex is unlocked



< Two threads use a mutex to share a counter >

Example 2-2: twordcount2.c (cont.)

- The threads system includes variables, called *mutual exclusion locks* (abbr. **mutex**)
 - Used by cooperating threads to prevent simultaneous access to variable, function, and resources
 - e.g., threads A and B
 - If thread A gets the lock first, thread B's call to `pthread_mutex_lock` waits until thread A calls `pthread_mutex_unlock`
 - This prevents simultaneous access to the shared counter

```
if ( !isalnum(c) && isalnum(prevc) ){  
    pthread_mutex_lock(&counter_lock);  
    total_words++;  
    pthread_mutex_unlock(&counter_lock);  
}
```

Example 2-2: twordcount2.c (cont.)

- Function Summaries: pthread_mutex_lock

| pthread_mutex_lock | | |
|--------------------|--|--|
| PURPOSE | Wait for and lock a mutex | |
| INCLUDE | #include <pthread.h> | |
| USAGE | int pthread_mutex_lock(pthread_mutex_t *mutex) | |
| ARGS | mutex | a pointer to a mutual exclusion object |
| RETURNS | 0 | for success |
| | errcode | for errors |

Example 2-2: twordcount2.c (cont.)

- Function Summaries: pthread_mutex_unlock

| pthread_mutex_unlock | | |
|----------------------|--|--|
| PURPOSE | Unlock a mutex | |
| INCLUDE | #include <pthread.h> | |
| USAGE | int pthread_mutex_unlock(pthread_mutex_t *mutex) | |
| ARGS | mutex | a pointer to a mutual exclusion object |
| RETURNS | 0 | for success |
| | errcode | for errors |

Example 2-2: twordcount2.c (cont.)

```
/* twordcount2.c - threaded word counter for two files.    */
/*                                     version 2: uses mutex to lock counter    */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <ctype.h>

int      total_words ;    /* the counter and its lock */
pthread_mutex_t counter_lock = PTHREAD_MUTEX_INITIALIZER;

int main(int ac, char *av[])
{
    pthread_t t1, t2;      /* two threads */
    void      *count_words(void *);

    if ( ac != 3 ){
        printf("usage: %s file1 file2\n", av[0]);
        exit(1);
    }
    total_words = 0;
    pthread_create(&t1, NULL, count_words, (void *) av[1]);
    pthread_create(&t2, NULL, count_words, (void *) av[2]);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("%5d: total words\n", total_words);

    return 0;
}
```

```
void *count_words(void *f)
{
    char *filename = (char *) f;
    FILE *fp;
    int c, prevc = '\0';

    if ( (fp = fopen(filename, "r")) != NULL ){
        while( ( c = getc(fp)) != EOF ){
            if ( !isalnum(c) && !isalnum(prevc) ){
                pthread_mutex_lock(&counter_lock);
                total_words++;
                pthread_mutex_unlock(&counter_lock);
            }
            prevc = c;
        }
        fclose(fp);
    } else
        perror(filename);
    return NULL;
}
```

Example 2-2: twordcount2.c (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc2 /etc/group /usr/share/dict/words
132005: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc2 /etc/group /usr/share/dict/words
132005: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc2 /etc/group /usr/share/dict/words
132005: total words
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc2 /etc/group /usr/share/dict/words
132005: total words
```

Is Mutex an *Ultimate* Solution?

- Mutex allows concurrent threads to be synchronized with a “lock”
 - Offers no inconsistent state of a variable
- But ... using a mutex makes the program *slower*
 - *A lot of operations* are added up due to:
 - Checking the lock, setting the lock, and releasing the lock for every word in both files
- What is a “more efficient” solution?

Example 2-3: twordcount3.c


- Version 3: “Two Threads, **Two Counters, Multiple Arguments to Threads**”
 - We’ll eliminate the need for a mutex
 - We’ll instead give each thread its own counter, which sounds more efficient
- Questions:
 - How do we get the threads to acquire those counters?
 - How do the threads pass their counts back?

Example 2-3: `twordcount3.c` (cont.)

- Our approach:
 - Makes the calling thread be able to pass to the function a ***pointer*** to a variable
 - Issue: `pthread_create` takes one argument, thereby no more room.
 - Solution: Let's use a (local) **struct** with two members of (i) *filename* and (ii) *number of words* and then have the struct handed over to each thread
 - Let the function ***increment*** that variable

Example 2-3: twordcount3.c (cont.)

- Two arguments (defined in a struct) are passed to the threads
 - The struct contains (i) file name and (ii) number of words.
 - Passing pointers to local structs
- This will remove the use of mutex and global variables
- Note
 - The argument to each thread must include a filename and a pointer to a counter



< Each thread has a pointer to its own struct >

Example 2-3:

twordcount3.c (cont.)

```
/* twordcount3.c - threaded word counter for two files.
 *                - Version 3: one counter per file
 */

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <ctype.h>

struct arg_set {
    char *fname;    /* two values in one arg */
    int count;      /* file to examine */
                  /* number of words */
};
```

```
int main(int ac, char *av[])
{
    pthread_t    t1, t2;        /* two threads */
    struct arg_set args1, args2; /* two argsets */
    void         *count_words(void *);

    if ( ac != 3 ){
        printf("usage: %s file1 file2\n", av[0]);
        exit(1);
    }
    args1.fname = av[1];
    args1.count = 0;
    pthread_create(&t1, NULL, count_words, (void *) &args1);

    args2.fname = av[2];
    args2.count = 0;
    pthread_create(&t2, NULL, count_words, (void *) &args2);

    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("%5d: %s\n", args1.count, av[1]);
    printf("%5d: %s\n", args2.count, av[2]);
    printf("%5d: total words\n", args1.count+args2.count);

    return 0;
}
```

Example 2-3: twordcount3.c (cont.)

```
void *count_words(void *a)
{
    struct arg_set *args = a;      /* cast arg back to correct type */
    FILE *fp;
    int c, prevc = '\0';

    if ( (fp = fopen(args->fname, "r")) != NULL ){
        while( ( c = getc(fp)) != EOF ){
            if ( !isalnum(c) && isalnum(prevc) )
                args->count++;
            prevc = c;
        }
        fclose(fp);
    } else
        perror(args->fname);
    return NULL;
}
```

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab13$ make
cc -o twc3 twordcount3.c -lpthread
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc3 /etc/group /usr/share/dict/words
249: /etc/group
131756: /usr/share/dict/words
132005: total words
```

Interthread Cooperation: Summary

- Three approaches to sharing values between threads
 - 1. Allowing threads to modify the same variable with no cooperation
 - `twordcount1.c`: works most of the time, which is NOT correct.
 - 2. Using a mutual exclusion object, a *mutex*,
 - To allow “only one thread” to increment the shared counter at a time
 - `twordcount2.c`: works most of the time
 - But makes many calls to the functions to check, set, and release the lock
 - 3. Creating separate counters for each thread while eliminating a shared counter
 - `twordcount3.c`: uses `pthread_join` to block until all threads finish
 - The threads no longer share a variable; then no need of cooperation
 - But can still cooperate with the original thread

Interthread Notification

- How can one thread notify another?
 - Consider a simple experiment
 - `twc really-big-file tiny-file`
 - The `twc` program would use `pthread_wait` to wait for the first thread to finish and then for the second to finish. Is this OK?
 - If the second thread finishes its work (e.g., counting) early, how can it notify the first thread that it has the counting results?
 - In processes, the `wait/exit` system calls are used to know termination and its status of any children.
 - However, threads don't work that way; they don't have a parent to notify

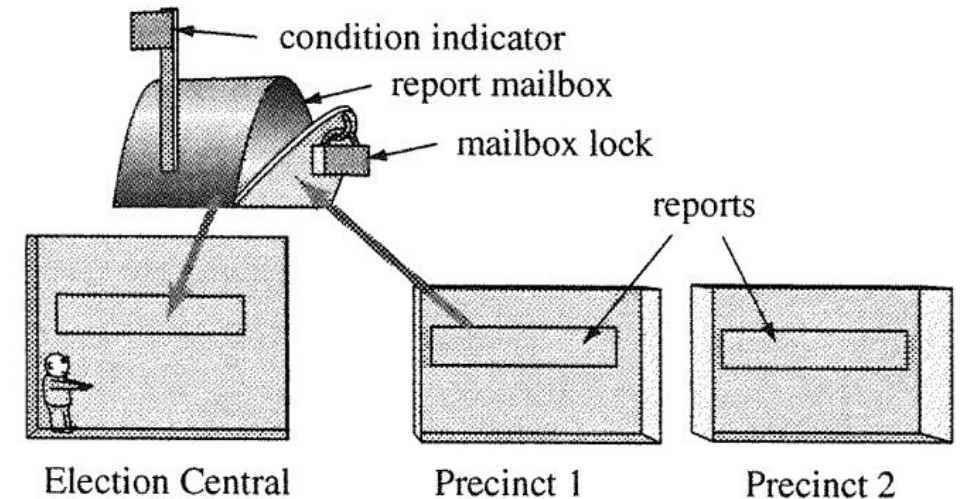
Notifying Election Central

- Analogy: Counting votes

- Each Precinct delivers a report to the mailbox and notifies EC
- EC then consumes the report
- EC waits for the flag to be signaled

while Precinct (선거구) can
raise—sends a signal to—the flag

- * precinct: an electoral district of
a city or town



< Using a locked mailbox to transfer data >

- This mailbox can only store one certified vote count at a time
 - The polling places produce vote counts and election central consumes vote counts
 - Only one vote count can be stored at a time

Notifying Election Central (cont.)

- (a) Election Central (EC) sets up a vote report mailbox
 - This mailbox has space for only one vote report at a time
 - This mailbox has a flag that can be raised and then snaps right back
 - This mailbox has a mutex that can be locked and unlocked
- (b) EC unlocks the box and waits until the flag is signaled
- (c) Voting place (VP) waits until it can lock the mailbox
 - If the mailbox is not empty, VP unlocks the mailbox and waits until the flag is signaled before it can lock the mailbox again
 - VP now puts the voting report in the mailbox

Notifying Election Central (cont.)

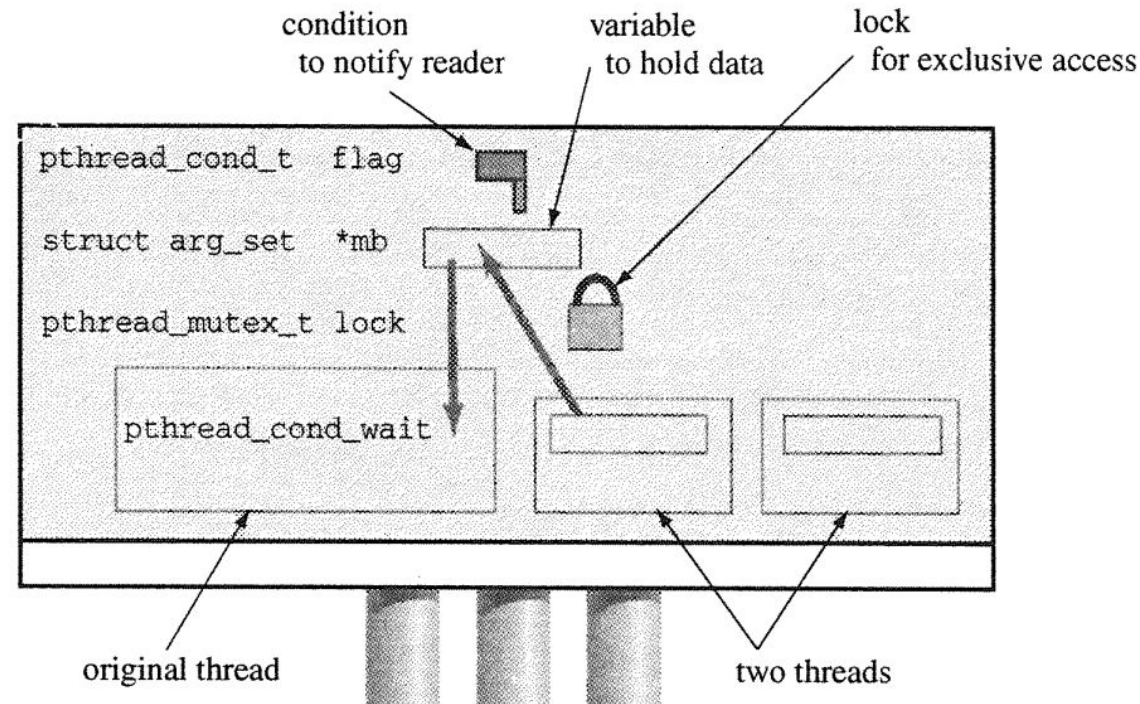
- (d) VP signals the flag on the mailbox
 - VP releases the lock on mailbox
- (e) EC stops waiting because the flag was signaled
 - EC locks the mailbox, takes the voting report from the mailbox, and processes the data on the voting report
 - EC then signals the flag in case a VP is waiting
 - EC returns to Step (b)

Condition Variables: wait and signal

- A condition variable is an object threads use to notify other threads
 - A thread ***waits for notification*** with `pthread_cond_wait()`
 - e.g., wait for signal on condition variable
 - A thread ***sends notification*** with `pthread_cond_signal()`
 - e.g., signal that you are done

Example 2-4: twordcount4.c

- Version 4: “Three Threads, Three Variables”
 - Programming with *conditional variables*
 - Let’s use a “locked (conditional) variable” to transfer data



< Using a locked variable to transfer data >

Example 2-4: `twordcount4.c` (cont.)

- Logic of the program
 - **The main program** (original thread) launches the “two counting threads” and then waits for results to come in.
 - Calls `pthread_cond_wait` to wait for the flag to be signaled
 - This call blocks the original thread.
 - When a **counting thread** finishes counting, it is ready to deliver the result by storing a “pointer” in the mailbox.
 - 1) That **counting thread** has to acquire a lock for the mailbox.
 - 2) It checks the mailbox once the lock is obtained.
 - If not empty, then the thread unlocks the mailbox and waits for the flag to be signaled before locking the mailbox again.
 - 3) The thread puts the result into the mailbox.
 - 4) The thread signals the condition variable flag by calling `pthread_cond_signal`

Example 2-4: `twordcount4.c` (cont.)

- Signaling this flag wakes up the original thread (or, the main program)
 - Note that it was blocked on that flag by calling `pthread_cond_wait`
 - The (awakened) original thread rushes to open the mailbox and tries to obtain a lock for that mailbox, but the lock is still held by the counting thread.
- When the **counting thread** releases the lock with `pthread_mutex_unlock`, the original thread gets the lock.
 - After holding the lock, the original thread now:
 - Takes the report out of the mailbox
 - Reports the result to the screen and adds the number to its total
 - Signals the flag in case a counting thread is waiting, and
 - Loops back to call `pthread_cond_wait`, which atomically unlocks the mutex and blocks the thread until the flag is signaled (by the counting thread) again

Example 2-4: twordcount4.c (cont.)

```
/* twordcount4.c - threaded word counter for two files.
 *
 *           - Version 4: condition variable allows counter
 *           functions to report results early
 */

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <ctype.h>

struct arg_set {
    /* two values in one arg*/
    char *fname; /* file to examine */
    int count; /* number of words */
};

struct arg_set *mailbox = NULL;
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t flag = PTHREAD_COND_INITIALIZER;
```

Example 2-4: twordcount4.c (cont.)

```
int main(int ac, char *av[])
{
    pthread_t    t1, t2;        /* two threads */
    struct arg_set args1, args2; /* two argsets */
    void         *count_words(void *);
    int          reports_in = 0;
    int          total_words = 0;

    if ( ac != 3 ){
        printf("usage: %s file1 file2\n", av[0]);
        exit(1);
    }

    args1.fname = av[1];
    args1.count = 0;
    pthread_create(&t1, NULL, count_words, (void *) &args1);

    args2.fname = av[2];
    args2.count = 0;
    pthread_create(&t2, NULL, count_words, (void *) &args2);
```

```
    while( reports_in < 2 ){
        printf("MAIN: waiting for flag to go up\n");
        pthread_cond_wait(&flag, &lock); /* wait for notify */
        printf("MAIN: Wow! flag was raised, I have the lock\n");
        printf("%7d: %s\n", mailbox->count, mailbox->fname);
        total_words += mailbox->count;
        if ( mailbox == &args1 )
            pthread_join(t1, NULL);
        if ( mailbox == &args2 )
            pthread_join(t2, NULL);
        mailbox = NULL;
        pthread_cond_signal(&flag); /* announce state change */
        reports_in++;
    }
    printf("%7d: total words\n", total_words);

    return 0;
}
```

Example 2-4: twordcount4.c (cont.)

```
void *count_words(void *a)
{
    struct arg_set *args = a;      /* cast arg back to correct type */
    FILE *fp;
    int c, prevc = '\0';

    if ( (fp = fopen(args->fname, "r")) != NULL ){
        while( ( c = getc(fp)) != EOF ){
            if ( !isalnum(c) && isalnum(prevc) )
                args->count++;
            prevc = c;
        }
        fclose(fp);
    } else
        perror(args->fname);
    printf("COUNT: waiting to get lock\n");
    pthread_mutex_lock(&lock);      /* get the mailbox */
    printf("COUNT: have lock, storing data\n");
    if ( mailbox != NULL ){
        printf("COUNT: oops..mailbox not empty. wait for signal\n");
        pthread_cond_wait(&flag,&lock);
    }
    mailbox = args;                 /* put ptr to our args there */
    printf("COUNT: raising flag\n");
    pthread_cond_signal(&flag);     /* raise the flag */
    printf("COUNT: unlocking box\n");
    pthread_mutex_unlock(&lock);    /* release the mailbox */
    return NULL;
}
```


Example 2-4: twordcount4.c (cont.)

- Execution

```
dynam@DESKTOP-Q4IJB7:~/lab13$ make
cc -o twc4 twordcount4.c -lpthread
dynam@DESKTOP-Q4IJB7:~/lab13$ ./twc4 /etc/group /usr/share/dict/words
MAIN: waiting for flag to go up
COUNT: waiting to get lock
COUNT: have lock, storing data
COUNT: raising flag
COUNT: unlocking box
MAIN: Wow! flag was raised, I have the lock
      249: /etc/group
MAIN: waiting for flag to go up
COUNT: waiting to get lock
COUNT: have lock, storing data
COUNT: raising flag
COUNT: unlocking box
MAIN: Wow! flag was raised, I have the lock
    131756: /usr/share/dict/words
    132005: total words
```


Example 2-4: Functions for Condition Variables - `pthread_cond_wait`

| pthread_cond_wait | | |
|--------------------------|---|---------------------------------|
| PURPOSE | Blocks a thread on a condition variable | |
| INCLUDE | #include <pthread.h> | |
| USAGE | int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex); | |
| ARGS | cond | pointer to a condition variable |
| | mutex | pointer to a mutex |
| RETURNS | 0 | if successful |
| | errcode | if not successful |

Example 2-4: Functions for Condition Variables - pthread_cond_signal

| pthread_cond_signal | |
|----------------------------|--|
| PURPOSE | Unblocks a thread waiting on a condition variable |
| INCLUDE | #include <pthread.h> |
| USAGE | int pthread_cond_signal(pthread_cond_t *cond); |
| ARGS | cond pointer to a condition variable |
| RETURNS | 0 if successful errcode if not successful |

Threads vs. Processes

| Processes | Threads |
|--|---|
| Supports <i>parallelism</i> ; <i>multitasking</i> | |
| <ul style="list-style-type: none">• Part of the Unix since the beginning | <ul style="list-style-type: none">• Added later |
| <ul style="list-style-type: none">• Model of the process is “clear and uniform.” | <ul style="list-style-type: none">• Evolved from variety of sources• Different type of threads with different attributes• <i>POSIX threads</i> used in our examples |
| <ul style="list-style-type: none">• Have “own” data space, file descriptors, and process ID number | <ul style="list-style-type: none">• “Share” one data space, set of file descriptors, and process ID number |

Threads vs. Processes (cont.)

| Processes | Threads |
|--|--|
| <ul style="list-style-type: none">• Heavyweight• Low throughput• Long response time• High IPC overhead via pipes, sockets, signal, exit/wait• Little synchronization overhead by no shared variables• Easier to debug | <ul style="list-style-type: none">• Lightweight• High throughput• Short response time• Low IPC overhead via global variables• Much synchronization overhead by mutex and conditional variables• Harder to debug |
| <ul style="list-style-type: none">• <code>fork()</code> can create a new process• <code>wait()</code> can wait for a child to complete. | <ul style="list-style-type: none">• <code>pthread_create()</code> can create a new thread• <code>pthread_join()</code> waits for a thread to be done. |
| <ul style="list-style-type: none">• <code>ps</code> can list processes | <ul style="list-style-type: none">• <code>ps</code> cannot list threads |

More Stuff about Threads

- 1) Shared Data Space
 - Multiple threads can read a large, complex tree-structured database in memory easily
 - Multiple queries from clients can be served from one process
 - Variables do not change; so no problems!
 - But consider another example: a program using `malloc` and `free` to manage memory.
 - One thread: allocates a chunk of memory to store a string
 - Another (different) thread: calls `free` to deallocate that chunk of memory.
 - Then what happens to the original thread?

More Stuff about Threads (cont.)

- 2) Shared File Descriptors
 - `fork` automatically duplicate the file descriptors to the child processes
 - So if the child process closes a file descriptor inherited from the parent, that file descriptor still open for the parent
 - In a multi-threaded program, it's possible to pass the same file descriptor to *two different* threads
 - Both those values refer to the same file descriptor
 - If one thread *closes* the file descriptor, then that file descriptor is *closed* for all threads in the same process although other threads need that connection

More Stuff about Threads (cont.)

- 3) `fork`, `exec`, `exit`, and `signals`:
 - All threads “share” the same process
 - If one thread calls `exec`, that thread will surprise the other threads by replacing the current program with a new one
 - If one thread calls `exit` or crashes, the “entire” process finishes or dies
 - If one thread calls `fork`, only that thread is running in the new process
 - No other threads get duplicated in the new process

Summary

- A thread of execution is the flow of control through a program
 - The pthreads library allows a program to run several functions at the same time
- Functions running at the same time have their own local variables
 - But they share all global variables and dynamically allocated data
- When threads share a variable, they need to make sure they don't get in each other's way
 - Threads may use a mutex lock to make sure only one thread is using a shared variable at a time

Summary (cont.)

- When threads need to coordinate or synchronize their action, they may use a condition variable.
 - One thread waits for the condition variable to change in a specific way
 - The other thread signals the variable to change.
- Threads need to use a mutex lock to prevent simultaneous access to functions that operate on shared resources.
 - Functions that are not reentrant must be protected this way

Appendix

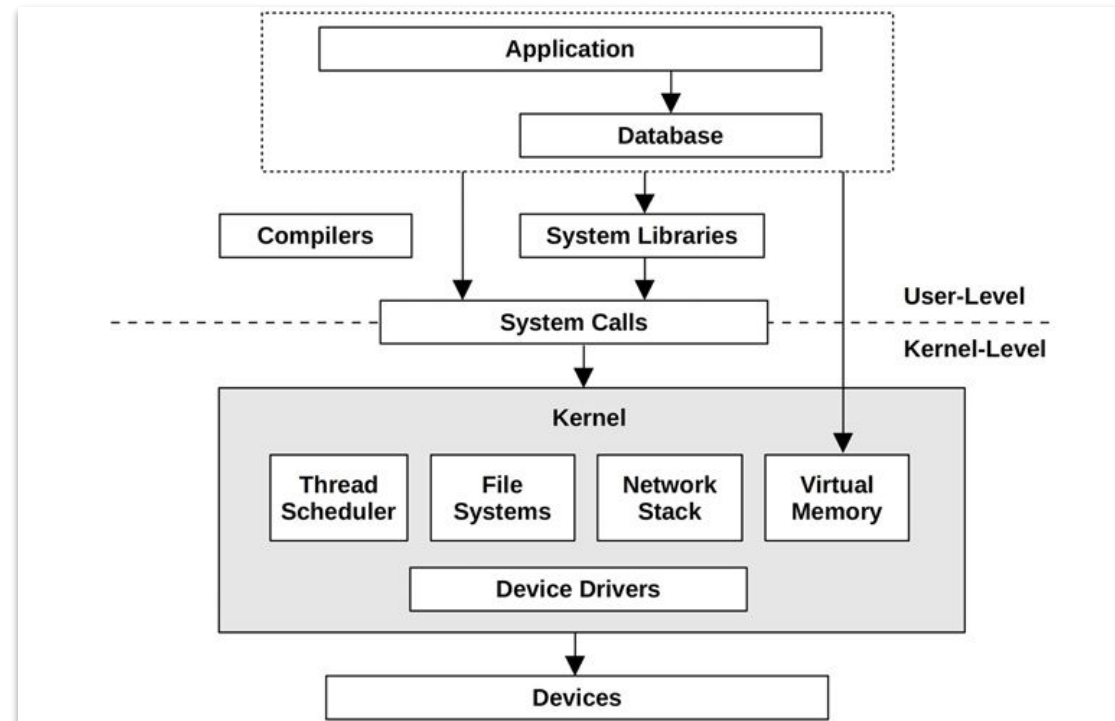
Systems Performance Intro

Introduction

- Systems performance
 - studies the performance of an entire computer system
 - including all major software and hardware components
- Data path from storage devices to application software
 - If you don't have a diagram of your environment showing the data path, find one or draw it yourself
- Goal
 - to improve the end-user experience by **reducing latency**
 - to reduce computing cost
 - can be achieved by eliminating inefficiencies, improving system throughput, and general tuning

Introduction (cont.)

- Full stack (in terms of systems performance)
 - The entire software stack from the *application* down to the *hardware*
 - including system libraries, the kernel, and the hardware itself



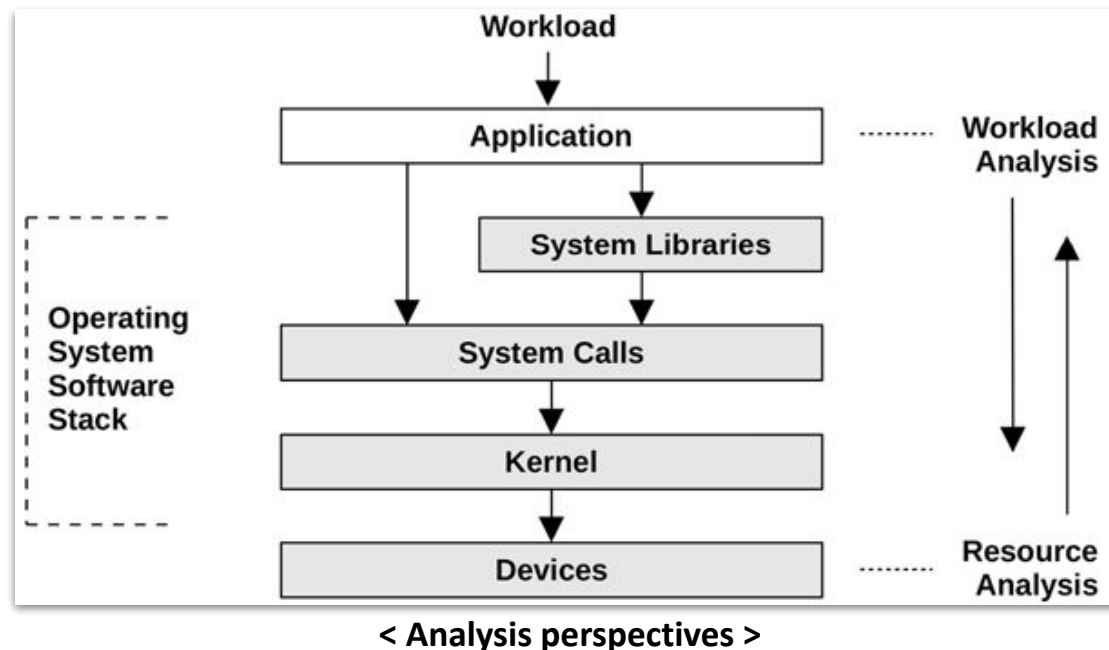
< Generic system software stack on a single server >

Introduction (cont.)

- Systems performance is done by a variety of job roles
 - including system administrators, site reliability engineers, application developers, network engineers, database administrators, web administrators, and other support staff
 - For some performance issues, finding the root cause or contributing factors requires a *cooperative effort* from more than one team

Two Perspectives

- Workload analysis
 - By application developers for the delivered performance of the workload
- Resource analysis
 - By system administrators, who are responsible for the system resources



Performance is Challenging

- Subjectivity (주관성)
 - e.g., The average disk I/O response time is 1 ms. Is this “good” or “bad”?
- Complexity
 - Complexity of systems and the lack of an obvious starting point for analysis
- Multiple causes
- Multiple performance issues
 - The real task *isn't finding* an issue; it's **identifying** which issue or issues matter *the most*